

Contents

Known bugs and v1.0 limitations	2
v1.0 known limitations (workarounds documented; v1.1 candidates)	2
.dta round-trip loses xtset state	2
ARIMA uses conditional likelihood, not exact ML (no Kalman filter)	2
Factor variable composition with TS ops	2
xtreg, be doesn't weight by T_i	2
MA(1) and ARMA SE precision	3
Variable-name length limit (32 chars)	3
Graphics: covered since v1.6.6	3
CHANGELOG (recently fixed)	3
Recently fixed	3
Tier 4 — econometric completion (NEW for v1.0)	3
Tier 3 — output and reproducibility (NEW)	4
Tier 2 — factor variables: i., c., #, ##, ib. (NEW)	5
Tier 1 polish — statistical functions, egen extensions, _b/_se, formats	6
margins — marginal effects after regress / logit / probit (NEW)	7
predict extended to logit, probit, and xtreg (NEW)	7
logit and probit — binary-outcome MLE (NEW)	8
test syntax extended (NEW)	9
lincom accepts TS-op coefficient names (NEW)	9
xtreg, re — random-effects (GLS) estimator (NEW)	9
hausman — FE vs RE specification test (NEW)	9
xtreg, fe — fixed-effects (within) estimator (NEW)	10
TS-op handling unified across all varlist-accepting commands (NEW)	10
xtdescribe / xtides (NEW)	11
TS operators in summarize and list (FIXED)	11
D-operator iterated differencing (FIXED — silent math bug)	11
Operator-list TS syntax (FIXED)	11
Multi-digit time-series operators read wrong lag (FIXED)	12
Time-series operators rejected in regress varlist (FIXED)	12
capture and _rc (FIXED in v1.4.0)	12
test 08 captures fewer lines than expected	12
Storage types are display-only	13
Numeric function family coverage	13
v0.6 .dta polish (this milestone)	13
Still deferred from v0.6	13
reported by user during real-data testing	13
reported by Claude during synthetic testing	14
reported by Mico during WB fertility replication (v1.6.x)	14
reported by Mico during WB fertility replication, round 2 (v1.6.2)	14
v1.6.3 — Stata-parity import naming + WASM refresh	15
v1.6.4 — quality-of-life extensions (not Stata, non-conflicting)	15
reported by Mico during WEO test_04 (v1.6.5)	15
v1.6.6 — graphics round two (the fertility-paper figures)	16
v1.6.7 — value-label attachments survive frame rebuilds (Mico's fig1)	17
v1.6.8 — quality of life: conversion-phase progress for import excel	17
v1.6.9 — period-date constructors + eval errors abort (OECD HPI round)	17
v1.6.10 — set more on off does something now	18
v1.6.11 — cellrange() + case() for import excel (WPP fertility round)	18
v1.6.12 — tabstat: value-labeled groups + format() (test_01_b round)	18

v1.6.13 — long variable names no longer wreck sum/describe/tabstat	19
v1.6.14 — documentation catch-up	19

Known bugs and v1.0 limitations

v1.0 known limitations (workarounds documented; v1.1 candidates)

These are not bugs in the strict sense — they’re scope decisions for v1.0 that we may revisit based on real-world testing.

.dta round-trip loses xtset state

save mydata.dta then use mydata.dta preserves variable types, formats, variable labels, and value labels, but does **not** preserve the xtset panel/time variable assignment. Stata writes this as a dataset characteristic (`_dta[iis]`, `_dta[tis]`); readstat doesn’t expose characteristic writing in its public API, and hand-writing characteristic blocks in raw .dta format is a v1.1 task.

Workaround: re-run `xtset country year` after use. This is the standard Stata-do-file convention anyway, so most users won’t notice.

ARIMA uses conditional likelihood, not exact ML (no Kalman filter)

tea’s `arima` command computes parameters by minimizing the SSR of the recursively-computed residuals, conditioning on $\varepsilon_{\{<0\}}=0$ (the conditional likelihood approach). This gives consistent estimates but slightly different finite-sample behavior than Stata’s default exact ML via the Kalman filter:

- Absolute log-likelihood values differ from Stata’s (we drop the initial-state contribution)
- AR coefficients near the unit-root boundary may converge less robustly
- No seasonal ARIMA (`sar`, `sma`) in v1.0

Workaround: for serious time-series work, escape to R’s `arima()` or Python’s `statsmodels.tsa.arima`.

Factor variable composition with TS ops

`i.country#c.L.gdp` (factor interaction with a lagged regressor) is **not** supported directly. The factor module operates on column names, not on materialized TS-op temps.

Workaround: pre-compute the lag manually: `gen L1_gdp = L.gdp`, then use `i.country#c.L1_gdp`.

xtreg, be doesn’t weight by T_i

Stata’s default `xtreg, be` uses a slight variance correction for unbalanced panels (down-weights groups with few observations). tea uses straight OLS on the panel means. For balanced panels the answers are identical. For unbalanced panels they can differ by a few percent in the SEs.

MA(1) and ARMA SE precision

ARIMA standard errors are computed from a numerical-differentiation Hessian (forward differences scaled by 1e-5). For models with $q > 0$ (MA components), the SEs can be off by 1-5% compared to analytical or exact-ML Hessians. Point estimates are unaffected.

Variable-name length limit (32 chars)

A residual Stata-12 limit. Factor-variable interaction names like `2010.year#1985.country#c.gdp` (28 chars) fit; longer 3-way interactions truncate. In practice not a problem.

Graphics: covered since v1.6.6

`scatter / line / histogram` (v1.3), plus multi-series `twoway` with `lowess`, `graph box` with `two-level over()`, and the `name()` registry with `graph combine` (v1.6.6). Remaining gaps: no `by()` faceting, no `bar/pie`, no graph editor.

CHANGELOG (recently fixed)

Recently fixed

Tier 4 — econometric completion (NEW for v1.0)

ivregress 2sls (in `src/regress.c` ~300 lines):

Syntax: `ivregress 2sls y [exog] (endog = instruments) [, vce(robust|cluster v)]`. Full first-stage regression of each endogenous regressor on the combined instrument matrix $W = [\text{exog}, Z, _cons]$ via `dgels`, second-stage $\beta = (\hat{X}'\hat{X})^{-1}\hat{X}'y$; residuals computed using ORIGINAL X (not fitted) for SE; sandwich $V = A \cdot B \cdot A$ for robust/cluster ($A = (\hat{X}'\hat{X})^{-1}$, B uses original- X residuals). HC1 / CR1 finite-sample adjustments. First-stage F-stat reported (minimum across endogenous regressors — the weak-instrument worry case).

Verified on textbook endogeneity case: with $\text{cov}(x, u) \neq 0$, OLS biased to 2.49; IV recovers 1.997 against true 2.0, first-stage $F=1364$.

poisson — count outcomes via the existing MLE driver:

Added `mle_family_poisson` to `src/mle.c` with `poisson_per_obs`: $\mu = \exp(\eta)$ clipped at $\eta < 700$; `score = y - \mu`; `weight = \mu`; `loglik = y\eta - \mu` (drops the $\log(y!)$ constant — absolute `loglik` differs from Stata but β identical). Added `count_validate` (rejects $y < 0$). Generalized `do_glm` binary header text to “Poisson regression”. Verified intercept-only: $\beta_cons = \log(\bar{y}) = 0.6931$ for $\bar{y} = 2$ (exact).

xtreg, be (between estimator):

Replaced the “not implemented” stub. After group-mean computation (already done in the FE/RE path), branch into BE: build $G_obs \times (K_orig + 1)$ matrix Xb (panel means + appended 1s column for cons), OLS via `dgels`, $\sigma^2_be = \text{RSS} / (G - K)$, $V = \sigma^2 \cdot (\hat{X}'\hat{X})^{-1}$ via Cholesky. Verified hand-checkable case: panel data with $\bar{y}_i = 5 + 2\bar{x}_i$ recovers $\beta_x = 2$, $_cons = 5$, $R^2_b = 1.0000$ exactly.

arima (new module `src/arima.c`, ~580 lines):

Conditional-likelihood ARIMA(p,d,q) with optional ARIMAX exog regressors. Syntax: `arima y [exog], arima(p d q) [noconstant]`. Algorithm: - Difference y in place d times

- Pack params as $[\mu, \beta_{K_ex}, \varphi_p, \theta_q]$ - Recursive residual computation with $\varepsilon_{\{<0\}}=0$ conditional assumption - Gauss-Newton iteration with step halving (max 50 iterations) - SSR via finite-difference gradient + Hessian - $V \approx 2\sigma^2 \cdot \text{Hess}^{-1}$ via `dgetrf/dgetri` with ridge

Verified: AR(1) on strong-signal data recovers $\varphi=0.7010$ (true 0.7). AR(1) on US WEO growth data matches OLS lag regression to 3 digits ($\hat{\varphi} = 0.154$ both). MA(1) Monte Carlo with `rnormal()` noise: $\theta=0.429$ (true 0.5) at $N=500$.

runiform() and rnormal() (in `eval.c`):

PCG32-based PRNG with Box-Muller normal. `set seed N` for reproducibility. Essential for Monte Carlo, simulation, and bootstrap workflows. Tested against expected distributional moments: mean 0.492, sd 0.291 for u (vs true 0.5, 0.289); mean 0.001, sd 1.015 for z (vs true 0, 1).

Tier 3 — output and reproducibility (NEW)

estimates command suite (new module `src/estcmd.c`):

```
estimates store NAME      clone last_est, save under name
estimates restore NAME    load NAME → last_est (for predict/margins/test)
estimates dir             list saved with cmd, N, depvar
estimates drop NAME...    remove named entries; `_all` clears
estimates table NAMES     compact side-by-side comparison table
                        (options: se, t, p, star, stats(N r2 rmse F ...))
```

Aliased as `est`. Workspace gains a linked-list field `stored_est` (in `dataset.h`); freed by `ws_free`. Built on the existing `est_clone` infrastructure from the `hausman` work.

estout — publication-ready tables (new module `src/estout.c`):

```
estout [names] [using FILE] [, format(latex|markdown|md|plain)
                        se|t|p stars stats(...) nogaps
                        title("...") label(tab:foo)
                        keep(varlist) drop(varlist)
                        nomtitles mtitles("M1" "M2" ...)]
```

Default format is LaTeX (per the IMF / academic-economics convention). LaTeX output produces a tabular block with significance stars in the superscript style ($\hat{\alpha}^{***}$), variable names properly escaped, and a footnote legend. Stata-style `\_cons` for the constant term. Markdown and plain text variants also supported.

Verified end-to-end on a three-model comparison with the standard options: stars apply z -distribution thresholds (10/5/1%), SEs go in parentheses below each coef, stats footer rows handle $N/r^2/r^2_a/rmse/F/df_r/df_m/sigma_u/sigma_e/rho$.

log audit and fix:

The `log` command captures output via a `printf`-macro override in `commands.c`. Problem: `display` lives in `interp.c` and was using bare `printf`, so its output went to `stdout` but not the log file. Fixed by making `g_logfp` extern and adding a manual tee from `do_display`. Round-trip verified: `log using file.log` now captures both command echoes (`. display ...`) AND their output (`_b[x] = 2`).

.dta round-trip audit:

Spot-checked save → use preservation: - Variable types: `%8.2f, %td`:
`%` - Variable labels (label variable x "GDP per capita"): `%` - Value labels (label define,

label values): `□` - Numeric values: `□` (down to floating-point precision) - String columns: `□`

Not preserved in `.dta`: `xtset` panel/time variable assignment. Stata writes this as a dataset characteristic (`_dta[iis]`, `_dta[tis]`); `readstat` doesn't expose characteristic writing in its public API, and hand-writing the characteristic block in raw `.dta` format is a v1.1 project. Workaround for v1.0: re-run `xtset country year` after use, which is also the standard Stata-do-file convention.

Tier 2 — factor variables: `i.`, `c.`, `#`, `##`, `ib.` (NEW)

The Stata factor-variable grammar, now supported across every estimator that goes through `tsop_expand_varlist` (`regress`, `xtreg fe/re`, `logit`, `probit`, `predict`, `margins`). Lives in new module `src/factor.[ch]`.

Grammar:

```
i.var          dummies for each non-base level (base = smallest by default)
ib<n>.var      same, but base = level n
c.var          explicit continuous (display prefix; mostly for use in #)
A#B           interaction: cross product of A's and B's columns
A##B          main effects + interaction: equivalent to "A B A#B"
```

Examples (assume `country` has levels {1,2,3}, `year` has {2010,2011,2012}):

```
i.country      → 2 cols: "2.country" "3.country"
ib2.country    → 2 cols: "1.country" "3.country"
i.country#c.gdp → 2 cols: "2.country#c.gdp" "3.country#c.gdp"
i.country#i.year → 4 cols: "2.country#2011.year" "2.country#2012.year"
                  "3.country#2011.year" "3.country#2012.year"
i.country##c.gdp → 5 cols total: 2.country, 3.country, gdp,
                  2.country#c.gdp, 3.country#c.gdp
```

Architecture:

- `factor.c` is a sibling to `tsop.c`: tokens are parsed into atoms (kind 'T' for `i.`, 'C' for `c.`, 'L' for coef-name "2.x", or bare variable in a `#` context). Atoms join into terms via `#`. `##` is unrolled to the power-set of single-term products (main effects + all interactions).
- `factor_is_factor_token` is the dispatch test; `factor_expand_token` produces a snapshot-able list of temp frame columns.
- Plumbed into `tsop_expand_varlist`: each token is tried as a TS-op first; if that returns "not mine," tried as a factor-var; if that returns "not mine," falls through to plain `varlist_expand`.

Predict/margins re-materialisation (the subtle bit):

When `predict` sees a coefficient named `2.country` or `2.country#c.gdp`, the column doesn't exist anymore — it was a temp dropped after the estimator finished. Solution: treat coef-name form `<level>.<varname>` as a new atom kind 'L' (level-pinned indicator). When `tsop_expand_varlist` is called with `"2.country"`, it dispatches to `factor`, which produces a single column with values `1{country==2}`. Same mechanism handles `2.country#c.gdp` (indicator × continuous).

Verified:

Spec	Hand-checkable result
<code>regress y x i.country on y = $\alpha_c + x$</code>	$\beta=1, _cons=\alpha_1, 2.country=\alpha_2-\alpha_1, \dots$

Spec	Hand-checkable result
<code>regress y x ib2.country</code>	<code>_cons=α₂, 1.country=α₁-α₂</code>
<code>regress y c.x#c.x on y = x²</code>	<code>coef = 1.0 exactly</code>
<code>regress y i.country##c.x on y = α_c + β_c·x</code>	<code>reproduces all α and β exactly</code>
<code>regress growth L.growth i.cid (LSDV) on WEO</code>	<code>matches xtreg fe exactly</code>

`predict after logit y x i.g` now works (was the open bug from the previous session). `margins, dydx(*)` after factor-var estimation works.

Restrictions in v1.0:

- `i.var` requires `var` to be numeric. Use `encode` for string ids.
- Composition with TS-ops (e.g., `i.country#c.L.gdp`) is not supported. Manually pre-compute the lag (`gen L1_gdp = L.gdp`) and use `c.L1_gdp`.
- Cell-count cap of 10000 columns per token (huge interactions blocked).
- Variable-name length limit of 32 chars means very long interaction names truncate. In practice not a problem for typical use.

Tier 1 polish — statistical functions, egen extensions, _b/_se, formats

A pass over the everyday-use surface in preparation for v1.0. An audit showed that tea's expression-language function library was already very comprehensive (math, trig, hyperbolic, strings, dates, distributions, glob+regex, value/variable labels, encode/decode/destring/tostring, recode, format). The gaps that mattered:

New statistical functions (in `src/eval.c`): - `normalden(x)` / `normalden(x,σ)` / `normalden(x,μ,σ)` — normal PDF - `lnnormal(x)` — `log Φ(x)`, stable for large negative `x` via the asymptotic expansion (reuses `tea_log_normal_cdf`) - `lnnormalden(x)` — `log φ(x)` closed-form

Better regex (in `src/eval.c`): - `regexs(n)` — extract `n`-th submatch from the last `regexm` call. Uses a file-scope `tea_regex_submatch[20]` buffer. Index 0 is the whole match; 1..N are capture groups. - `regexpr(s, pat, rep)` — replace first regex match with a string

Extended egen (in `src/commands.c`): - New aggregators: `median`, `iqr`, `p25`, `p75`, `p50`, `pc` (with option `p(N)` to specify percentile), `rank`, `tag` - Multi-argument row functions now work: `rowtotal(v1 v2 v3)`, `rowmean`, `rowmin`, `rowmax`, `rowsd`, `rowmiss`, `rownonmiss` (the previous implementation was single-arg only) - `tag(g1 g2)` marks the first row of each group as 1, others 0 - `rank(x)` uses average-rank-for-ties

Format application (in `src/commands.c::fmt_cell`): Custom printf-style numeric formats set via `format var %w.dF` are now honored by `list` (previously only date formats like `%td` were).

Coefficient/SE macros after estimation (in `src/regress.c` + `src/interp.c`):

`_b[name]` and `_se[name]` are now accessible everywhere `e(N)` is. The estimators (`regress`, `xtreg`, `logit`, `probit`) store macros keyed by `_b[varname]` and `_se[varname]` (including bracketed names like `_b[L.growth]` for TS-op regressors). `macro_expand` recognises the pattern and substitutes outside double quotes; inside quotes the text stays literal. Standard Stata behaviour.

This unlocks idioms like:

```
quietly regress y x
display "t-stat = " _b[x] / _se[x]
gen yhat = _b[_cons] + _b[x]*x
```

Bug fix — quoted-string macro substitution (in src/interp.c): display "e(N) = " e(N) used to produce "3 = 3" because the macro expander substituted both occurrences of e(N), including the one inside the quoted string. Fixed by tracking in_dquote state and skipping e() / r() / _b[] / _se[] substitution inside strings. Stata's behaviour preserved: backtick-locals and \$globals are still substituted inside strings, since those are Stata's quote-aware substitution syntax.

margins — marginal effects after regress / logit / probit (NEW)

```
margins, dydx(*)           AME for all continuous regressors
margins, dydx(x1 x2)       AME for specified regressors
margins, dydx(*) atmeans   MEM (effects at sample means)
margins, dydx(x1)          AME for x1 only
```

After OLS, the AME of x_k is exactly β_k and the SE is exactly $SE(\beta_k)$ — the margins output reproduces the coefficient table. Useful as a sanity check.

After logit / probit, AMEs are the economically meaningful number: how the predicted *probability* changes with x , averaged across the sample. Coefficients themselves aren't directly interpretable (they're log-odds for logit, latent-index units for probit).

Math:

Logit: $m_k = (1/N) \sum_i \Lambda(X_i\beta) \cdot (1 - \Lambda(X_i\beta)) \cdot \beta_k$ Probit: $m_k = (1/N) \sum_i \varphi(X_i\beta) \cdot \beta_k$ OLS: $m_k = \beta_k$

Delta-method SE: $V(m_k) = G_k \cdot V(\beta) \cdot G_k'$, where G_k is the K-vector $\partial m_k / \partial \beta_j = (1/N) \sum_i [g'(\eta_i) X_{ij} \beta_k + g(\eta_i) 1\{j=k\}]$. Uses the full $V(\beta)$ matrix from the previous estimation (so robust/cluster SEs flow through automatically).

Verified: - Logit and probit on the same data produce nearly identical AMEs (a known result — they only differ at the tails, which contribute little to the average). In one test case $AME(\text{logit}, x1) = 0.12237$, $AME(\text{probit}, x1) = 0.12253$ — different in the fifth decimal. - AME for OLS reproduces the regress coefficient table exactly. - TS-op coefficient names (L.x etc.) work via the same snapshot-then-drop trick predict uses.

Deferred for v2: - margins x1 (predicted outcomes at varying x1 — useful for plots) - margins, at(x1=(0 1) x2=10) (predictions at specified values) - Discrete-change effects for indicator variables - Marginal effects after xtreg (would need within-effect-aware formulas)

predict extended to logit, probit, and xtreg (NEW)

The old predict worked only after regress, with options xb and residuals. Rewritten to dispatch on the previous estimator's cmd:

```
After regress:
  predict yhat           xb (default)
  predict r, residuals   y - Xβ
  predict s, stdp        standard error of the linear prediction

After logit / probit:
  predict p              predicted probability (default)
  predict idx, xb        linear index Xβ
```

After `xtreg (fe or re):`

<code>predict yhat</code>	<code>xb</code> (default; does NOT include <code>u_i</code>)
<code>predict u, u</code>	$\hat{\alpha}_i$ (panel effect)
<code>predict e, e</code>	idiosyncratic residual $y - X\beta - \hat{\alpha}_i$
<code>predict ue, ue</code>	$\hat{\alpha}_i + e$ ($= y - X\beta$)
<code>predict xbu, xbu</code>	$X\beta + \hat{\alpha}_i$ (in-sample prediction)

The dispatch lives in `predict_resolve_kind` and the per-row evaluation is one big switch. Sensible errors when options don't match the preceding command (e.g. `predict p, pr` after `regress` fails cleanly).

TS-op coefficient names: when `e->xnames[j]` is `L.growth` (TS-op form), the column doesn't exist in the frame anymore at predict time, but we can re-materialize it via the shared `tsop_expand_varlist`. A subtlety: we have to **snapshot** the materialized data into our own buffer, drop the temps, THEN add the new prediction variable — otherwise `tsop_drop_temps` would drop the just-added prediction column instead of the temp. This is what an earlier draft got wrong (a regression test would have shown a list-not-found error).

xtreg specifics: $\hat{\alpha}_i$ is computed in one pass over the in-sample rows (`e->used` tells us which) by streaming through the panel-sorted data and accumulating $\hat{y} - \bar{x}'\beta$ per panel. For predictions on rows that weren't in the estimation sample, $\hat{\alpha}_i$ is unknown so `u/e/ue/xbu` return missing.

Verified on the hand-checkable FE case (3 panels $\times$ 4 obs, perfect fit) where the math is exact: $\beta=1$, $\alpha=10/20/30$, $e=0$, $xbu=y$. Also tested on WEO AR(1) growth where `predict` reproduces the fit `yhat + r = growth` exactly.

logit and probit — binary-outcome MLE (NEW)

```
logit y x1 x2 ... [if] [in] [weight], [vce(robust|cluster v)]
probit y x1 x2 ... [if] [in] [weight], [vce(robust|cluster v)]
```

Implemented via a Newton-Raphson MLE driver in `src/mle.[ch]` with family-specific callbacks for the score, weight, and log-likelihood contributions. Same driver should extend to `poisson`, `cloglog`, etc. with no changes to the iteration logic — just new `MleFamily` structs.

Output matches Stata's layout: iteration log, LR χ^2 , pseudo R^2 , coefficient table with z-stats, optional cluster adjustment note.

Standard errors: - (default): classical $(X'WX)^{-1}$ (asymptotic inverse-Hessian) - `vce(robust)`: HC1 sandwich - `vce(cluster v)`: CR1 cluster-robust on score

Numerical care: - $\log(1+e^\eta)$ computed stably as $\max(\eta, 0) + \log(1+e^{-|\eta|})$ — no overflow. - Probit's ϕ/Φ uses asymptotic Mills bound $\lambda(\eta) \rightarrow -\eta$ when $\Phi(\eta)$ underflows. Per-obs Fisher weight clipped at $1e8$. - Starting values: OLS solve, which also detects rank deficiency (collinear regressors marked omitted, held at 0 through iteration). - Step halving (up to 8 halvings) when full Newton step decreases $\square$. - Perfect-separation detected when $\square \rightarrow 0$; warning printed.

Verified: - Intercept-only on 50/50 sample: $\beta_{\text{cons}} = 0$ exactly. SE = 0.2 for logit, 0.1253 for probit (matches closed forms). - Real-effect data: classical $\beta_{\text{logit}}/\beta_{\text{probit}} \approx 1.81$ ratio reproduced. - Newton converges in 3-5 iterations for well-conditioned data.

Postestimation: `test` and `lincom` work on logit/probit estimates. (`test` uses F with $df_r = N-K$; the asymptotic test would be χ^2 , but the difference is negligible for typical sample sizes.)

test syntax extended (NEW)

The test command now supports the full Stata-style hypothesis syntax:

```
test var                H0:  $\beta_{\text{var}} = 0$ 
test var = 0            same, explicit
test var = 0.5          H0:  $\beta_{\text{var}} = 0.5$  (numeric RHS)
test var1 = var2        H0:  $\beta_{\text{var1}} = \beta_{\text{var2}}$  (equality)
test var1 var2 ...      joint H0: all listed = 0
```

Internally constructs the restriction matrix R and value vector r , then computes the Wald $F = (R\beta - r)'(RVR')^{-1}(R\beta - r)/q$ against $F(q, df_r)$. Pre-existing test L.growth L2.growth (joint zero) is a special case and still works.

lincom accepts TS-op coefficient names (NEW)

The lincom parser now treats . as part of a coefficient name when followed by a name character, so lincom L.growth + L2.growth works correctly. Previously the parser stopped at the first . and tried to look up “L” as a coefficient.

xtreg, re — random-effects (GLS) estimator (NEW)

Implements panel RE via feasible GLS with quasi-demeaning:

```
xtreg y x1 x2 ... [if] [in], re [vce(robust|cluster v)]
```

Method: the model is $y_{it} = \alpha_i + x_{it}'\beta + \varepsilon_{it}$ with $\alpha_i \sim (0, \sigma_u^2)$ assumed independent of x_{it} . The estimator transforms each variable by subtracting θ_i times the panel mean, where:

$$\theta_i = 1 - \sigma_e / \sqrt{(T_i \cdot \sigma_u^2 + \sigma_e^2)}$$

Then runs OLS on $(y, X, 1-\theta_i)$ — note an explicit constant column, unlike FE. σ_e^2 comes from the within (FE) regression; σ_u^2 comes from the variance components computation ($\max(0, \text{var}(\hat{\alpha}_i) - \sigma_e^2/\bar{T})$). The transform interpolates between OLS ($\theta=0$) and FE ($\theta=1$).

Output matches Stata’s xtreg, re layout: - “Random-effects GLS regression” header - R-within, R-between, R-overall — all computed with $\beta_{\text{RE}} - \sigma_u, \sigma_e, \rho$ — variance components (shared with FE math) - θ — quasi-demeaning factor; min/avg/max for unbalanced panels - $_cons$ coefficient and SE (unlike FE) - Wald χ^2 for the slopes - $\text{corr}(u_i, X) = 0$ noted as the RE assumption

Verified on: - Degenerate hand-checkable case (perfect within fit $\rightarrow \theta=1 \rightarrow$ RE collapses to FE; $\beta_{\text{RE}} = \beta_{\text{FE}} = 1.0$ exactly; $_cons$ omitted since the $1-\theta$ column is identically zero). - WEO AR(1) growth: $\beta_{\text{RE},L} = 0.245$ sits between pooled OLS and FE $\beta = 0.220$. $\sigma_u/\sigma_e/\rho = 0.97/5.37/0.032$ — most variance within- panel. θ varies 0.19-0.39 across 210 countries.

v1 limitations: - Uses $\sigma_u^2 = \max(0, \sigma_{\text{BE}}^2 - \sigma_e^2/\bar{T})$ with $\bar{T} = N/n$ for unbalanced. Stata’s exact Swamy-Arora uses a slightly different $\bar{T}$ correction; for balanced panels the estimators agree.

hausman — FE vs RE specification test (NEW)

```
xtreg y x1 x2 ..., fe
xtreg y x1 x2 ..., re
hausman
```

Computes $H = (\beta_{FE} - \beta_{RE})' [V_{FE} - V_{RE}]^{-1} (\beta_{FE} - \beta_{RE})$, distributed $\chi^2(K_{common})$ under H_0 : $cov(\alpha_i, x_{it}) = 0$. Rejection means RE is inconsistent; use FE.

Tea saves the last xtreg fe and last xtreg re into dedicated workspace slots automatically, so the user just types hausman with no arguments after running both. Order doesn't matter.

Implementation notes: - Only compares slope coefficients present (and not omitted) in both estimates. `_cons` is dropped since FE doesn't have one. - The difference matrix $V_{FE} - V_{RE}$ should be PSD under H_0 , but sampling noise (especially with robust/cluster SEs) can make it not strictly PD. Cholesky inverse is tried first; if it fails, falls back to an SVD pseudoinverse. When the resulting statistic comes out negative, we report $|H|$ with a "sign flipped" warning. - Verified on WEO AR(1) growth: $\chi^2(1) = 189$, $p < 0.0001 \rightarrow$ strongly rejects RE (matches the FE-reported $corr(u_i, Xb) = 0.97$ warning sign). On a constructed case where $\alpha_i \perp x$ by construction, $\chi^2 \approx 0$ and the test correctly fails to reject.

xtreg, fe — fixed-effects (within) estimator (NEW)

Implements panel FE OLS via within transformation:

```
xtreg y x1 x2 ... [if] [in], fe [vce(robust|cluster v)]
```

Reports R-within, R-between, R-overall, σ_u , σ_e , ρ , $corr(u_i, X\beta)$, and the F-test that all $u_i = 0$. Verified on a hand-checkable case (tests/regression/29_xtreg_fe.do): 3 panels $\times$ 4 obs, $y = \alpha_i + x$ with $\alpha \in \{10, 20, 30\}$ gives $\beta_{FE} = 1.0$ exactly, $R_{within} = 1.0$, $\sigma_u = 10$ exactly. Verified against the WEO AR(1) growth dynamic showing internally consistent results (β from xtreg D.growth L.growth, fe = $-(1 - \beta$ from xtreg growth L.growth, fe)).

Standard error options: | , fe | classical: $\sigma^2 \cdot (X_w' X_w)^{-1}$ with $df = N - n_{groups} - K$
 | | , fe vce(robust) | HC1 sandwich on within-transformed data | | , fe vce(cluster v)
 | CR1 cluster-robust on within-transformed data |

v1 limitations (deferred to v2): - `_cons` is not displayed. Its point estimate (mean of α_i) is computable but the strictly correct SE requires LSDV-equivalent variance propagation we haven't implemented. Users wanting inference on `_cons` can run `regress y x i.panel` for now. - Stata defaults `vce(robust)` to cluster-by-panel for xtreg fe; tea's `vce(robust)` is HC1. Use explicit `vce(cluster panelvar)` for cluster-by-panel. - xtreg, be (between effects) is not yet implemented — returns a clean error suggesting , fe or , re. (, re ships in this version — see above.) - Weights work the same way as in regress (build_design handles them), but the FE-specific weight semantics have not been carefully verified yet.

TS-op handling unified across all varlist-accepting commands (NEW)

Time-series operator support is now centralized in `src/tsop.[ch]`. The parser handles every Stata form once (`L.x`, `L2.x`, `L(1/2).x`, `L.(x y)`, `L2.(x y)`, `L(1/2).(x y)`, step form `L(1(2)9).x`), and every command that accepts a varlist routes through `tsop_expand_varlist`. Concretely:

Command	TS ops	Mechanism
regress	□	tsop_expand_varlist (depvar + xspec)
summarize	□	tsop_expand_varlist
list	□	tsop_expand_varlist
tabulate	□	tsop_expand_varlist
tabstat	□	tsop_expand_varlist
collapse	□	tsop_expand_varlist (paren-aware tokenizer)

regress lost ~400 lines of duplicated parser code in this refactor (from 1198 down to 796 lines). Future commands that accept varlists — xtreg, ivregress, correlate, etc. — get TS-op support for free by calling `tsop_expand_varlist`. Test 28 exercises the universal coverage; tests 21-27 cover individual command surfaces.

Side fix discovered: keep if / drop if used to call `frame_unsort` which over-eagerly cleared `ts_panel / ts_time`. Row deletion doesn't invalidate panel structure (the columns are still there and still sorted; gaps are fine for tea's gap-aware `tsidx` lookup), so the ts state now survives those operations.

xtdescribe / xtdes (NEW)

Added `xtdescribe` (with `xtdes` abbreviation) to summarize the panel structure declared by `xtset`. Shows number of panels, time range, delta, span, min/p25/p50/p75/max distribution of obs per panel, and balanced/unbalanced determination. Test 27 exercises both balanced and unbalanced cases.

TS operators in summarize and list (FIXED)

`summarize L.x` and `summarize D.x` failed with “variable not found” because those commands called `varlist_expand` directly and didn't recognize TS-op tokens. The earlier TS-op fix only touched `regress`.

Fixed by adding a shared `tsop` module (`src/tsop.[ch]`). Both `summarize` and `list` now route their varlist through `tsop_expand_varlist`, which materializes TS-op tokens as temporary frame columns (with canonical display names like `L.x`, `D2.y`) and appends them to the frame for the duration of the command, then drops them via `tsop_drop_temps` before returning. Test 26 covers this.

D-operator iterated differencing (FIXED — silent math bug)

`D2.x`, `D3.x`, etc. all returned the same as `D.x` because the eval code hardcoded `back = -1` for any D-kind operator. This was a **silent** numerical bug: regressions including `D2.x` would compute correctly for `D.x` but the higher-order differences were all collinear with `D.x` and got “omitted” — looking like a model specification issue rather than a math error.

Stata's convention is that `D#.x` is the iterated #-th difference:

- $D.x = x[t] - x[t-1]$
- $D2.x = D(D.x) = x[t] - 2 \cdot x[t-1] + x[t-2]$
- $D3.x = x[t] - 3 \cdot x[t-1] + 3 \cdot x[t-2] - x[t-3]$
- $D^k x[t] = \sum_{j=0..k} (-1)^j \cdot C(k,j) \cdot x[t-j]$

Fixed in `src/eval.c` using a binomial-recurrence loop. The `S` operator (separate semantics: simple gap-aware seasonal difference, $S#.x = x[t] - x[t-#]$) is unchanged. Test 25 verifies on $x = t^2$ where `D.x` is linear, `D2.x` is constant 2, `D3.x` is constant 0.

Action item for users: any analysis using `D2.x`, `D3.x`, or higher in past sessions silently produced wrong values. Re-run any such code.

Operator-list TS syntax (FIXED)

Stata's compact form `regress y L(1/2).x` (expanding to `L1.x L2.x`) was rejected with “independent variables not found”. Tea now handles the full set of Stata TS-op forms in `regress` varlists:

Syntax	Expands to
<code>L.x</code>	<code>L.x</code>
<code>L2.x</code>	<code>L2.x</code>
<code>L0.x</code>	<code>x</code> (no shift)
<code>L(1/3).x</code>	<code>L.x L2.x L3.x</code> (range)
<code>L(1 3).x</code>	<code>L.x L3.x</code> (explicit list)
<code>L(1(2)9).x</code>	<code>L.x L3.x L5.x L7.x L9.x</code> (step form)
<code>L.(x y)</code>	<code>L.x L.y</code>
<code>L2.(x y)</code>	<code>L2.x L2.y</code>
<code>L(1/2).(x y)</code>	<code>L.x L.y L2.x L2.y</code> (cross product)

Implemented in `src/regress.c` via the unified `try_tsop_form` parser plus `parse_numlist` (range / explicit list / step) and `find_matching_paren` for nested parens. The regressor tokenizer is paren-aware so spaces inside (...) don't break tokenization. Tests 21, 23, and 24 cover the forms.

Multi-digit time-series operators read wrong lag (FIXED)

`L2.x` was actually computing the 12th lag instead of the 2nd; `L3.x` the 13th, `F2.x` the 12th lead, `S3.x` a 13-period seasonal difference. Affected expressions, `regress`, and any other consumer.

The bug was in `try_tsop` in `src/parse.c`: the lag-magnitude accumulator was initialized to 1 instead of 0, so reading “L2” gave `num = 1*10 + 2 = 12`. Lag-1 (`L.x` / `L1.x` / no digits) accidentally worked because `num` stayed at 1 either way.

Fixed by initializing `num = 0`. The existing `if (!have) num = 1`; fallback still gives bare `L.` the right value of 1. Test 22 pins down correct multi-digit operator behavior.

Time-series operators rejected in regress varlist (FIXED)

`regress growth L.growth if country_id == "USA"` failed with “regress: independent variables not found” because the regress code called `varlist_expand` directly on each regressor token; that helper only knew about plain names, wildcards, ranges, and `_all` — not the `L.x` / `F.x` / `D.x` / `S.x` syntax.

Fixed by refactoring `build_design` in `src/regress.c` to detect TS-op tokens and route them through the existing expression evaluator (`expr_parse` → `N_TSOP`). Plain names still go through `varlist_expand`. Works for both the dependent variable and the regressors (`regress D.growth L.growth` is supported). Test 21 covers `regress` with TS operators.

capture and _rc (FIXED in v1.4.0)

Fixed: `capture` now suppresses all output of the captured command (stdout and stderr, at the file-descriptor level) and `_rc` returns the captured return code — 0 on success, the error code otherwise — exactly as in Stata. One small remaining deviation: a successful command that is NOT under capture leaves `_rc` unchanged rather than resetting it to 0; branch on `_rc` only immediately after a capture.

test 08 captures fewer lines than expected

The reshape-error regression test passes, but if you read the expected output you'll notice some error messages print to stderr while others to stdout, and the line ordering across

the merged stream depends on line-buffering. This is OK as of v0.5.20 because we set the stdout buffer to line-buffered in `main.c`, but the harness assumes the ordering is stable — if a future change reorders error/output emission between commands, the test may need updating.

Storage types are display-only

`gen byte x = 1` accepts the byte qualifier but stores the value as double internally; `gen strl0 country = "USA"` accepts `strl0` but doesn't track width. This is a deliberate design decision documented in `COMPATIBILITY.md` — tea is double + string internally regardless of declared type, and the `.dta` writer auto-compresses on save. Not a bug.

Numeric function family coverage

Many Stata numeric functions are not yet implemented:

- Distribution functions: `normal`, `normalden`, `invnormal`, `chi2`, `chi2tail`, `invchi2`, `t`, `ttail`, `invttail`, `F`, `Ftail`, `binomial`, `binomialp`, `poisson`, `poissonp`.
- Matrix functions (we don't have a matrix type yet).
- Random-number generators: `runiform`, `rnormal`, etc. (Tied to `set seed`.)

The existing function set covers basic arithmetic, transcendentals, date/time, string manipulation, and the common aggregator functions used in `egen`.

v0.6 .dta polish (this milestone)

- **Value labels** — `label define` / `label values` sets are now written into `.dta` and read back via `use`. Round-trips cleanly (test 16).
- **Dataset label** — `label data "text"` now persists. Visible in `describe`, round-trips through `save/use` (test 17).
- **save FILE.dta, version(NUM)** — emit any DTA format 104-119 (test 18). Defaults to 118 (Stata 14).
- **Workspace label-set cleanup** — `clear` and `use`, `clear` now also drop the workspace's value-label sets, matching Stata semantics.

Still deferred from v0.6

- **Notes** — Stata's note: text on a dataset or variable. Not yet read or written.
- **strL strings > 2045 bytes** — currently strings longer than the `str#` limit are truncated at 2045 on write. Real-world risk is essentially zero for econ data; will add when needed.
- **.dta fweight metadata writes are silently dropped** — `readstat 1.1.9` accepts `readstat_writer_set_fweight_variable` but does not serialize the resulting Stata `_fweight` characteristic. Tea's read-side captures the metadata correctly when present (so real Stata-saved files round-trip through tea correctly *if* the user doesn't save through tea), and the write-side wiring is in place for when `readstat` fixes this. Track upstream.

reported by user during real-data testing

- Bug 1 — `_N` in display context returned 1 → **FIXED in v0.5.20**
- Bug 2 — `tabstat columns()` ignored in `by`-mode → **FIXED in v0.5.19**
- Bug 3 — `tabstat` with `if` + `by` showed empty `by`-groups → **FIXED in v0.5.18**

- Bug 4 — list * segfaulted on 147-var dataset → **FIXED in v0.5.14**
- Bug 5 — quoted filename in import not opened → **FIXED in v0.5.12**

reported by Claude during synthetic testing

- Bug 6 — commas in string literals broke parser → **FIXED in v0.5.20**
- Bug 7 — \$1 in strings expanded as undefined macro → **FIXED in v0.5.20**
- Bug 8 — egen silently skipped strings → **FIXED in v0.5.20**

reported by Mico during WB fertility replication (v1.6.x)

- Bug 9 — import excel, firstrow header naming didn't match Stata: invalid chars were underscored (Country Code → Country_Code) instead of removed (CountryCode), and digit-leading headers (1960) got a v prefix instead of Stata's Excel column-letter fallback (F...BQ). Do-files written against Stata failed with variable not found. → **FIXED**: excel/ods import now uses Stata's exact rule; duplicate headers fall back to the column letter.
- Bug 10 — foreach v of varlist BAD was a **silent no-op**: an unknown variable made varlist_expand return -1, which the loop treated as an empty list. The body ran zero times with no error, and scripts failed much later with a misleading message. → **FIXED**: hard error, rc=111. The same -1-as-empty flaw was guarded at every remaining silent site: by BAD: (ran ungrouped!), egen ..., by(BAD) (ungrouped), collapse, by(BAD) (would collapse the whole dataset to one row), format % BAD (silent no-op). All now rc=111.
- Bug 11 — local ++yr parsed ++yr as the macro *name*, silently defining a macro literally called ++yr while yr never changed — loop counters froze and rename loops died with “already exists” on the second iteration. → **FIXED**: local ++x / local --x now increment/decrement per Stata; ++ on an undefined or non-numeric macro errors (rc=198) instead of silently minting a new macro.

reported by Mico during WB fertility replication, round 2 (v1.6.2)

- Bug 12 — reshape long silently **DROPPED every non-stub variable**: after reshape long y, i(CountryCode IndicatorCode) j(year) the carried columns (IndicatorName, CountryName, ...) simply vanished, and the do-file died six lines later with a misleading keep: variable not found. → **FIXED**: carried variables are replicated across the j-rows of each wide observation, exactly as Stata does. Formats and variable labels survive on i() and carried columns.
- Bug 13 — reshape had a fixed 512-level cap that **silently dropped** level 513 onward (real WB pulls have ~1,900 indicator codes). → **FIXED**: dynamic sorted level arrays, binary-search everywhere.
- Bug 14 — reshape wide read j as numeric unconditionally, so j(code_safe) string was impossible. → **FIXED**: string j in both directions. New loud errors, never silent mangling: string j without the string option (rc 109), generated names that aren't valid identifiers (rc 198, with a strtoname() hint), missing/empty j values (rc 498), carried variables not constant within i() (rc 9), duplicate j within an i() group (rc 9 — previously last-write-wins, silently).
- Bug 15 — merge dropped variable labels (formats were kept). → **FIXED**: labels survive merge on master and using columns.
- Bug 16 — frame_set_nobs **leaked every string cell on shrink**: any row-deleting operation (drop if, keep if, duplicates drop) leaked the dropped rows' strings, and a later regrow overwrote the stale pointers, making it permanent. Invisible until

reshape began carrying string columns into frames that then get row-pruned; caught by ASan. → **FIXED** in dataset.c.

- NEW — preserve / restore implemented (was on the roadmap): single depth like Stata, snapshot on disk as a native .tea tempfile so large frames don't double peak memory. restore, not discards; restore, preserve reloads but keeps the snapshot. A preserve still pending when a do-file concludes — normally or via abort — is restored automatically, with a note.
- NEW — strtoname() string function, Stata-exact: invalid characters become _, a leading digit gets a _ prefix, 32-char cap.

v1.6.3 — Stata-parity import naming + WASM refresh

- import excel WITHOUT firstrow now matches Stata exactly: columns are named by their Excel letters (A, B, ...) and row 1 is kept as a DATA row. (Previously tea always consumed row 1 as headers, silently losing the first data row.) Mixed columns (text header over numbers) become string columns, as in Stata.
- import delimited now applies Stata's naming rule: names are LOWERCASED by default and invalid characters are REMOVED ("Country Code" -> countrycode, not Country_Code); a header that is empty or starts with a digit falls back to the position name v#; duplicate headers also resolve to v#. New option case(preserve|lower|upper) for the Stata-compatible opt-out. NOTE: this is a deliberate behavior change — do-files that relied on tea's old underscored, case-preserving CSV names need case(preserve) or updated names. The bundled sysuse datasets are unaffected (already lowercase-clean).
- web/tea.wasm rebuilt — the browser engine is current again (it had been frozen at the v1.6.0 feature set while native moved ahead).
- Bug 17 — tempfile names were deterministic (/tmp/tea_tmpN_name), so any do-file using the tempfile + file open, write idiom **failed on its own second run** with rc=602 ("file already exists") — and leftover files accumulated in /tmp. → **FIXED**: paths now embed the pid (unique per session, like Stata's), and every tempfile handed out is deleted when tea exits.

v1.6.4 — quality-of-life extensions (not Stata, non-conflicting)

- NEW — status: one-line dataset summary — source filename, obs, vars, exact in-memory size, sort and xtset state. Source is tracked through use/import/sysuse, updated by save, cleared by clear (new Frame.source field, set at the command layer so internal frame operations like the reshape swap or restore never disturb it).
- NEW — progress indicator on long operations (import, reshape long, reshape wide, merge). Time-gated: nothing is drawn until ~1s of wall time, redraws at most every 250ms, stderr-and-TTY-only, erased on completion — so logs, pipes, the regression harness, capture, and short operations see nothing, ever. set progress off to disable. Compiled to no-ops in the WASM build (no TTY; a JS hook can come later).
- Bug 17 refinement: the tempfile uniqueness token is now pid ^ nanosecond-startup-time, because pid alone fails under emscripten (constant pid, and the node harness mounts the HOST /tmp) — caught by the WASM rig colliding with its own previous run.

reported by Mico during WEO test_04 (v1.6.5)

- Bug 18 — if COND { evaluated against an empty scratch frame with _N hardwired to 1 — so if _N > 0 { was ALWAYS true and if _N == 0 { ALWAYS false, independent of

the data. The diagnostic guard blocks in the WEO script ran unconditionally. Silently wrong control flow — the worst class. → **FIXED**: conditions evaluate against the active frame (`_n=1`, `_N` = real obs count); variable references like `x[1] == 5` now work in if-conditions too.

- Bug 19 — for values `i = 1/_N` ran the body `**ZERO times silently**`: the range was parsed before macro expansion, `sscanf` failed on the backtick, and `1/0` looped zero times with no error. → **FIXED**: the range is macro-expanded first, and an unparseable range is a loud `rc=198` (a legitimately empty range like `1/0` still runs zero times, as in Stata).
- Bug 20 — compound quotes `"..."` were unsupported in macro expansion: the expander treated `"` as a macro reference and swallowed the text. → **FIXED** per Stata's rules: `"` opens, `"` closes, nesting tracked, macros still expand inside, a lone `"` inside is literal. `display, local x "..."`, and label variable accept them; plain `"..."` strings in `display` and label variable also accept the `""` doubled-quote escape.
- Bug 21 — label variable stripped ALL trailing quotes from the label, mangling labels that contain quotes. → **FIXED** with a real quoted-string parser (see Bug 20).
- NEW — quietly `{ ... }` / capture `{ ... }` / noisily `{ ... }` block forms (prefixes chain: capture quietly `{ ... }`). quietly now also suppresses display, as in Stata (prefix and block forms both).
- NEW — extended macro function `local x : subinstr local|global y "from" "to" [, all]`, with compound-quoted arguments (the way to say a literal double quote). Other extended functions error loudly with the supported list rather than silently assigning the text.
- NEW — error `#` (abort with a return code; the do-file assertion idiom `if (bad) { ... error 459 }`) and `compress` (accepted for do-file compatibility; tea storage is already minimal — reports `"(0 bytes saved)"`).
- NEW — `display as error|as text|as result|as input` and old-style `display in red|green|...` are recognized as style directives and no longer printed as literal text.
- Auto-restore of a pending `preserve` now also fires for do-files run from the command line (tea `script.do`), not just via the `do` command; `preserve` snapshots joined the `atexit` cleanup net.

v1.6.6 — graphics round two (the fertility-paper figures)

- NEW — **multi-series twoway**: parenthesized series each with their own `if` and options; plot types `scatter / line / connected / lowess`. Per-series `lcolor()` `lpattern()` `msymbol(i)` `mlabel()` `mlabcolor()` `mlabposition()` (clock); globals `ylines(#,...)` (repeatable), `legend(off)` (simple legend otherwise), `yscale(range())`, `ylabel(a(s)b)` / `xlabel(a(s)b)`, `note()`. Axis titles inside a series merge upward, as in Stata. Tick rules and plot range are separate: a `ylabel()` rule never clips data.
- NEW — **lowess**: tricube kernel, running-line (default) or mean, `bwidth()` (0.8), `adjust`. Pure C, byte-identical across rigs.
- NEW — **graph box**: Stata percentile interpolation, adjacent-value whiskers (1.5 IQR), outside values (`noout` hides), one- or two-level `over()` (first varies fastest inside bands of the second), value labels honored, `relabel()`, `label(angle() labsize())`.
- NEW — **named-graph registry + graph combine**: `name(NAME[,replace])` stores the SVG in-session AND writes `NAME.svg` (documented deviation); `graph combine N1 N2, cols()/rows()` nests panels as scaled SVG; `graph dir / graph drop NAME|_all`. Name collision without `replace` is `rc 110`.
- Everything renders on the same deterministic SVG engine: `%2f` coordinates with `-0` snapped, golden-file diffed byte-identically on native, ASan, and WASM rigs (tests 56, 57). `plot.c` and the test-40 goldens are untouched.
- Documented deviations (COMPATIBILITY.md): `name()` writes a file; `twoway`

line/lowess series are sorted by x; unknown cosmetic graphics suboptions are accepted-and-ignored (strict everywhere else).

v1.6.7 — value-label attachments survive frame rebuilds (Mico’s fig1)

- Bug 22 — **encoded group variables silently reverted to raw numerics** in graph box band labels: reshape (all four carried-variable paths), merge (master and using), and frame copy copied format+vlabel but dropped Variable.vallab. All eight copy sites fixed.
- Bug 23 — the native format stored only the per-variable attachment NAME, not the label-set contents, so save→clear→use via .tea (and the preserve snapshot format) lost definitions. The format is now **TEA2**: value-label sets referenced by attached variables travel with the dataset, exactly like Stata’s .dta. Legacy TEA1 files still read (without attachments); old tea versions cannot read TEA2.
- Bug 24 — merge ... using FILE defaulted a missing extension to .tea while use/save default to .dta — so savef'+merge ... using f' on an extensionless tempfile failed with “cannot readtea”. Merge now defaults to .dta, matching use/save and Stata.
- Test 58 locks all five survival paths: reshape wide+long, preserve/restore, .tea round-trip, .dta round-trip, and merge — asserted via rendered graph box band labels.

v1.6.8 — quality of life: conversion-phase progress for import excel

- import excel on a large workbook was silent for most of its wall time: the ssconvert/libreoffice child ran under a blocking system() with no feedback, and only the subsequent CSV-parse phase was byte-instrumented — so the progress line appeared “around 70%”. The converter now runs under a polled fork/exec with a new progress ACTIVITY mode: spinner + elapsed seconds (“converting WB_data.xlsx / 34s”), same contract as the percentage line (stderr TTY only, 1s activation gate, 250ms redraws, erased on completion, set progress off disables, no-op on WASM). Batch output is byte-identical.

v1.6.9 — period-date constructors + eval errors abort (OECD HPI round)

- NEW — string period-date constructors: quarterly(s, "YQ"), monthly(s, "YM"), halfyearly(s, "YH"), weekly(s, "YW"), yearly(s, "Y"), and daily(s, mask) as the Stata alias of date(). Integer tokens are extracted left-to-right and assigned by the mask; two-digit years get +2000 (matching date()); an out-of-range period (“2020-Q7”) yields missing, never a wrong date. The rest of the date family (mdy/dofq/yofd/qofd/... and %tq/%td/%tm display formats) already existed.
- Bug 25 — **a runtime evaluation error in gen/replace did not abort**: the error printed once per row, missing was stored anyway, each store counted as a “real change”, and the command returned 0. A do-file with gen qdate = quarterly(...) under a build lacking the function marched on: collapse grouped on an all-missing year, merge matched ZERO rows, and the garbage panel was saved without a nonzero rc anywhere. Now: an eval error at the type probe fails before the variable exists; an error mid-loop aborts with rc=133 and rolls back completely — gen leaves no trace of the new variable, replace restores the column from a snapshot. Partial writes never survive.

v1.6.10 — set more on|off does something now

- `set more` had been accepted-and-ignored since v1.0 — a silent no-op: `set more on` armed nothing, and a stray `list` on a 400k-row panel flooded the terminal. Now it is Stata's output pager for real: with `set more on`, long `list` and `describe` output pauses at a screenful with `--more--`; space/any key = next page, Enter = one more line, q = stop with `--Break--`. Terminal height is read live (TIOCGWINSZ, fallback 24 rows). It engages ONLY in the interactive REPL with stdout a TTY — do-files, pipes, capture, logs, and the test suite never see it, so batch output is byte-identical. Default is off, matching modern Stata. `set more banana errors (198)`, as does any other malformed argument — same strictness as `set progress`.

v1.6.11 — cellrange() + case() for import excel (WPP fertility round)

- NEW — `import excel ..., cellrange(A17:AF22000)`: restricts the import to a sheet rectangle. Real workbooks (the UN WPP file) carry title junk above the table; without `cellrange`, junk row 1 became the header. With `firstrow`, the range's first row is the header row. The end corner is optional (`cellrange(A17)` = from A17 to the sheet's end). The slice happens on the converted CSV with a fully CSV-aware state machine — quoted fields containing commas, doubled quotes, or embedded newlines slice correctly.
- NEW — `import delimited ..., rowrange(r1[:r2]) colrange(c1[:c2])`: Stata's rectangle restriction for text files, sharing the same slicer (and giving the regression suite a WASM-testable path — the `xlsx` conversion needs `ssconvert`, absent on that rig).
- Bug 26 — `case("lower")` with quotes, legal Stata syntax, ERRORED on import delimited and was **silently ignored** on import excel, which had no `case()` handling at all (names hardcoded to preserve). Both branches now accept quoted or bare arguments; the excel naming rule (invalid chars removed, column-letter fallback) folds case AFTER cleaning, and the excel default remains preserve, matching Stata.
- Test 60 locks range slicing (bounded + start-only), quoted-field survival, quoted `case()` args, and loud errors on malformed ranges.
- WASM-rig lesson #3 (joins Bug 17 and the test-57 note): emscripten NODEFS misbehaves after `rename()` — subsequent opens in the same directory can fail on stale cache state. The slicer therefore never renames; it writes a `.rng` sibling that callers load and unlink. Slice temp names are per-call unique (`pid ^ nanotime`), not per-pid.

v1.6.12 — tabstat: value-labeled groups + format() (test_01_b round)

- Bug 27 — `tabstat ..., by(ctr_group)` named groups by raw numeric value (1/2/3) instead of the attached value labels (Advanced/Emerging/LIC). Same lookup rule as `list` and `graph box` now: labels when attached, %g fallback otherwise; string by-variables unchanged.
- Bug 28 — `format(%5.2f)` was ignored entirely — `tabstat` had no `format()` parsing and every cell went through a hardcoded `%.7g`. Now: the option applies to every cell INCLUDING the Total rows (the Total block carries a second copy of the cell-format macro, which the first fix pass missed — caught because the verification run showed a formatted table with an unformatted Total). Quoted `format("%5.2f")` accepted. Only printf-safe numeric formats pass validation (%w.d followed by f/g/e/F/G/E); anything else — `format(%s)`, `format(banana)` — is a loud 198, since an arbitrary string reaching `snprintf` as a format is undefined behavior.
- Test 61 locks labeled groups, formatted group and Total rows in both table orientations, the quoted form, and both rejection paths.

v1.6.13 — long variable names no longer wreck sum/describe/tabstat

- Bug 29 — the WPP panel’s 26-character names (totalpopulationasof1januar, ...) overflowed the fixed 12-char name column in `summarize` and the name columns/headers in `tabstat`, and shoved `describe`’s type/format columns out of alignment. Fixes: `summarize` and `tabstat` use Stata’s `abbrev()` display rule (first n-2 characters + ‘~’ + the last character: totalpopul~r), applied to variable names, column headers, and by-group labels alike; `describe` pads its name column dynamically to the longest name in view (16..32), header row included. Frames with short names render byte-identically to before — the whole existing suite passes against unchanged goldens.
- Test 62 locks the abbreviation rule in `sum`, `tabstat` headers and by-labels (including long `VALUE` labels), and `describe`’s dynamic alignment.

v1.6.14 — documentation catch-up

- The manual and the generated command reference had lagged three releases of features: the `import` chapter now documents `cellrange()` (with the header-at-row-17 semantics), `rowrange()/colrange()`, `quoted-and-bare case()`; the `date-function` table gains the period constructors (quarterly/monthly/halfyearly/weekly/yearly and the daily alias); `import`’s and `tabstat`’s dispatch help strings (the source of the generated reference) now describe their full current option surface, including `tabstat`’s `format()` and value-labeled `by()` groups. No behavior changes.