

Contents

Known bugs and v1.0 limitations	2
v1.0 known limitations (workarounds documented; v1.1 candidates)	2
.dta round-trip loses xtset state	2
ARIMA uses conditional likelihood, not exact ML (no Kalman filter)	3
Factor variable composition with TS ops	3
xtreg, be doesn't weight by T_i	3
MA(1) and ARMA SE precision	3
Variable-name length limit (32 chars)	3
Graphics: covered since v1.6.6	3
CHANGELOG (recently fixed)	3
Recently fixed	3
Tier 4 — econometric completion (NEW for v1.0)	4
Tier 3 — output and reproducibility (NEW)	4
Tier 2 — factor variables: i., c., #, ##, ib. (NEW)	5
Tier 1 polish — statistical functions, egen extensions, _b/_se, formats	6
margins — marginal effects after regress / logit / probit (NEW)	7
predict extended to logit, probit, and xtreg (NEW)	8
logit and probit — binary-outcome MLE (NEW)	9
test syntax extended (NEW)	9
lincom accepts TS-op coefficient names (NEW)	9
xtreg, re — random-effects (GLS) estimator (NEW)	9
hausman — FE vs RE specification test (NEW)	10
xtreg, fe — fixed-effects (within) estimator (NEW)	10
TS-op handling unified across all varlist-accepting commands (NEW)	11
xtdescribe / xtides (NEW)	11
TS operators in summarize and list (FIXED)	11
D-operator iterated differencing (FIXED — silent math bug)	12
Operator-list TS syntax (FIXED)	12
Multi-digit time-series operators read wrong lag (FIXED)	12
Time-series operators rejected in regress varlist (FIXED)	13
capture and _rc (FIXED in v1.4.0)	13
test 08 captures fewer lines than expected	13
Storage types are display-only	13
Numeric function family coverage	13
v0.6 .dta polish (this milestone)	14
Still deferred from v0.6	14
reported by user during real-data testing	14
reported by Claude during synthetic testing	14
reported by Mico during WB fertility replication (v1.6.x)	14
reported by Mico during WB fertility replication, round 2 (v1.6.2)	15
v1.6.3 — Stata-parity import naming + WASM refresh	15
v1.6.4 — quality-of-life extensions (not Stata, non-conflicting)	16
reported by Mico during WEO test_04 (v1.6.5)	16
v1.6.6 — graphics round two (the fertility-paper figures)	17
v1.6.7 — value-label attachments survive frame rebuilds (Mico's fig1)	17
v1.6.8 — quality of life: conversion-phase progress for import excel	18
v1.6.9 — period-date constructors + eval errors abort (OECD HPI round)	18
v1.6.10 — set more on off does something now	18
v1.6.11 — cellrange() + case() for import excel (WPP fertility round)	18
v1.6.12 — tabstat: value-labeled groups + format() (test_01_b round)	19

v1.6.13 — long variable names no longer wreck sum/describe/tabstat	19
v1.6.14 — documentation catch-up	19
v1.6.15 — command echo for the browser editor (colleague feedback round) . . .	20
v1.6.16 — estimation tables line up (the xtreg output round)	20
v1.6.17 — macOS build fails with words, not a garbage include path	20
v1.6.17 — macOS build: honest errors for missing dependencies	21
v1.6.18 — macOS: vendored ReadStat (there is no brew formula)	21
v1.6.19 — make deps-readstat can actually run	21
v1.6.20 — macOS: iconv is not free (the last mile of deps-readstat)	21
v1.6.21 — clang-clean (Apple clang round, swept in one pass)	22
v1.6.22 — deep-clean under clang -Weverything	22
v1.6.23 — Bug 35: (long)missing is undefined behavior, and ARM proved it	22
v1.6.24 — Apple-silicon test-suite portability (58/64 -> 64/64 expected)	23
v1.6.25 — sysuse datasets carry variable labels	23
v1.6.26 — native Excel import: the browser reads xlsx now	23
v1.6.27 — Bug 37: browser downloads silently cancelled	24
v1.6.28 — Bug 38: multi-file download discarded by the permission prompt . . .	24
v1.6.29 — the workspace download includes the data in memory	24
v1.6.30 — save buttons with live state: memory is the deliverable	24
v1.6.31 — dir/ls resurrected; save buttons never lock you out	25
v1.6.32 — dir output sorted (readdir order is filesystem-dependent)	25
v1.6.33 — the unified VCE module (DESIGN_VCE.md); Bug 39	25
v1.6.34 — hygiene tier: correlate and ttest	26
v1.6.35 — the macro time-series inference tier	26
v1.6.36 — the state-space engine and ucm (tier release one)	26
v1.6.37 — arima re-grounded on exact ML (the announced golden change)	27
v1.6.38 — test 72 hardened against optimizer landing-point spread	27
v1.6.39 — dfactor, the constraints subsystem, and dsyev (tier release three) . . .	27
v1.6.40 — the sspace subset completes the state-space tier; Bug 40	28
v1.6.41 — sspace diffuse, dfactor idio AR, seasonal arima, ucm cycles	28
v1.6.42 — the panel depth tier	29
v1.6.43 — manual refresh and thematic help	29
v1.6.44 — styled help headings, everywhere a human is watching	29
v1.6.45 — build-system hygiene: exact incrementality, and a latent bug	30

Known bugs and v1.0 limitations

v1.0 known limitations (workarounds documented; v1.1 candidates)

These are not bugs in the strict sense — they’re scope decisions for v1.0 that we may revisit based on real-world testing.

.dta round-trip loses xtset state

save mydata.dta then use mydata.dta preserves variable types, formats, variable labels, and value labels, but does **not** preserve the xtset panel/time variable assignment. Stata writes this as a dataset characteristic (`_dta[iis]`, `_dta[tis]`); readstat doesn’t expose characteristic writing in its public API, and hand-writing characteristic blocks in raw .dta format is a v1.1 task.

Workaround: re-run `xtset country year` after use. This is the standard Stata-do-file convention anyway, so most users won’t notice.

ARIMA uses conditional likelihood, not exact ML (no Kalman filter)

tea's `arima` command computes parameters by minimizing the SSR of the recursively-computed residuals, conditioning on $\varepsilon_{\{<0\}}=0$ (the conditional likelihood approach). This gives consistent estimates but slightly different finite-sample behavior than Stata's default exact ML via the Kalman filter:

- Absolute log-likelihood values differ from Stata's (we drop the initial-state contribution)
- AR coefficients near the unit-root boundary may converge less robustly
- No seasonal ARIMA (`sar`, `sma`) in v1.0

Workaround: for serious time-series work, escape to R's `arima()` or Python's `statsmodels.tsa.arima`.

Factor variable composition with TS ops

`i.country#c.L.gdp` (factor interaction with a lagged regressor) is **not** supported directly. The factor module operates on column names, not on materialized TS-op temps.

Workaround: pre-compute the lag manually: `gen L1_gdp = L.gdp`, then use `i.country#c.L1_gdp`.

xtreg, be doesn't weight by T_i

Stata's default `xtreg, be` uses a slight variance correction for unbalanced panels (down-weights groups with few observations). tea uses straight OLS on the panel means. For balanced panels the answers are identical. For unbalanced panels they can differ by a few percent in the SEs.

MA(1) and ARMA SE precision

ARIMA standard errors are computed from a numerical-differentiation Hessian (forward differences scaled by $1e-5$). For models with $q>0$ (MA components), the SEs can be off by 1-5% compared to analytical or exact-ML Hessians. Point estimates are unaffected.

Variable-name length limit (32 chars)

A residual Stata-12 limit. Factor-variable interaction names like `2010.year#1985.country#c.gdp` (28 chars) fit; longer 3-way interactions truncate. In practice not a problem.

Graphics: covered since v1.6.6

`scatter` / `line` / `histogram` (v1.3), plus multi-series `twoway` with `lowess`, `graph box` with `two-level over()`, and the `name()` registry with `graph combine` (v1.6.6). Remaining gaps: `no by()` faceting, `no bar/pie`, `no graph editor`.

CHANGELOG (recently fixed)

Recently fixed

Tier 4 — econometric completion (NEW for v1.0)

ivregress 2sls (in `src/regress.c` ~300 lines):

Syntax: `ivregress 2sls y [exog] (endog = instruments) [, vce(robust|cluster v)]`. Full first-stage regression of each endogenous regressor on the combined instrument matrix $W = [\text{exog}, Z, _cons]$ via `dgels`, second-stage $\beta = (\hat{X}'\hat{X})^{-1}\hat{X}'y$; residuals computed using ORIGINAL X (not fitted) for SE; sandwich $V = A \cdot B \cdot A$ for robust/cluster ($A = (\hat{X}'\hat{X})^{-1}$, B uses original- X residuals). HC1 / CR1 finite-sample adjustments. First-stage F-stat reported (minimum across endogenous regressors — the weak-instrument worry case).

Verified on textbook endogeneity case: with $\text{cov}(x, u) \neq 0$, OLS biased to 2.49; IV recovers 1.997 against true 2.0, first-stage $F=1364$.

poisson — count outcomes via the existing MLE driver:

Added `mle_family_poisson` to `src/mle.c` with `poisson_per_obs`: $\mu = \exp(\eta)$ clipped at $\eta < 700$; `score = y - \mu`; `weight = \mu`; `loglik = y\eta - \mu` (drops the $\log(y!)$ constant — absolute `loglik` differs from Stata but β identical). Added `count_validate` (rejects $y < 0$). Generalized `do_glm_binary` header text to “Poisson regression”. Verified intercept-only: $\beta_cons = \log(\bar{y}) = 0.6931$ for $\bar{y} = 2$ (exact).

xtreg, be (between estimator):

Replaced the “not implemented” stub. After group-mean computation (already done in the FE/RE path), branch into BE: build $G_obs \times (K_orig + 1)$ matrix Xb (panel means + appended 1s column for cons), OLS via `dgels`, $\sigma^2_be = \text{RSS} / (G - K)$, $V = \sigma^2 \cdot (\hat{X}'\hat{X})^{-1}$ via Cholesky. Verified hand-checkable case: panel data with $\bar{y}_i = 5 + 2\bar{x}_i$ recovers $\beta_x = 2$, $_cons = 5$, $R^2_b = 1.0000$ exactly.

arima (new module `src/arima.c`, ~580 lines):

Conditional-likelihood ARIMA(p,d,q) with optional ARIMAX exog regressors. Syntax: `arima y [exog], arima(p d q) [noconstant]`. Algorithm: - Difference y in place d times - Pack params as $[\mu, \beta_K_ex, \phi_p, \theta_q]$ - Recursive residual computation with $\varepsilon_{\{<0\}} = 0$ conditional assumption - Gauss-Newton iteration with step halving (max 50 iterations) - SSR via finite-difference gradient + Hessian - $V \approx 2\sigma^2 \cdot \text{Hess}^{-1}$ via `dgetrf/dgetri` with ridge

Verified: AR(1) on strong-signal data recovers $\phi = 0.7010$ (true 0.7). AR(1) on US WEO growth data matches OLS lag regression to 3 digits ($\hat{\phi} = 0.154$ both). MA(1) Monte Carlo with `rnormal()` noise: $\theta = 0.429$ (true 0.5) at $N=500$.

runiform() and **rnormal()** (in `eval.c`):

PCG32-based PRNG with Box-Muller normal. `set seed N` for reproducibility. Essential for Monte Carlo, simulation, and bootstrap workflows. Tested against expected distributional moments: mean 0.492, sd 0.291 for u (vs true 0.5, 0.289); mean 0.001, sd 1.015 for z (vs true 0, 1).

Tier 3 — output and reproducibility (NEW)

estimates command suite (new module `src/estcmd.c`):

<code>estimates store NAME</code>	clone <code>last_est</code> , save under name
<code>estimates restore NAME</code>	load <code>NAME</code> → <code>last_est</code> (for <code>predict</code> / <code>margins</code> / <code>test</code>)
<code>estimates dir</code>	list saved with <code>cmd</code> , <code>N</code> , <code>depvar</code>
<code>estimates drop NAME...</code>	remove named entries; <code>`_all`</code> clears
<code>estimates table NAMES</code>	compact side-by-side comparison table (options: <code>se</code> , <code>t</code> , <code>p</code> , <code>star</code> , <code>stats(N r2 rmse F ...)</code>)

Aliased as `est`. Workspace gains a linked-list field `stored_est` (in `dataset.h`); freed by `ws_free`. Built on the existing `est_clone` infrastructure from the `hausman` work.

estout — publication-ready tables (new module `src/estout.c`):

```
estout [names] [using FILE] [, format(latex|markdown|md|plain)
                                se|t|p stars stats(...) nogaps
                                title("...") label(tab:foo)
                                keep(varlist) drop(varlist)
                                nomtitles mtitles("M1" "M2" ...)]
```

Default format is LaTeX (per the IMF / academic-economics convention). LaTeX output produces a tabular block with significance stars in the superscript style ($\text{\textasciitimes}^{\text{\textasciitimes}}$), variable names properly escaped, and a footnote legend. Stata-style `\_cons` for the constant term. Markdown and plain text variants also supported.

Verified end-to-end on a three-model comparison with the standard options: stars apply z-distribution thresholds (10/5/1%), SEs go in parentheses below each coef, stats footer rows handle $N/r^2/r^2_a/rmse/F/df_r/df_m/sigma_u/sigma_e/rho$.

log audit and fix:

The `log` command captures output via a `printf`-macro override in `commands.c`. Problem: `display` lives in `interp.c` and was using bare `printf`, so its output went to `stdout` but not the log file. Fixed by making `g_logfp` extern and adding a manual tee from `do_display`. Round-trip verified: `log using file.log` now captures both command echoes (`. display ...`) AND their output (`_b[x] = 2`).

.dta round-trip audit:

Spot-checked save → use preservation: - Variable types: `%8.2f`, `%td`:
`%` - Variable labels (label variable x "GDP per capita"): `%` - Value labels (label define, label values): `%` - Numeric values: `%` (down to floating-point precision) - String columns: `%`

Not preserved in `.dta`: `xtset` panel/time variable assignment. Stata writes this as a dataset characteristic (`_dta[iis]`, `_dta[tis]`); `readstat` doesn't expose characteristic writing in its public API, and hand-writing the characteristic block in raw `.dta` format is a v1.1 project. Workaround for v1.0: re-run `xtset country year` after use, which is also the standard Stata-do-file convention.

Tier 2 — factor variables: `i.`, `c.`, `#`, `##`, `ib.` (NEW)

The Stata factor-variable grammar, now supported across every estimator that goes through `tsop_expand_varlist` (`regress`, `xtreg fe/re`, `logit`, `probit`, `predict`, `margins`). Lives in new module `src/factor.[ch]`.

Grammar:

<code>i.var</code>	dummies for each non-base level (base = smallest by default)
<code>ib<n>.var</code>	same, but base = level n
<code>c.var</code>	explicit continuous (display prefix; mostly for use in <code>#</code>)
<code>A#B</code>	interaction: cross product of A's and B's columns
<code>A##B</code>	main effects + interaction: equivalent to "A B A#B"

Examples (assume `country` has levels {1,2,3}, `year` has {2010,2011,2012}):

<code>i.country</code>	→ 2 cols: "2.country" "3.country"
<code>ib2.country</code>	→ 2 cols: "1.country" "3.country"

```

i.country#c.gdp      → 2 cols: "2.country#c.gdp" "3.country#c.gdp"
i.country#i.year      → 4 cols: "2.country#2011.year" "2.country#2012.year"
                        "3.country#2011.year" "3.country#2012.year"
i.country##c.gdp      → 5 cols total: 2.country, 3.country, gdp,
                        2.country#c.gdp, 3.country#c.gdp

```

Architecture:

- `factor.c` is a sibling to `tsop.c`: tokens are parsed into atoms (kind 'T' for i., 'C' for c., 'L' for coef-name "2.x", or bare variable in a # context). Atoms join into terms via #. ## is unrolled to the power-set of single-term products (main effects + all interactions).
- `factor_is_factor_token` is the dispatch test; `factor_expand_token` produces a snapshot-able list of temp frame columns.
- Plumbed into `tsop_expand_varlist`: each token is tried as a TS-op first; if that returns "not mine," tried as a factor-var; if that returns "not mine," falls through to plain `varlist_expand`.

Predict/margins re-materialisation (the subtle bit):

When `predict` sees a coefficient named `2.country` or `2.country#c.gdp`, the column doesn't exist anymore — it was a temp dropped after the estimator finished. Solution: treat coef-name form `<level>.<varname>` as a new atom kind 'L' (level-pinned indicator). When `tsop_expand_varlist` is called with `"2.country"`, it dispatches to `factor`, which produces a single column with values `1{country==2}`. Same mechanism handles `2.country#c.gdp` (indicator × continuous).

Verified:

Spec	Hand-checkable result
<code>regress y x i.country on y = $\alpha_c + x$</code>	$\beta=1$, <code>_cons=α_1</code> , <code>2.country=$\alpha_2-\alpha_1$</code> , ...
<code>regress y x ib2.country</code>	<code>_cons=α_2</code> , <code>1.country=$\alpha_1-\alpha_2$</code>
<code>regress y c.x#c.x on y = x^2</code>	<code>coef = 1.0 exactly</code>
<code>regress y i.country##c.x on y = $\alpha_c + \beta_c \cdot x$</code>	reproduces all α and β exactly
<code>regress growth L.growth i.cid (LSDV) on WEO</code>	matches <code>xtreg fe</code> exactly

`predict` after `logit y x i.g` now works (was the open bug from the previous session). `margins, dydx(*)` after factor-var estimation works.

Restrictions in v1.0:

- `i.var` requires `var` to be numeric. Use `encode` for string ids.
- Composition with TS-ops (e.g., `i.country#c.L.gdp`) is not supported. Manually pre-compute the lag (`gen L1_gdp = L.gdp`) and use `c.L1_gdp`.
- Cell-count cap of 10000 columns per token (huge interactions blocked).
- Variable-name length limit of 32 chars means very long interaction names truncate. In practice not a problem for typical use.

Tier 1 polish — statistical functions, egen extensions, _b/_se, formats

A pass over the everyday-use surface in preparation for v1.0. An audit showed that tea's expression-language function library was already very comprehensive (math, trig, hyperbolic, strings, dates, distributions, glob+regex, value/variable labels, encode/decode/destring/tostring, recode, format). The gaps that mattered:

New statistical functions (in `src/eval.c`): - `normalden(x)` / `normalden(x,σ)` / `normalden(x,μ,σ)` — normal PDF - `lnnormal(x)` — $\log \Phi(x)$, stable for large negative x via the asymptotic expansion (reuses `tea_log_normal_cdf`) - `lnnormalden(x)` — $\log \varphi(x)$ closed-form

Better regex (in `src/eval.c`): - `regexs(n)` — extract n -th submatch from the last `regexm` call. Uses a file-scope `tea_regex_submatch[20]` buffer. Index 0 is the whole match; 1.. N are capture groups. - `regexr(s, pat, rep)` — replace first regex match with a string

Extended egen (in `src/commands.c`): - New aggregators: `median`, `iqr`, `p25`, `p75`, `p50`, `pc` (with option `p(N)` to specify percentile), `rank`, `tag` - Multi-argument row functions now work: `rowtotal(v1 v2 v3)`, `rowmean`, `rowmin`, `rowmax`, `rowsd`, `rowmiss`, `rownonmiss` (the previous implementation was single-arg only) - `tag(g1 g2)` marks the first row of each group as 1, others 0 - `rank(x)` uses average-rank-for-ties

Format application (in `src/commands.c::fmt_cell`): Custom `printf`-style numeric formats set via `format var %w.dF` are now honored by `list` (previously only date formats like `%td` were).

Coefficient/SE macros after estimation (in `src/regress.c` + `src/interp.c`):

`_b[name]` and `_se[name]` are now accessible everywhere `e(N)` is. The estimators (`regress`, `xtreg`, `logit`, `probit`) store macros keyed by `_b[varname]` and `_se[varname]` (including bracketed names like `_b[L.growth]` for TS-op regressors). `macro_expand` recognises the pattern and substitutes outside double quotes; inside quotes the text stays literal. Standard Stata behaviour.

This unlocks idioms like:

```
quietly regress y x
display "t-stat = " _b[x] / _se[x]
gen yhat = _b[_cons] + _b[x]*x
```

Bug fix — quoted-string macro substitution (in `src/interp.c`): `display "e(N) = "` `e(N)` used to produce `"3 = 3"` because the macro expander substituted both occurrences of `e(N)`, including the one inside the quoted string. Fixed by tracking `in_dquote` state and skipping `e()` / `r()` / `_b[]` / `_se[]` substitution inside strings. Stata's behaviour preserved: backtick-locals and `$globals` are still substituted inside strings, since those are Stata's quote-aware substitution syntax.

margins — marginal effects after regress / logit / probit (NEW)

<code>margins, dydx(*)</code>	AME for all continuous regressors
<code>margins, dydx(x1 x2)</code>	AME for specified regressors
<code>margins, dydx(*) atmeans</code>	MEM (effects at sample means)
<code>margins, dydx(x1)</code>	AME for <code>x1</code> only

After OLS, the AME of x_k is exactly β_k and the SE is exactly $SE(\beta_k)$ — the margins output reproduces the coefficient table. Useful as a sanity check.

After logit / probit, AMEs are the economically meaningful number: how the predicted *probability* changes with x , averaged across the sample. Coefficients themselves aren't directly interpretable (they're log-odds for logit, latent-index units for probit).

Math:

Logit: $m_k = (1/N) \sum_i \Lambda(X_i\beta) \cdot (1 - \Lambda(X_i\beta)) \cdot \beta_k$ Probit: $m_k = (1/N) \sum_i \varphi(X_i\beta) \cdot \beta_k$ OLS: $m_k = \beta_k$

Delta-method SE: $V(m_k) = G_k \cdot V(\beta) \cdot G_k'$, where G_k is the K -vector $\partial m_k / \partial \beta_j = (1/N) \sum_i [g'(\eta_i) X_{ij} \beta_k + g(\eta_i) 1\{j=k\}]$. Uses the full $V(\beta)$ matrix from the previous estimation (so robust/cluster SEs flow through automatically).

Verified: - Logit and probit on the same data produce nearly identical AMEs (a known result — they only differ at the tails, which contribute little to the average). In one test case $AME(\text{logit}, x1) = 0.12237$, $AME(\text{probit}, x1) = 0.12253$ — different in the fifth decimal. - AME for OLS reproduces the regress coefficient table exactly. - TS-op coefficient names (L.x etc.) work via the same snapshot-then-drop trick predict uses.

Deferred for v2: - margins x1 (predicted outcomes at varying x1 — useful for plots) - margins, at(x1=(0 1) x2=10) (predictions at specified values) - Discrete-change effects for indicator variables - Marginal effects after xtreg (would need within-effect-aware formulas)

predict extended to logit, probit, and xtreg (NEW)

The old predict worked only after regress, with options xb and residuals. Rewritten to dispatch on the previous estimator's cmd:

After regress:

predict yhat	xb (default)
predict r, residuals	y - Xβ
predict s, stdp	standard error of the linear prediction

After logit / probit:

predict p	predicted probability (default)
predict idx, xb	linear index Xβ

After xtreg (fe or re):

predict yhat	xb (default; does NOT include u_i)
predict u, u	$\hat{\alpha}_i$ (panel effect)
predict e, e	idiosyncratic residual y - Xβ - $\hat{\alpha}_i$
predict ue, ue	$\hat{\alpha}_i + e$ (= y - Xβ)
predict xbu, xbu	Xβ + $\hat{\alpha}_i$ (in-sample prediction)

The dispatch lives in predict_resolve_kind and the per-row evaluation is one big switch. Sensible errors when options don't match the preceding command (e.g. predict p, pr after regress fails cleanly).

TS-op coefficient names: when e->xnames[j] is L.growth (TS-op form), the column doesn't exist in the frame anymore at predict time, but we can re-materialize it via the shared tsop_expand_varlist. A subtlety: we have to **snapshot** the materialized data into our own buffer, drop the temps, THEN add the new prediction variable — otherwise tsop_drop_temps would drop the just-added prediction column instead of the temp. This is what an earlier draft got wrong (a regression test would have shown a list-not-found error).

xtreg specifics: $\hat{\alpha}_i$ is computed in one pass over the in-sample rows (e->used tells us which) by streaming through the panel-sorted data and accumulating $\bar{y} - \bar{x}'\beta$ per panel. For predictions on rows that weren't in the estimation sample, $\hat{\alpha}_i$ is unknown so u/e/ue/xbu return missing.

Verified on the hand-checkable FE case (3 panels × 4 obs, perfect fit) where the math is exact: $\beta=1$, $\alpha=10/20/30$, $e=0$, $xbu=y$. Also tested on WEO AR(1) growth where predict reproduces the fit yhat + r = growth exactly.

logit and probit — binary-outcome MLE (NEW)

```
logit y x1 x2 ... [if] [in] [weight], [vce(robust|cluster v)]
probit y x1 x2 ... [if] [in] [weight], [vce(robust|cluster v)]
```

Implemented via a Newton-Raphson MLE driver in `src/mle.[ch]` with family-specific callbacks for the score, weight, and log-likelihood contributions. Same driver should extend to poisson, cloglog, etc. with no changes to the iteration logic — just new `MleFamily` structs.

Output matches Stata’s layout: iteration log, LR χ^2 , pseudo R^2 , coefficient table with z-stats, optional cluster adjustment note.

Standard errors: - (default): classical $(X'WX)^{-1}$ (asymptotic inverse-Hessian) - `vce(robust)`: HC1 sandwich - `vce(cluster v)`: CR1 cluster-robust on score

Numerical care: - $\log(1+e^\eta)$ computed stably as $\max(\eta, 0) + \log(1+e^{-|\eta|})$ — no overflow. - Probit’s ϕ/Φ uses asymptotic Mills bound $\lambda(\eta) \rightarrow -\eta$ when $\Phi(\eta)$ underflows. Per-obs Fisher weight clipped at $1e8$. - Starting values: OLS solve, which also detects rank deficiency (collinear regressors marked omitted, held at 0 through iteration). - Step halving (up to 8 halvings) when full Newton step decreases $\square$. - Perfect-separation detected when $\square \rightarrow 0$; warning printed.

Verified: - Intercept-only on 50/50 sample: $\beta_{\text{cons}} = 0$ exactly. SE = 0.2 for logit, 0.1253 for probit (matches closed forms). - Real-effect data: classical $\beta_{\text{logit}}/\beta_{\text{probit}} \approx 1.81$ ratio reproduced. - Newton converges in 3-5 iterations for well-conditioned data.

Postestimation: test and lincom work on logit/probit estimates. (test uses F with $df_r = N-K$; the asymptotic test would be χ^2 , but the difference is negligible for typical sample sizes.)

test syntax extended (NEW)

The test command now supports the full Stata-style hypothesis syntax:

```
test var                H0:  $\beta_{\text{var}} = 0$ 
test var = 0            same, explicit
test var = 0.5          H0:  $\beta_{\text{var}} = 0.5$  (numeric RHS)
test var1 = var2        H0:  $\beta_{\text{var1}} = \beta_{\text{var2}}$  (equality)
test var1 var2 ...      joint H0: all listed = 0
```

Internally constructs the restriction matrix R and value vector r , then computes the Wald $F = (R\beta - r)'(RVR')^{-1}(R\beta - r)/q$ against $F(q, df_r)$. Pre-existing test `L.growth` `L2.growth` (joint zero) is a special case and still works.

lincom accepts TS-op coefficient names (NEW)

The lincom parser now treats `.` as part of a coefficient name when followed by a name character, so `lincom L.growth + L2.growth` works correctly. Previously the parser stopped at the first `.` and tried to look up “L” as a coefficient.

xtreg, re — random-effects (GLS) estimator (NEW)

Implements panel RE via feasible GLS with quasi-demeaning:

```
xtreg y x1 x2 ... [if] [in], re [vce(robust|cluster v)]
```

Method: the model is $y_{it} = \alpha_i + x_{it}'\beta + \varepsilon_{it}$ with $\alpha_i \sim (0, \sigma_u^2)$ assumed independent of x_{it} . The estimator transforms each variable by subtracting θ_i times the panel mean, where:

$$\theta_i = 1 - \sigma_e / \sqrt{(T_i \cdot \sigma_u^2 + \sigma_e^2)}$$

Then runs OLS on $(y, X, 1-\theta_i)$ — note an explicit constant column, unlike FE. σ_e^2 comes from the within (FE) regression; σ_u^2 comes from the variance components computation $(\max(0, \text{var}(\hat{\alpha}_i) - \sigma_e^2/\bar{T}))$. The transform interpolates between OLS ($\theta=0$) and FE ($\theta=1$).

Output matches Stata's xtreg, re layout: - “Random-effects GLS regression” header - R-within, R-between, R-overall — all computed with $\beta_{RE} - \sigma_u, \sigma_e, \rho$ — variance components (shared with FE math) - θ — quasi-demeaning factor; min/avg/max for unbalanced panels - $_cons$ coefficient and SE (unlike FE) - Wald χ^2 for the slopes - $\text{corr}(u_i, X) = 0$ noted as the RE assumption

Verified on: - Degenerate hand-checkable case (perfect within fit $\rightarrow \theta=1 \rightarrow$ RE collapses to FE; $\beta_{RE} = \beta_{FE} = 1.0$ exactly; $_cons$ omitted since the $1-\theta$ column is identically zero). - WEO AR(1) growth: $\beta_{RE,L} = 0.245$ sits between pooled OLS and FE $\beta = 0.220$. $\sigma_u/\sigma_e/\rho = 0.97/5.37/0.032$ — most variance within- panel. θ varies 0.19-0.39 across 210 countries.

v1 limitations: - Uses $\sigma_u^2 = \max(0, \sigma_{BE}^2 - \sigma_e^2/\bar{T})$ with $\bar{T} = N/n$ for unbalanced. Stata's exact Swamy-Arora uses a slightly different $\bar{T}$ correction; for balanced panels the estimators agree.

hausman — FE vs RE specification test (NEW)

```
xtreg y x1 x2 ..., fe
xtreg y x1 x2 ..., re
hausman
```

Computes $H = (\beta_{FE} - \beta_{RE})' [V_{FE} - V_{RE}]^{-1} (\beta_{FE} - \beta_{RE})$, distributed $\chi^2(K_{\text{common}})$ under H_0 : $\text{cov}(\alpha_i, x_{it}) = 0$. Rejection means RE is inconsistent; use FE.

Tea saves the last xtreg fe and last xtreg re into dedicated workspace slots automatically, so the user just types hausman with no arguments after running both. Order doesn't matter.

Implementation notes: - Only compares slope coefficients present (and not omitted) in both estimates. $_cons$ is dropped since FE doesn't have one. - The difference matrix $V_{FE} - V_{RE}$ should be PSD under H_0 , but sampling noise (especially with robust/cluster SEs) can make it not strictly PD. Cholesky inverse is tried first; if it fails, falls back to an SVD pseudoinverse. When the resulting statistic comes out negative, we report $|H|$ with a “sign flipped” warning. - Verified on WEO AR(1) growth: $\chi^2(1) = 189$, $p < 0.0001 \rightarrow$ strongly rejects RE (matches the FE-reported $\text{corr}(u_i, Xb) = 0.97$ warning sign). On a constructed case where $\alpha_i \perp x$ by construction, $\chi^2 \approx 0$ and the test correctly fails to reject.

xtreg, fe — fixed-effects (within) estimator (NEW)

Implements panel FE OLS via within transformation:

```
xtreg y x1 x2 ... [if] [in], fe [vce(robust|cluster v)]
```

Reports R-within, R-between, R-overall, $\sigma_u, \sigma_e, \rho, \text{corr}(u_i, X\beta)$, and the F-test that all $u_i = 0$. Verified on a hand-checkable case (tests/regression/29_xtreg_fe.do): 3 panels $\times$ 4 obs, $y = \alpha_i + x$ with $\alpha \in \{10, 20, 30\}$ gives $\beta_{FE} = 1.0$ exactly, R-within = 1.0, $\sigma_u = 10$ exactly. Verified against the WEO AR(1) growth dynamic showing internally consistent results (β from xtreg D.growth L.growth, fe = $-(1 - \beta$ from xtreg growth L.growth, fe)).

Standard error options: | , fe | classical: $\sigma^2 \cdot (X_w'X_w)^{-1}$ with $df = N - n_groups - K$
 | | , fe vce(robust) | HC1 sandwich on within-transformed data | | , fe vce(cluster v)
 | CR1 cluster-robust on within-transformed data |

v1 limitations (deferred to v2): - `_cons` is not displayed. Its point estimate (mean of α_i) is computable but the strictly correct SE requires LSDV-equivalent variance propagation we haven't implemented. Users wanting inference on `_cons` can run `regress y x i.panel` for now. - Stata defaults `vce(robust)` to cluster-by-panel for `xtreg fe`; tea's `vce(robust)` is HC1. Use explicit `vce(cluster panelvar)` for cluster-by-panel. - `xtreg, be` (between effects) is not yet implemented — returns a clean error suggesting , fe or , re. (, re ships in this version — see above.) - Weights work the same way as in `regress` (build `design` handles them), but the FE-specific weight semantics have not been carefully verified yet.

TS-op handling unified across all varlist-accepting commands (NEW)

Time-series operator support is now centralized in `src/tsop.[ch]`. The parser handles every Stata form once (`L.x`, `L2.x`, `L(1/2).x`, `L.(x y)`, `L2.(x y)`, `L(1/2).(x y)`, step form `L(1(2)9).x`), and every command that accepts a varlist routes through `tsop_expand_varlist`. Concretely:

Command	TS ops	Mechanism
<code>regress</code>	☐	<code>tsop_expand_varlist (depvar + xspec)</code>
<code>summarize</code>	☐	<code>tsop_expand_varlist</code>
<code>list</code>	☐	<code>tsop_expand_varlist</code>
<code>tabulate</code>	☐	<code>tsop_expand_varlist</code>
<code>tabstat</code>	☐	<code>tsop_expand_varlist</code>
<code>collapse</code>	☐	<code>tsop_expand_varlist (paren-aware tokenizer)</code>

`regress` lost ~400 lines of duplicated parser code in this refactor (from 1198 down to 796 lines). Future commands that accept varlists — `xtreg`, `ivregress`, `correlate`, etc. — get TS-op support for free by calling `tsop_expand_varlist`. Test 28 exercises the universal coverage; tests 21-27 cover individual command surfaces.

Side fix discovered: `keep if / drop if` used to call `frame_unsort` which over-eagerly cleared `ts_panel / ts_time`. Row deletion doesn't invalidate panel structure (the columns are still there and still sorted; gaps are fine for tea's gap-aware `tsidx` lookup), so the `ts` state now survives those operations.

xtdescribe / xt des (NEW)

Added `xtdescribe` (with `xt des` abbreviation) to summarize the panel structure declared by `xtset`. Shows number of panels, time range, delta, span, min/p25/p50/p75/max distribution of obs per panel, and balanced/unbalanced determination. Test 27 exercises both balanced and unbalanced cases.

TS operators in summarize and list (FIXED)

`summarize L.x` and `summarize D.x` failed with “variable not found” because those commands called `varlist_expand` directly and didn't recognize TS-op tokens. The earlier TS-op fix only touched `regress`.

Fixed by adding a shared `tsop` module (`src/tsop.[ch]`). Both summarize and list now route their varlist through `tsop_expand_varlist`, which materializes TS-op tokens as temporary frame columns (with canonical display names like `L.x`, `D2.y`) and appends them to the frame for the duration of the command, then drops them via `tsop_drop_temps` before returning. Test 26 covers this.

D-operator iterated differencing (FIXED — silent math bug)

`D2.x`, `D3.x`, etc. all returned the same as `D.x` because the eval code hardcoded `back = -1` for any D-kind operator. This was a **silent** numerical bug: regressions including `D2.x` would compute correctly for `D.x` but the higher-order differences were all collinear with `D.x` and got “omitted” — looking like a model specification issue rather than a math error.

Stata’s convention is that `D#.x` is the iterated `#`-th difference:

- $D.x = x[t] - x[t-1]$
- $D2.x = D(D.x) = x[t] - 2 \cdot x[t-1] + x[t-2]$
- $D3.x = x[t] - 3 \cdot x[t-1] + 3 \cdot x[t-2] - x[t-3]$
- $D^k x[t] = \sum_{j=0..k} (-1)^j \cdot C(k,j) \cdot x[t-j]$

Fixed in `src/eval.c` using a binomial-recurrence loop. The `S` operator (separate semantics: simple gap-aware seasonal difference, $S#.x = x[t] - x[t-#]$) is unchanged. Test 25 verifies on $x = t^2$ where `D.x` is linear, `D2.x` is constant 2, `D3.x` is constant 0.

Action item for users: any analysis using `D2.x`, `D3.x`, or higher in past sessions silently produced wrong values. Re-run any such code.

Operator-list TS syntax (FIXED)

Stata’s compact form `regress y L(1/2).x` (expanding to `L1.x L2.x`) was rejected with “independent variables not found”. Tea now handles the full set of Stata TS-op forms in `regress` varlists:

Syntax	Expands to
<code>L.x</code>	<code>L.x</code>
<code>L2.x</code>	<code>L2.x</code>
<code>L0.x</code>	<code>x</code> (no shift)
<code>L(1/3).x</code>	<code>L.x L2.x L3.x</code> (range)
<code>L(1 3).x</code>	<code>L.x L3.x</code> (explicit list)
<code>L(1(2)9).x</code>	<code>L.x L3.x L5.x L7.x L9.x</code> (step form)
<code>L.(x y)</code>	<code>L.x L.y</code>
<code>L2.(x y)</code>	<code>L2.x L2.y</code>
<code>L(1/2).(x y)</code>	<code>L.x L.y L2.x L2.y</code> (cross product)

Implemented in `src/regress.c` via the unified `try_tsop_form` parser plus `parse_numlist` (range / explicit list / step) and `find_matching_paren` for nested parens. The regressor tokenizer is paren-aware so spaces inside (...) don’t break tokenization. Tests 21, 23, and 24 cover the forms.

Multi-digit time-series operators read wrong lag (FIXED)

`L2.x` was actually computing the 12th lag instead of the 2nd; `L3.x` the 13th, `F2.x` the 12th lead, `S3.x` a 13-period seasonal difference. Affected expressions, `regress`, and any other consumer.

The bug was in `try_tstop` in `src/parse.c`: the lag-magnitude accumulator was initialized to 1 instead of 0, so reading “L2” gave $\text{num} = 1 \times 10 + 2 = 12$. Lag-1 (`L.x / L1.x / no digits`) accidentally worked because `num` stayed at 1 either way.

Fixed by initializing `num = 0`. The existing `if (!have) num = 1`; fallback still gives bare `L`. the right value of 1. Test 22 pins down correct multi-digit operator behavior.

Time-series operators rejected in regress varlist (FIXED)

`regress growth L.growth if country_id == "USA"` failed with “regress: independent variables not found” because the `regress` code called `varlist_expand` directly on each regressor token; that helper only knew about plain names, wildcards, ranges, and `_all` — not the `L.x / F.x / D.x / S.x` syntax.

Fixed by refactoring `build_design` in `src/regress.c` to detect TS-op tokens and route them through the existing expression evaluator (`expr_parse` → `N_TSOP`). Plain names still go through `varlist_expand`. Works for both the dependent variable and the regressors (`regress D.growth L.growth` is supported). Test 21 covers `regress` with TS operators.

capture and _rc (FIXED in v1.4.0)

Fixed: `capture` now suppresses all output of the captured command (stdout and stderr, at the file-descriptor level) and `_rc` returns the captured return code — 0 on success, the error code otherwise — exactly as in Stata. One small remaining deviation: a successful command that is NOT under capture leaves `_rc` unchanged rather than resetting it to 0; branch on `_rc` only immediately after a capture.

test 08 captures fewer lines than expected

The reshape-error regression test passes, but if you read the expected output you’ll notice some error messages print to stderr while others to stdout, and the line ordering across the merged stream depends on line-buffering. This is OK as of v0.5.20 because we set the stdout buffer to line-buffered in `main.c`, but the harness assumes the ordering is stable — if a future change reorders error/output emission between commands, the test may need updating.

Storage types are display-only

`gen byte x = 1` accepts the byte qualifier but stores the value as double internally; `gen strl0 country = "USA"` accepts `strl0` but doesn’t track width. This is a deliberate design decision documented in `COMPATIBILITY.md` — tea is double + string internally regardless of declared type, and the `.dta` writer auto-compresses on save. Not a bug.

Numeric function family coverage

Many Stata numeric functions are not yet implemented:

- Distribution functions: `normal`, `normalden`, `invnormal`, `chi2`, `chi2tail`, `invchi2`, `t`, `ttail`, `invttail`, `F`, `Ftail`, `binomial`, `binomialp`, `poisson`, `poissonp`.
- Matrix functions (we don’t have a matrix type yet).
- Random-number generators: `runiform`, `rnormal`, etc. (Tied to `set seed`.)

The existing function set covers basic arithmetic, transcendentals, date/time, string manipulation, and the common aggregator functions used in `egen`.

v0.6 .dta polish (this milestone)

- **Value labels** — label define / label values sets are now written into .dta and read back via use. Round-trips cleanly (test 16).
- **Dataset label** — label data "text" now persists. Visible in describe, round-trips through save/use (test 17).
- **save FILE.dta, version(NUM)** — emit any DTA format 104-119 (test 18). Defaults to 118 (Stata 14).
- **Workspace label-set cleanup** — clear and use, clear now also drop the workspace's value-label sets, matching Stata semantics.

Still deferred from v0.6

- **Notes** — Stata's note: text on a dataset or variable. Not yet read or written.
- **strL strings > 2045 bytes** — currently strings longer than the str# limit are truncated at 2045 on write. Real-world risk is essentially zero for econ data; will add when needed.
- **.dta fweight metadata writes are silently dropped** — readstat 1.1.9 accepts readstat_writer_set_fweight_variable but does not serialize the resulting Stata _fweight characteristic. Tea's read-side captures the metadata correctly when present (so real Stata-saved files round-trip through tea correctly *if* the user doesn't save through tea), and the write-side wiring is in place for when readstat fixes this. Track upstream.

reported by user during real-data testing

- Bug 1 — _N in display context returned 1 → **FIXED in v0.5.20**
- Bug 2 — tabstat columns() ignored in by-mode → **FIXED in v0.5.19**
- Bug 3 — tabstat with if + by showed empty by-groups → **FIXED in v0.5.18**
- Bug 4 — list * segfaulted on 147-var dataset → **FIXED in v0.5.14**
- Bug 5 — quoted filename in import not opened → **FIXED in v0.5.12**

reported by Claude during synthetic testing

- Bug 6 — commas in string literals broke parser → **FIXED in v0.5.20**
- Bug 7 — \$1 in strings expanded as undefined macro → **FIXED in v0.5.20**
- Bug 8 — egen silently skipped strings → **FIXED in v0.5.20**

reported by Mico during WB fertility replication (v1.6.x)

- Bug 9 — import excel, firstrow header naming didn't match Stata: invalid chars were underscored (Country Code → Country_Code) instead of removed (CountryCode), and digit-leading headers (1960) got a v prefix instead of Stata's Excel column-letter fallback (F...BQ). Do-files written against Stata failed with variable not found. → **FIXED**: excel/ods import now uses Stata's exact rule; duplicate headers fall back to the column letter.
- Bug 10 — foreach v of varlist BAD was a **silent no-op**: an unknown variable made varlist_expand return -1, which the loop treated as an empty list. The body ran zero times with no error, and scripts failed much later with a misleading message. → **FIXED**: hard error, rc=111. The same -1-as-empty flaw was guarded at every remaining silent site: by BAD: (ran ungrouped!), egen ..., by(BAD) (ungrouped),

`collapse, by(BAD)` (would collapse the whole dataset to one row), `format % BAD` (silent no-op). All now `rc=111`.

- Bug 11 — `local ++yr` parsed `++yr` as the macro *name*, silently defining a macro literally called `++yr` while `yr` never changed — loop counters froze and `rename` loops died with “already exists” on the second iteration. → **FIXED**: `local ++x / local --x` now increment/decrement per Stata; `++` on an undefined or non-numeric macro errors (`rc=198`) instead of silently minting a new macro.

reported by Mico during WB fertility replication, round 2 (v1.6.2)

- Bug 12 — `reshape long` silently **DROPPED every non-stub variable**: after `reshape long y, i(CountryCode IndicatorCode) j(year)` the carried columns (`IndicatorName`, `CountryName`, ...) simply vanished, and the do-file died six lines later with a misleading `keep: variable not found`. → **FIXED**: carried variables are replicated across the `j`-rows of each wide observation, exactly as Stata does. Formats and variable labels survive on `i()` and carried columns.
- Bug 13 — `reshape` had a fixed 512-level cap that **silently dropped** level 513 onward (real WB pulls have ~1,900 indicator codes). → **FIXED**: dynamic sorted level arrays, binary-search everywhere.
- Bug 14 — `reshape wide read j` as numeric unconditionally, so `j(code_safe)` string was impossible. → **FIXED**: string `j` in both directions. New loud errors, never silent mangling: string `j` without the `string` option (`rc 109`), generated names that aren't valid identifiers (`rc 198`, with a `strtoname()` hint), missing/empty `j` values (`rc 498`), carried variables not constant within `i()` (`rc 9`), duplicate `j` within an `i()` group (`rc 9` — previously last-write-wins, silently).
- Bug 15 — `merge` dropped variable labels (formats were kept). → **FIXED**: labels survive `merge` on master and using columns.
- Bug 16 — `frame_set_nobs` **leaked every string cell on shrink**: any row-deleting operation (`drop if`, `keep if`, `duplicates drop`) leaked the dropped rows' strings, and a later `regrow` overwrote the stale pointers, making it permanent. Invisible until `reshape` began carrying string columns into frames that then get row-pruned; caught by ASan. → **FIXED** in `dataset.c`.
- NEW — `preserve / restore` implemented (was on the roadmap): single depth like Stata, snapshot on disk as a native `.tea` tempfile so large frames don't double peak memory. `restore`, not `discards`; `restore`, `preserve` reloads but keeps the snapshot. A `preserve` still pending when a do-file concludes — normally or via `abort` — is restored automatically, with a note.
- NEW — `strtoname()` string function, Stata-exact: invalid characters become `_`, a leading digit gets a `_` prefix, 32-char cap.

v1.6.3 — Stata-parity import naming + WASM refresh

- `import excel WITHOUT` `firstrow` now matches Stata exactly: columns are named by their Excel letters (A, B, ...) and row 1 is kept as a DATA row. (Previously tea always consumed row 1 as headers, silently losing the first data row.) Mixed columns (text header over numbers) become string columns, as in Stata.
- `import delimited` now applies Stata's naming rule: names are LOWERCASED by default and invalid characters are REMOVED (“Country Code” -> `countrycode`, not `Country_Code`); a header that is empty or starts with a digit falls back to the position name `v#`; duplicate headers also resolve to `v#`. New option `case(preserve|lower|upper)` for the Stata-compatible opt-out. NOTE: this is a deliberate behavior change — do-files that relied on tea's old underscored, case-preserving CSV names need

case(preserve) or updated names. The bundled sysuse datasets are unaffected (already lowercase-clean).

- web/tea.wasm rebuilt — the browser engine is current again (it had been frozen at the v1.6.0 feature set while native moved ahead).
- Bug 17 — tempfile names were deterministic (/tmp/tea_tmpN_name), so any do-file using the tempfile + file open, write idiom **failed on its own second run** with rc=602 (“file already exists”) — and leftover files accumulated in /tmp. → **FIXED**: paths now embed the pid (unique per session, like Stata’s), and every tempfile handed out is deleted when tea exits.

v1.6.4 — quality-of-life extensions (not Stata, non-conflicting)

- NEW — status: one-line dataset summary — source filename, obs, vars, exact in-memory size, sort and xtset state. Source is tracked through use/import/sysuse, updated by save, cleared by clear (new Frame.source field, set at the command layer so internal frame operations like the reshape swap or restore never disturb it).
- NEW — progress indicator on long operations (import, reshape long, reshape wide, merge). Time-gated: nothing is drawn until ~1s of wall time, redraws at most every 250ms, stderr-and-TTY-only, erased on completion — so logs, pipes, the regression harness, capture, and short operations see nothing, ever. set progress off to disable. Compiled to no-ops in the WASM build (no TTY; a JS hook can come later).
- Bug 17 refinement: the tempfile uniqueness token is now pid ^ nanosecond-startup-time, because pid alone fails under emscripten (constant pid, and the node harness mounts the HOST /tmp) — caught by the WASM rig colliding with its own previous run.

reported by Mico during WEO test_04 (v1.6.5)

- Bug 18 — **if COND { evaluated against an empty scratch frame with _N hardwired to 1** — so if _N > 0 { was ALWAYS true and if _N == 0 { ALWAYS false, independent of the data. The diagnostic guard blocks in the WEO script ran unconditionally. Silently wrong control flow — the worst class. → **FIXED**: conditions evaluate against the active frame (_n=1, _N = real obs count); variable references like x[1] == 5 now work in if-conditions too.
- Bug 19 — forvalues i = 1/_N ran the body **ZERO times silently**: the range was parsed before macro expansion, sscanf’ failed on the backtick, and 1/0 looped zero times with no error. → **FIXED**: the range is macro-expanded first, and an unparseable range is a loud rc=198 (a legitimately empty range like 1/0 still runs zero times, as in Stata).
- Bug 20 — compound quotes "...“ were unsupported in macro expansion: the expander treated “ as a macro reference and swallowed the text. → **FIXED** per Stata’s rules: “ opens, “ closes, nesting tracked, macros still expand inside, a lone “ inside is literal. display, local x "...“, and label variable accept them; plain “...” strings in display and label variable also accept the “” doubled-quote escape.
- Bug 21 — label variable stripped ALL trailing quotes from the label, mangling labels that contain quotes. → **FIXED** with a real quoted- string parser (see Bug 20).
- NEW — quietly { ... } / capture { ... } / noisily { ... } block forms (prefixes chain: capture quietly { ... }). quietly now also suppresses display, as in Stata (prefix and block forms both).
- NEW — extended macro function local x : subinstr local|global y “from” “to” [, all], with compound-quoted arguments (the way to say a literal double quote). Other extended functions error loudly with the supported list rather than

silently assigning the text.

- NEW — `error #` (abort with a return code; the do-file assertion idiom `if (bad) { ... error 459 }`) and `compress` (accepted for do-file compatibility; tea storage is already minimal — reports “(0 bytes saved)”).
- NEW — `display as error|as text|as result|as input` and old-style `display in red|green|...` are recognized as style directives and no longer printed as literal text.
- Auto-restore of a pending `preserve` now also fires for do-files run from the command line (tea `script.do`), not just via the `do` command; `preserve` snapshots joined the `atexit` cleanup net.

v1.6.6 — graphics round two (the fertility-paper figures)

- NEW — **multi-series twoway**: parenthesized series each with their own `if` and options; plot types `scatter / line / connected / lowess`. Per-series `lcolor()` `lpattern()` `msymbol(i)` `mlabel()` `mlabcolor()` `mlabposition()` (`clock`); globals `ylines(#,...)` (repeatable), `legend(off)` (simple legend otherwise), `yscale(range())`, `ylabel(a(s)b)` / `xlabel(a(s)b)`, `note()`. Axis titles inside a series merge upward, as in Stata. Tick rules and plot range are separate: a `ylabel()` rule never clips data.
- NEW — **lowess**: tricube kernel, running-line (default) or mean, `bwidth()` (0.8), `adjust`. Pure C, byte-identical across rigs.
- NEW — **graph box**: Stata percentile interpolation, adjacent-value whiskers (1.5 IQR), outside values (`noout` hides), one- or two-level `over()` (first varies fastest inside bands of the second), value labels honored, `relabel()`, `label(angle() lsize())`.
- NEW — **named-graph registry + graph combine**: `name(NAME[,replace])` stores the SVG in-session AND writes `NAME.svg` (documented deviation); `graph combine N1 N2, cols()/rows()` nests panels as scaled SVG; `graph dir / graph drop NAME|_all`. Name collision without `replace` is rc 110.
- Everything renders on the same deterministic SVG engine: `%2f` coordinates with -0 snapped, golden-file diffed byte-identically on native, ASan, and WASM rigs (tests 56, 57). `plot.c` and the test-40 goldens are untouched.
- Documented deviations (COMPATIBILITY.md): `name()` writes a file; `twoway line/lowess` series are sorted by `x`; unknown cosmetic graphics suboptions are accepted-and-ignored (strict everywhere else).

v1.6.7 — value-label attachments survive frame rebuilds (Mico’s fig1)

- Bug 22 — **encoded group variables silently reverted to raw numerics** in `graph box` band labels: `reshape` (all four carried-variable paths), `merge` (master and using), and `frame copy` copied `format+vlabel` but dropped `Variable.vallab`. All eight copy sites fixed.
- Bug 23 — the native format stored only the per-variable attachment `NAME`, not the label-set contents, so `save→clear→use` via `.tea` (and the `preserve` snapshot format) lost definitions. The format is now **TEA2**: value-label sets referenced by attached variables travel with the dataset, exactly like Stata’s `.dta`. Legacy `TEA1` files still read (without attachments); old tea versions cannot read `TEA2`.
- Bug 24 — `merge ... using FILE` defaulted a missing extension to `.tea` while `use/save` default to `.dta` — so `savef'+merge ... using f'` on an extensionless tempfile failed with “cannot readtea”. `Merge` now defaults to `.dta`, matching `use/save` and Stata.
- Test 58 locks all five survival paths: `reshape wide+long`, `preserve/restore`, `.tea round-trip`, `.dta round-trip`, and `merge` — asserted via rendered `graph box` band labels.

v1.6.8 — quality of life: conversion-phase progress for import excel

- `import excel` on a large workbook was silent for most of its wall time: the `ssconvert/libreoffice` child ran under a blocking `system()` with no feedback, and only the subsequent CSV-parse phase was byte-instrumented — so the progress line appeared “around 70%”. The converter now runs under a polled `fork/exec` with a new progress ACTIVITY mode: `spinner + elapsed seconds` (“converting WB_data.xlsx / 34s”), same contract as the percentage line (`stderr` TTY only, 1s activation gate, 250ms redraws, erased on completion, `set progress off` disables, no-op on WASM). Batch output is byte-identical.

v1.6.9 — period-date constructors + eval errors abort (OECD HPI round)

- NEW — string period-date constructors: `quarterly(s, "YQ")`, `monthly(s, "YM")`, `halfyearly(s, "YH")`, `weekly(s, "YW")`, `yearly(s, "Y")`, and `daily(s, mask)` as the Stata alias of `date()`. Integer tokens are extracted left-to-right and assigned by the mask; two-digit years get +2000 (matching `date()`); an out-of-range period (“2020-Q7”) yields missing, never a wrong date. The rest of the date family (`mdy/dofq/yofd/qofd/...` and `%tq/%td/%tm` display formats) already existed.
- Bug 25 — **a runtime evaluation error in gen/replace did not abort**: the error printed once per row, missing was stored anyway, each store counted as a “real change”, and the command returned 0. A do-file with `gen qdate = quarterly(...)` under a build lacking the function marched on: collapse grouped on an all-missing year, merge matched ZERO rows, and the garbage panel was saved without a nonzero `rc` anywhere. Now: an eval error at the type probe fails before the variable exists; an error mid-loop aborts with `rc=133` and rolls back completely — `gen` leaves no trace of the new variable, `replace` restores the column from a snapshot. Partial writes never survive.

v1.6.10 — set more on|off does something now

- `set more` had been accepted-and-ignored since v1.0 — a silent no-op: `set more on` armed nothing, and a stray list on a 400k-row panel flooded the terminal. Now it is Stata’s output pager for real: with `set more on`, long list and describe output pauses at a screenful with `--more--`; space/any key = next page, Enter = one more line, q = stop with `--Break--`. Terminal height is read live (`TIOCGWINSZ`, fallback 24 rows). It engages ONLY in the interactive REPL with `stdout` a TTY — do-files, pipes, capture, logs, and the test suite never see it, so batch output is byte-identical. Default is off, matching modern Stata. `set more banana errors (198)`, as does any other malformed argument — same strictness as `set progress`.

v1.6.11 — cellrange() + case() for import excel (WPP fertility round)

- NEW — `import excel ..., cellrange(A17:AF22000)`: restricts the import to a sheet rectangle. Real workbooks (the UN WPP file) carry title junk above the table; without `cellrange`, junk row 1 became the header. With `firstrow`, the range’s first row is the header row. The end corner is optional (`cellrange(A17)` = from A17 to the sheet’s end). The slice happens on the converted CSV with a fully CSV-aware state machine — quoted fields containing commas, doubled quotes, or embedded newlines slice correctly.

- NEW — `import delimited ..., rowrange(r1[:r2]) colrange(c1[:c2])`: Stata's rectangle restriction for text files, sharing the same slicer (and giving the regression suite a WASM-testable path — the `xlsx` conversion needs `ssconvert`, absent on that rig).
- Bug 26 — `case("lower")` with quotes, legal Stata syntax, `ERRORED` on `import delimited` and was **silently ignored** on `import excel`, which had no `case()` handling at all (names hardcoded to preserve). Both branches now accept quoted or bare arguments; the excel naming rule (invalid chars removed, column-letter fallback) folds case AFTER cleaning, and the excel default remains preserve, matching Stata.
- Test 60 locks range slicing (bounded + start-only), quoted-field survival, quoted `case()` args, and loud errors on malformed ranges.
- WASM-rig lesson #3 (joins Bug 17 and the test-57 note): emscripten `NODEFS` misbehaves after `rename()` — subsequent opens in the same directory can fail on stale cache state. The slicer therefore never renames; it writes a `.rng` sibling that callers load and unlink. Slice temp names are per-call unique (`pid ^ nanotime`), not per-pid.

v1.6.12 — `tabstat`: value-labeled groups + `format()` (test_01_b round)

- Bug 27 — `tabstat ..., by(ctr_group)` named groups by raw numeric value (1/2/3) instead of the attached value labels (Advanced/Emerging/LIC). Same lookup rule as `list` and `graph box` now: labels when attached, `%g` fallback otherwise; string by-variables unchanged.
- Bug 28 — `format(%5.2f)` was ignored entirely — `tabstat` had no `format()` parsing and every cell went through a hardcoded `%.7g`. Now: the option applies to every cell INCLUDING the Total rows (the Total block carries a second copy of the cell-format macro, which the first fix pass missed — caught because the verification run showed a formatted table with an unformatted Total). Quoted `format("%5.2f")` accepted. Only printf-safe numeric formats pass validation (`%w.d` followed by `f/g/e/F/G/E`); anything else — `format(%s)`, `format(banana)` — is a loud 198, since an arbitrary string reaching `snprintf` as a format is undefined behavior.
- Test 61 locks labeled groups, formatted group and Total rows in both table orientations, the quoted form, and both rejection paths.

v1.6.13 — long variable names no longer wreck `sum/describe/tabstat`

- Bug 29 — the WPP panel's 26-character names (`totalpopulationasof1januar, ...`) overflowed the fixed 12-char name column in `summarize` and the name columns/headers in `tabstat`, and shoved `describe`'s type/format columns out of alignment. Fixes: `summarize` and `tabstat` use Stata's `abbrev()` display rule (first n-2 characters + `'~'` + the last character: `totalpopul~r`), applied to variable names, column headers, and by-group labels alike; `describe` pads its name column dynamically to the longest name in view (16..32), header row included. Frames with short names render byte-identically to before — the whole existing suite passes against unchanged goldens.
- Test 62 locks the abbreviation rule in `sum`, `tabstat` headers and by-labels (including long VALUE labels), and `describe`'s dynamic alignment.

v1.6.14 — documentation catch-up

- The manual and the generated command reference had lagged three releases of features: the `import` chapter now documents `cellrange()` (with the header-at-row-17 semantics), `rowrange()/colrange()`, quoted-and-bare `case()`; the `date-function` table gains the period constructors (`quarterly/monthly/halfyearly/weekly/yearly` and the daily alias); `import`'s and `tabstat`'s dispatch help strings (the source of the generated

reference) now describe their full current option surface, including `tabstat's format()` and value-labeled by() groups. No behavior changes.

v1.6.15 — command echo for the browser editor (colleague feedback round)

- NEW — set `echo on|off`: echoes each do-file line Stata-style (“.”) before its output. Default OFF — real Stata echoes do by default (run is the silent variant), but tea’s batch-output contract and every golden test predate the feature; documented deviation. Interactive input never echoes (already visible). The browser editor’s Run wraps the script in `echo on/off`, so the terminal now shows each command with its results beneath it.
- Browser editor hardening, from a colleague’s “select-all + run deleted my script, leaving one stray character” report. The run shortcut itself was safe (preventDefault present); the real mechanism is a mistimed chord — Ctrl+A, then d landing after Ctrl is released, which replaces the whole selection with the letter d, and the 400ms autosave then persists the wreckage. Mitigations: the run shortcut now matches D case-insensitively (Shift/CapsLock used to defeat it and leak the key to the browser); every Run first snapshots the buffer to a backup slot; and a “restore last run” link swaps the backup in (reversibly). Native textarea undo (Ctrl+Z) also survives throughout.

v1.6.16 — estimation tables line up (the xtreg output round)

- Bug 30 — coefficient tables misaligned whenever a value needed more than its column: the cell format was %10.6g — six significant digits, MINIMUM width 10, no maximum — so -0.000367782 rendered 12 wide and sheared the row. All estimation tables (regress, xtreg fe/re/be, logit/probit, margins, hausman, arima — 19 sites) now format cells with `gfit()`, an emulation of Stata’s %#.0g: the most significant digits that FIT the column, with the leading zero of $|x| < 1$ dropped (-.0003678, .0000669) — exactly how Stata keeps coefficient columns rigid at any magnitude.
- The header stat blocks (xtreg’s Number of obs / groups / R-squared / Obs per group / F / corr rows, and classic regress’s ANOVA right column) now put every ‘=’ in one fixed column with right-aligned values, Stata’s layout.
- 22 goldens re-blessed; every changed line inspected — coefficient rows, header stats, and ANOVA cells only; no numeric drift (gfit reformats the same doubles).
- Table display precision is capped at SIX significant digits: the WASM rig immediately caught the 7th digit of longley’s ill-conditioned CI bound differing between OpenBLAS and the reference BLAS (-5496.530 vs -5496.529) — the old 6-sig format had been hiding it. Displayed precision now equals the cross-rig reproducible bound; stored doubles keep full precision (e(b)/e(V) and exports are unaffected).

v1.6.17 — macOS build fails with words, not a garbage include path

- On macOS, a missing Homebrew dependency made `brew --prefix` return empty, the Makefile emitted `-I/include`, and the first thing the user saw was a mystifying “readstat.h file not found” deep into the compile. The Darwin block now checks every resolved prefix at parse time and fails immediately with the exact `brew install` command. Also: README’s short macOS line was missing `readline` and `readstat` (the full line elsewhere was correct) — the likely root cause of the report. Linux/WASM builds untouched.

v1.6.17 — macOS build: honest errors for missing dependencies

- Bug 31 — on macOS with readstat not installed, make produced the mystifying readstat.h file not found behind a garbage -I/include flag. The Makefile's guard checked that brew --prefix readstat returned a non-empty string — but brew --prefix prints the would-be opt path even for a KNOWN-BUT-UNINSTALLED formula (exit 0), and comes back empty under other PATH/environment variants, so a non-empty prefix proved nothing. The guards now test for the actual header files (readstat.h, cblas.h, lapacke.h, gsl_cdf.h, readline.h) under each resolved prefix and fail with the exact brew command to run — including a from-source ReadStat recipe in case brew lacks the formula — plus a make READSTAT_PREFIX=/path override. House-keeping targets (clean, distclean, showpaths) are exempt from the checks.

v1.6.18 — macOS: vendored ReadStat (there is no brew formula)

- Correction to v1.6.17's error message and README: **Homebrew has no readstat formula** (verified against homebrew-core) — brew install readstat was advice that cannot work, which is why the header guard fired with an empty prefix on a machine with everything else installed. New make deps-readstat target: clones WizardMac/ReadStat (-depth 1), builds it static into vendor/readstat, one command, ~a minute; the Makefile auto-detects the vendored copy afterwards on every platform and links libreadstat.a statically (one less runtime dependency). vendor/ is excluded from dist tarballs. macOS setup is now: brew install readline openblas lapack gsl && make deps-readstat && make. All 63 tests pass against ReadStat HEAD via the vendored path on the Linux rig.

v1.6.19 — make deps-readstat can actually run

- Bug 32 (chicken-and-egg, one release old): the v1.6.17 header guards fire at Makefile PARSE time, so on a machine without readstat, make deps-readstat — the cure the error message prescribes — died of the same error before its recipe could run. The exemption list (clean/distclean/showpaths) now also covers deps-readstat and check-deps: the cure and the diagnosis must both work precisely when the dependency is missing. Verified from a pristine extracted tarball: deps-readstat → make → full test suite, no system readstat involved.

v1.6.20 — macOS: iconv is not free (the last mile of deps-readstat)

- Bug 33 — on macOS, make deps-readstat built libreadstat.a successfully and then died linking ReadStat's bundled extract_metadata tool: undefined _iconv/_iconv_open/_iconv_close for arm64. Root cause: iconv lives inside glibc on Linux but is a SEPARATE libiconv on macOS, and ReadStat's configure does not add it. Two-part fix: the deps-readstat recipe passes LIBS=-liconv to ReadStat's configure on Darwin, and — the failure one step further that this preempts — tea's own final link now appends -liconv (and -lz for readstat's compression paths) on Darwin, since the static libreadstat.a defers those symbols to the final link. -lz joined the common link line on all platforms (harmless where redundant); -liconv stays Darwin-only, as Linux glibc provides iconv natively and has no libiconv.so to link. Linux revalidated from scratch: vendor rebuild, clean build, 63/63.

v1.6.21 — clang-clean (Apple clang round, swept in one pass)

- Bug 34 — Apple clang’s -Wunused-but-set-variable (an error under -Werror; Linux gcc does not flag these) stopped the macOS build at estcmd.c. Rather than fix one file per report, the whole tree was swept with clang-15 locally; three distinct issues, all fixed properly rather than suppressed: two genuinely dead counters in estimates dir/drop (deleted — Stata is silent there anyway), and the printf/fprintf log-tee macros redefining glibc’s _FORTIFY_SOURCE macros without #undef (clang’s -Wmacro-redefined).
- The Makefile is now compiler-aware: -Wno-format-truncation is GCC-only and upstream clang under -Werror rejects unknown warning options (Apple clang merely tolerates them) — the flag is applied only when CC is gcc.
- The full golden suite passes byte-identically on the clang-built binary — a fifth verification axis alongside gcc-native, ASan, WASM, and (pending Mico’s machine) Apple silicon.

v1.6.22 — deep-clean under clang -Weverything

- Proactive sweep with clang-15 at -Weverything (minus deliberately pedantic families) to over-approximate any Apple clang generation. The tree came back almost clean; three findings, all fixed on merit: the version command embedded **DATE/TIME** — a compile timestamp that made builds non-reproducible byte-for-byte, gone (the version string carries release identity); one genuinely confusing dense one-liner in keep/drop reformatted (correct code, but -Wmisleading-indentation and any human reader agreed it looked wrong); the remaining family (-Wgnu-statement-expression in a macro) is a deliberate GNU-ism outside every compiler’s default set.
- The golden suite passes byte-identically under the clang-built binary at the strictest flag set. If the macOS build still shows errors beyond v1.6.21’s fixes, send `make 2>&1 | grep error: | sort -u` — the last paste arrived empty.

v1.6.23 — Bug 35: (long)missing is undefined behavior, and ARM proved it

- Mico’s Apple-silicon run — the fourth rig — failed test 11: `substr("aaa bbb aaa","aaa","X",.)` returned the ORIGINAL string on macOS while replacing all on Linux. Root cause: the function evaluator cast raw argument doubles to long, and casting a MISSING (NaN-class) double is undefined behavior that diverges by ISA — x86’s `cvtsd2si` yields `LONG_MIN` (negative, accidentally meaning “all”), arm64’s `fcvtzs` yields 0 (limit zero: replace nothing). Same source, opposite results, no diagnostic anywhere. Notably the WASM rig never caught this: LLVM’s `wasm lowering` happened to match x86.
- Every numeric argument coercion in the evaluator (30 sites) now checks missing FIRST and implements Stata’s documented semantics: `substr/substr n=.` means all occurrences; `substr n1=.` gives “”, `substr(s,n,.)` gives the remainder; `word(s,.)`/`char(.)` give “”; the entire date family propagates missing. The outputs are now ISA-independent by construction.
- Test 64 locks the defined semantics for all of it — when it passes on Apple silicon it is the ARM proof, since x86 accidentally produced most of these answers already.

v1.6.24 — Apple-silicon test-suite portability (58/64 -> 64/64 expected)

- Mico’s first full suite run on Apple silicon: 58/64, test 64 GREEN — the ARM proof of Bug 35’s fix. The six failures reduce to two causes, both suite-side, no engine changes:
- Five tests (56/57/58/60/63) printed `cd /tmp`’s output; macOS’s `/tmp` is a symlink to `/private/tmp` and `cd` prints the resolved `getcwd()`. Those tests never meant to assert the path — quietly `cd /tmp` now, goldens re-blessed without the platform-dependent line.
- Test 41: longley’s `_se[year]` differs in the SIXTH significant digit between x86 OpenBLAS and Apple-silicon OpenBLAS (.455479 vs .455478) — the near-singular matrix reproduces across BLAS implementations only to ~5-6 digits, one digit under the general display bound set in v1.6.16. Ruling: the general 6-digit table bound stays (it is right for well-conditioned problems); the TEST now smoke-checks longley with quietly `regress + round()ed _b/_se/e()` assertions at the precision this matrix actually reproduces on every rig. Longley keeps doing its job — probing near-singularity — without holding the byte-identity promise hostage to its own pathology.
- With these, all 64 goldens are expected green on all five platforms: x86 gcc, x86 clang, ASan/UBSan, WASM, and Apple silicon.

v1.6.25 — sysuse datasets carry variable labels

- The bundled datasets loaded with bare column names: they are embedded as raw CSVs, a format that cannot carry labels. Each dataset now embeds a companion variable-label table (`data/NAME.lbl`, plain “varlabel” lines; weo’s 145 indicator labels generated from `data/weo_codes.txt`, truncated at the 80-char vlabel bound) which sysuse applies after the CSV load — describe on grunfeld now reads like the textbook, and the WEO extract documents itself.
- Fixed in passing: sysuse’s missing-braces bug printed the “obs loaded” message even when the load failed (the `if(rc==0)` guarded only the source-string line) — the misleading-indentation class, caught by reading the code it sat next to.
- Test 46 re-blessed: `outreg` has ALWAYS preferred variable labels over names in export tables (publication behavior); it was simply unobservable until bundled data had labels. Test 65 locks labels across all six datasets.

v1.6.26 — native Excel import: the browser reads xlsx now

- NEW — `src/xlsx.c`: a native `.xlsx` reader (zip parsing + zlib inflate + sheet XML), replacing the `ssconvert` shell-out for `.xlsx`. No gnumeric or libreoffice needed anywhere, and — the point — `import excel` works IN THE BROWSER: drag a workbook into the page and import it. Handled: shared strings (incl. rich-text runs), `CACHED FORMULA VALUES` (a cell holding `=‘Sheet’!B2` imports its last evaluated value), inline strings, booleans, XML entities, gap cells, `sheet()` selection by name with a loud 601 for a missing sheet. Cells are emitted as CSV to the same temp path as before, so `firstrow/case()/cellrange()` behave identically. `ssconvert` remains the `.ods` path and an exotic- file fallback. The WASM build links an inflate-only zlib built from madler/zlib source.
- Bug 36 — the `ssconvert`-era cleanup `rmdir`’d the temp CSV’s PARENT directory. Correct when that parent was the per-import `/tmp/tea-xlsx-PID/` dir; with the native path writing straight into `/tmp` it became `rmdir(“/tmp”) — silently a no-op on Linux (non-empty) but SUCCESSFUL on the browser’s MEMFS once our own unlinks emptied it: /tmp deleted, every later temp operation (sysuse!) broken. The rmdir now fires only`

for directories named tea-xlsx*. Found because the demo rehearsal ran import-then-sysuse in one session — exactly the order a real workflow uses.

- Test 66 (with a stored-zip fixture in tests/fixtures/) locks the native reader on all rigs — including WASM, where import excel was previously impossible.

v1.6.27 — Bug 37: browser downloads silently cancelled

- “Download workspace files” counted the files and delivered none, and “Save .do” did nothing, in some browsers. Two independent defects with one visible symptom: the download anchor was never appended to the document (Firefox ignores programmatic clicks on detached anchors), and `URL.revokeObjectURL()` ran synchronously after `click()` — but `click()` only QUEUES the download, so the immediate revoke could destroy the blob before the browser read it. Whether anything landed depended on engine and timing; tea’s own counter was honest (“3 new files downloaded”) because the filesystem side succeeded — the browser then dropped the baton. Both handlers now share one `triggerDownload()`: append to DOM, click, remove, revoke deferred 10 s. Web UI only; the engine is unchanged.

v1.6.28 — Bug 38: multi-file download discarded by the permission prompt

- v1.6.27 fixed the detached-anchor and premature-revoke defects, and single-file downloads (Save .do) work. But “Download workspace files” with SEVERAL new files still delivered nothing: Chromium gates multiple programmatic downloads behind a permission prompt, and the attempts that TRIGGERED the prompt are discarded — “Allow” only covers future clicks. Diagnosed live on Edge the night before an IMF demo.
- Architectural fix: never fire more than one download. One new file downloads directly; several are bundled client-side into a single `tea-workspace.zip` (stored entries, CRC32, plain JS — ~50 lines, validated byte-for-byte against unzip). One file, one download, no permission prompt, every browser. Web UI only; engine unchanged.

v1.6.29 — the workspace download includes the data in memory

- Design gap, called by Mico: “Download workspace files” shipped only files the session had explicitly saved — but the workspace, as a user understands it, INCLUDES the dataset being worked on. `sysuse + modify + download` yielded nothing unless you remembered save.
- New quiet engine export `tea_web_save_memory(path)`: serializes the current frame to .dta (format 118, labels and all) with no terminal output; returns 1 when memory is empty. The download handler calls it first and adds `data_in_memory.dta` to the bundle. Verified by round-trip: modified `grunfeld` (25 obs + generated variable + labels) and the demo’s collapsed end-state both re-open natively, exact.
- The explicit-save workflow is unchanged; this only makes the download button’s promise true.

v1.6.30 — save buttons with live state: memory is the deliverable

- Redesign, specified by Mico after demo rehearsals: the workspace-zip download conflated too much (script-saved intermediates, the editor autosave, charts) and had no notion of “unsaved”. New semantics: “Save data (.dta)” downloads exactly the CURRENT IN-MEMORY dataset, and both save buttons carry live state — lit when there are unsaved changes, grayed otherwise.

- Engine: `tea_web_data_hash()`, an FNV-1a fingerprint of the frame (names, types, labels, all data bytes; 0 = empty). The UI re-hashes after every executed command and compares against the last-saved hash: reads (summarize, list) keep the button gray, writes (gen, drop, label, keep) light it, and a value-identical replace correctly stays gray — the fingerprint knows better than a dirty flag would.
- Editor: input events set a dirty flag; Save .do clears it. Charts: each plot in the Plots tab now carries its own “save NAME.svg” link (single-file downloads, no permission prompts).

v1.6.31 — dir/ls resurrected; save buttons never lock you out

- Sandbox-filesystem visibility (Mico): in the browser you could drop a file in and then not SEE it anywhere. The answer was already in the tree: a complete `do_dir` — Stata’s `dir/ls` with glob patterns — implemented but never registered in the command table. Dead code, now wired: `dir` lists the working directory (sizes deterministic, courtesy of the byte-identical writers), `ls *.dta` filters, and status continues to describe the data in memory. Test 67.
- Save-button policy bug (v1.6.30 was too literal): after saving, the data button DISABLED itself until the data changed, so a re-download — mid-demo, after a lost file, after saving a chart in between — was impossible. New policy: disabled only when there is NO data; the highlight, not the disabled state, carries “unsaved changes”. Save .do likewise: enabled whenever the editor has text, highlighted when dirty.

v1.6.32 — dir output sorted (readdir order is filesystem-dependent)

- Test 67 failed on its second filesystem within the hour: the resurrected `do_dir` printed raw `readdir()` order — alphabetical by luck on one filesystem, ext4 hash order on another. The listing is now collected and `qsorted` before printing. A tidy demonstration of why the goldens exist: the dead code carried a latent nondeterminism, and the suite caught it on the first cross-machine run. Golden unchanged (it was authored on the lucky filesystem).

v1.6.33 — the unified VCE module (DESIGN_VCE.md); Bug 39

- Foundation work for the econometrics roadmap: sandwich VCEs existed as FOUR independent implementations (regress, the GLM family, ivregress, xtreg). New `src/vce.c` owns the abstraction — $V = c*BMB$ from estimator-supplied bread and score rows — plus the option parser, the finite-sample policy table, and the output furniture. `DESIGN_VCE.md` is the contract, including the migration plan and the Stata-verification protocol (`tools/stata_check_vce.do`).
- Bug 39 (semantic): `xtreg, fe robust` computed plain HC1 on the within regression. Stata has promoted robust to `vce(cluster panelvar)` since Stata 10 because HC1 is inconsistent under the within transformation (Stock & Watson 2008) — tea printed a different number than Stata for the same command line. Fixed: robust promotes, the two spellings now agree exactly, and the cluster note prints. Cosmetic contract fix alongside: the SE column header under clustering reads “Robust” (Stata’s label), not “Cluster”.
- `regress` and `xtreg` now route through the module (their shared `robust_V/cluster_V` are thin wrappers; numbers verified unchanged). The GLM family and `ivregress` keep their in-place implementations this release and port next, per the design note’s staging.

- Test 68 locks the estimator x {default, robust, cluster} matrix cross-rig, with xtreg's robust==cluster equality as two identical golden lines.

v1.6.34 — hygiene tier: correlate and ttest

- Two daily-use commands were missing while their siblings existed (pwcrr but no correlate; regress inference but no ttest). Both now implemented per [R] correlate / [R] ttest: correlate = LISTWISE deletion (the defining difference from pwcrr), default _all, honors if/in via the house Sel idiom (which pwcrr, noted for later, does not), means and covariance options, Stata's matrix layout with the "(obs=N)" line. ttest = one-sample against a constant, paired, and two-sample by() with pooled or Welch (unequal) variances, Stata's full table including the three-alternative p-value footer; t probabilities via the GSL t CDF.
- Verification: correlations and all four t statistics (incl. the Welch-Satterthwaite df) reproduced independently to displayed precision; tools/stata_check_vce.do gained a companion section so the Stata run pins layout as well as numbers. Test 69.

v1.6.35 — the macro time-series inference tier

- The largest single addition since graphics (DESIGN_TSINFER.md, new src/tsinfer.c): newey (HAC via the vce module's new Bartlett meat — "a different M from the same s_i", exactly as DESIGN_VCE.md promised), dfuller (constant/trend/drift/noconstant, lags(), ts-operators, regress option), pperron, tsfilter hp/bk/hamilton (frequency-aware defaults; hp by exact banded solve), var (up to 8 variables, 8 lags; per-equation tables, LL and information criteria), vargranger (per-pair and ALL Wald tests), irf create/table (simple + Cholesky-orthogonalized, in memory), lpirf (Jorda local projections, Newey SEs with lag = horizon), and vecrank (Johansen trace with Osterwald-Lenum critical values).
- Verification: ADF t statistics, MacKinnon critical values, HP filter moments, the Newey-West SE, all VAR coefficients, and all Johansen trace statistics reproduce statsmodels EXACTLY at displayed precision on the airline and WEO-USA test problems. Documented convention differences (unit-root p-value approximation, Sigma divisor, CV table editions) recorded in COMPATIBILITY.md; tools/stata_check_ts.do pins the tier against Stata.
- Tests 70 (unit roots + filters) and 71 (var/granger/irf/lpirf/ vecrank) lock the tier cross-rig. Staged out per the design note: vec estimation, fevd, irf files, multivariate lpirf, svar, dfgls, xtunitroot.

v1.6.36 — the state-space engine and ucm (tier release one)

- The Kalman tier arrives per the signed DESIGN_SSPACE.md: src/kalman.c is the engine — DK-form models, UNIVARIATE (sequential) filtering, EXACT diffuse initialization carried as (Pstar, Pinf) until rank exhaustion, the DK backward smoother with the (r0, r1) diffuse extension, Lyapunov initialization for stationary blocks — and src/sspace.c is the first front-end: ucm with ntrend/llevel/ltrend/rwalk/rwdrift, dummy seasonal(#), OIM standard errors on the variance scale by delta method, sm-state(NEW) for the smoothed level, and ML by BFGS2 with central-difference gradients from three deterministic starting points (flat UC likelihoods have boundary-swapped local optima; airline finds the global one on the third start).
- Verification: Nile local level reproduces THE canonical numbers of the literature — var(e)=15098.5, var(level)=1469.18 against Durbin-Koopman's 15099/1469.1 — and matches statsmodels' exact-diffuse log-likelihood to 5 decimals (-633.46456),

including on a sample with a missing decade (-567.14230 vs -567.14231), where the smoothed level inside the hole agrees to display precision. Airline lltrend+seasonal(12) matches statsmodels' optimum (ll 217.42013 vs 217.4203) with the same variance decomposition. sysuse nile ships embedded so the literature's example is a regression test.

- Next per the design note: exact-ML arima re-grounded on the engine (a documented golden change), then dfactor, then the sspace subset.

v1.6.37 — arima re-grounded on exact ML (the announced golden change)

- arima now evaluates the EXACT Gaussian likelihood through the state-space engine (Harvey companion form, Lyapunov initialization, univariate Kalman filter) instead of the conditional-SSR Gauss-Newton of v1.0, per DESIGN_SSPACE.md §8.2. Stationarity and invertibility are enforced by Monahan's partial-autocorrelation transform; optimization is BFGS2 with central-difference gradients from three deterministic starts; standard errors are honest OIM.
- This is the documented Bug-39-class numbers change announced in the design note. In the end no golden actually moved: the only arima regression assertion was already tolerance-based. The stale "conditional likelihood" caveats were removed from the manual, and COMPATIBILITY.md gained the OIM-vs-OPG paragraph — including the finding that statsmodels' reported 'oim' for ARMA terms disagrees with the finite-difference Hessian of its own likelihood, which reproduces tea's numbers instead.
- Verification (exact ML vs statsmodels on identical series): AR(1) ll -218.0418 both; airline ARIMA(0,1,1) ll 121.7537 both; ARMA(1,1) with constant ll -575.8401 both, tea's _cons the mean form (4.80252 vs implied 4.80247); regression with AR(1) errors ll -213.2064 both, with the intercept identity _cons*(1-phi) = sm's c to 5 decimals. The OIM constant SE matches statsmodels' mean-form ARIMA class to 6 digits (.136483 vs .136484).

v1.6.38 — test 72 hardened against optimizer landing-point spread

- Mico's X1 (ext4, OpenBLAS) failed test 72 on the airline lltrend+seasonal(12) problem: var(level) .0006993 vs the container's .0006994, and correspondingly in var(seas) and var(e) — 4th significant digit, exceeding the golden's 1e-7 rounding. The seventh documented substrate instance, and the second delivery on DESIGN_SSPACE.md Decision 6's prediction that an iterative optimizer is where backend arithmetic first becomes visible: the airline UC likelihood is famously flat (it is the problem that forced the three-start strategy), and different BLAS backends walk its ridge to points ~2e-7 apart. Nile, rwalk, and rwdrift — sharp likelihoods — were byte-identical across all rigs.
- Per the design's own doctrine (goldens lock signal, not dust), the airline trio is now asserted with tolerances against the statsmodels optimum (the signal: tea agrees with the reference to ~4 digits) instead of printing digits whose last place is backend-dependent. Mico's observed values pass with 60x margin. No engine or front-end code changed.

v1.6.39 — dfactor, the constraints subsystem, and dsyev (tier release three)

- dfactor arrives per the signed DESIGN_DFACTOR.md with Mico's amended D2: identification through a GENERAL constraints subsystem (constraint define/list/drop +

constraints() on the estimator), because generic sspace models will need constraints even more often than dfactor. Multi-factor from day one per the same-name-same-functionality compatibility ruling: $k \leq 4$ factors, full interacting AR($p \leq 4$) dynamics, Stata's $\text{Var}(v)=I$ normalization, constants and exog in the observation equations, smfactor(stub) for smoothed factors, ragged edges and interior missing values handled by the univariate filter.

- The constraints machinery reparameterizes $R\theta = r$ exactly: minimum-norm particular solution (dgelsd) plus SVD null basis (dgesdd), $\theta = \theta_p + N\psi$, and the BFGS2/multistart/OIM pipeline runs unchanged on ψ with $V_\theta = NV_\psi N'$.
- Verification: $k=1$ demeaned/noconstant loglik equals statsmodels DynamicFactor to 5 decimals (-930.83738 both); $k=2$ identified by $[y1]f2=0$ reproduces the UNCONSTRAINED statsmodels optimum exactly (-1753.32571 both) — the rotation-invariance theorem holding numerically; ragged-edge scalar-observation accounting exact.
- One real bug found and fixed before release: an explosive factor transition can still solve the Lyapunov equation algebraically, yielding a non-PD “covariance” and a garbage likelihood the optimizer exploited (observed companion root 1.96). The evaluation guard now requires positive definiteness (dpotrf) of the Lyapunov solution. arima was never exposed (Monahan's transform keeps its T stationary by construction).
- dsyev entered the wasm LAPACKE shim for the PCA starting values — instance eight, predicted in writing in the design note before the linker could complain.

v1.6.40 — the sspace subset completes the state-space tier; Bug 40

- sspace arrives per DESIGN_SSPACE.md Addendum A (appended this release): state equations in lag-1 states, observation equations in contemporaneous states, covstate(identity|diagonal), covobs(diagonal), noerror/noconstant per equation, smstates(stub), identification via the constraints subsystem — the v1.6.39 machinery this command was the stated reason for. Estimation is the dfactor pipeline verbatim. The tier's delivery order (ucm -> exact-ML arima -> dfactor -> sspace) is complete.
- Verification: (1) internal exactness — the AR(1) as sspace reproduces arima's ll (-218.04181), coefficient (.673299), and OIM standard error (.0431407) identically: the same exact likelihood reached through a different front door; (2) a noisy-AR(1) matches a hand-written statsmodels custom MLEModel to 5 decimals (-310.16419 both); (3) identity-vs-diagonal covstate normalizations reach the same likelihood with the loading absorbing the scale (.506174 = $\sqrt{.256214}$).
- Bug 40, found BY the sspace regression test: tea_srand did not clear the Box-Muller spare, so set seed after an odd number of normal draws produced a stream shifted by one stale value. Test 74's in-session data refused to match the same commands run from a fresh session — the discrepancy that unmasked it. Fixed: the spare is file-scope state reset by tea_srand. Every mid-file re-seed in the shipped suite sat at even parity, so no golden moved; the fix is pure correctness.

v1.6.41 — sspace diffuse, dfactor idio AR, seasonal arima, ucm cycles

- Four staged items land in one release: sspace diffuse initialization (nonstationary user models; local level via sspace == ucm llevel EXACTLY, the third internal-exactness proof); dfactor observation-equation ar(1) (idiosyncratic AR(1) states; ll matches statsmodels error_order=1 to 5 decimals, -969.47356 both); arima sarima(P D Q s) (multiplicative seasonal polynomials expanded onto the engine; the canonical airline model matches to reported precision, ll 244.6965); and ucm cycle (Harvey

damped stochastic cycle, logit-bounded damping/frequency, mixed diffuse-trend + stationary-cycle initialization).

- Two findings for the record: (1) statsmodels' exact-diffuse log-likelihood disagrees with the large-kappa limit by ~ 2.09 on the mixed-initialization cycle model while tea equals the limit to 5 decimals at three kappa magnitudes — and statsmodels' optimizer also failed the cycle problem outright (69 ll units below tea's optimum, collapsing the cycle); (2) ucm did not post e(ll) (masked until test 75 asserted on it) — fixed.

v1.6.42 — the panel depth tier

- areg, xtivreg (fe), xtabond arrive per DESIGN_PANEL.md. areg is internally exact against regress-with-dummies (coefficients AND classical SEs to every printed digit, $df = N - K - G$); xtivreg fe matches a numpy within+2SLS referee exactly; xtabond one-step, robust, and two-step match an independent numpy GMM referee implementing the same Arellano-Bond formulas digit-for-digit (L.y .449738, conventional se .061213, panel-robust .036321, Sargan 16.8021 on 20 df, 22 instruments). Estimated rho .4497 on a true-.5 DGP ($T=8$, $N=200$) sits within one SE — the AB estimator behaving as the textbook says it should.
- Posted coefficient names are dot-free (Ly, Dx1) so `_b[_se]` can address them, while table labels keep Stata's L./D. dress — the sspace paren-free lesson, applied on arrival rather than after a bug report.

v1.6.43 — manual refresh and thematic help

- The manual catches up with five releases of estimation commands: a new State-space models chapter (ucm with cycles, sspace with diffuse, dfactor with idiosyncratic AR, the constraint subsystem, and the engine/verification story), a filled-out time-series chapter (sarima, newey, unit roots, tsfilter, VARs/IRFs/lpirf, vecrank), panel-tier sections (areg, xtivreg fe, xtabond with its staged-features caveats), correlate/pwcorr/ttest under Exploring data, and a cleaned-up intro. Three stale claims excised: “no xtabond”, “no VAR/VECM”, and a commands-that-don't-exist list that still included graphics, var, and dfuller — all shipped releases ago.
- help is now thematic: help prints eleven themed groups, help TOPIC (data, create, explore, reshape, graphics, estimation, panel, timeseries, statespace, postest, files) lists one theme with each command's one-line syntax, and help CMD is unchanged. Two machine-facing subcommands: help _list (one name per line) now feeds tools/gen_cmdref.sh, making the generated manual reference independent of the human-facing grid format; and help _check prints any command missing from every theme and fails rc 9 — test 77 goldens it, so registering a command without categorizing it breaks the suite (the silent-failure doctrine applied to documentation). Four stale one-liners refreshed along the way (arima still said “conditional ML”; ucm/sspace/dfactor lacked their v1.6.41 options).

v1.6.44 — styled help headings, everywhere a human is watching

- Theme headings in help (and the banner and footer mini-headers) are now bold. The vocabulary is deliberately the VT100-vintage SGR subset (1/0) that renders identically on every terminal in the release matrix and in the browser, whose terminal is xterm.js — a full ANSI emulator that needed no work at all.
- The analysis that mattered was WHEN to emit, not how: native tea styles only when stdout is a real terminal, TERM is not dumb, and NO_COLOR is unset (the no-color.org convention), so pipes, logs, and test goldens remain byte-clean. Under Emscripten,

stdout claims to be a tty even inside the node test harness, which would have leaked escapes into the WASM golden comparisons and broken the byte-identical promise — so the WASM default is unstyled and the browser page, the one layer that knows a rendering terminal is attached, flips styling on through the exported `tea_web_set_style(1)` after boot. The 77 goldens double as the regression test for this cleanliness: any leaked escape anywhere diffs loudly.

v1.6.45 — build-system hygiene: exact incrementality, and a latent bug

- Mico observed the build “recompiles everything”. The audit: the native build was properly incremental all along, but `make wasm` was a single monolithic `emcc` rule (any trigger recompiled every source, and `VERSION` was not a prerequisite — the staleness footgun that forced `wasm-clean` into the release gate), `make dist` deleted the object tree to keep artifacts out of the tarball (forcing a full rebuild after every release), and `make debug` is a whole-program build by design. Underneath sat a latent under-compilation bug: `TEA_VERSION` arrived as a `-D` flag, which `make` cannot track, so `echo new > VERSION`; `make` rebuilt nothing and shipped a binary reporting the old version — masked for its whole life by the clean-build release discipline.
- The fix is one mechanism: `src/tea_version.h` is generated from `VERSION` with a content compare (`mtime` moves only when the version changes), included by exactly the three units that print the version; `-MMD` does the rest. A version bump now recompiles three units and links — measured, and the stale-version repro now produces the right answer without a clean. The `wasm` build is per-object with `emcc`’s own `-MMD` tracking (touch one source: one recompile plus link), which retires `wasm-clean` as a correctness crutch. `dist` now uses `tar` excludes instead of `rm`, so releasing no longer destroys the build tree. The release gate keeps its clean builds as deliberate policy — belt and braces for shipped artifacts — but day-to-day iteration is exact.
- Noted, pre-existing, `wasm-only`: a `-Wconstant-conversion` warning in a string-hash constant (unsigned long is 32-bit under `wasm32`); functionally harmless truncation, surfaced only now that per-object logs are readable. Left as-is deliberately.