

Contents

tea — tiny econometric assistant	1
What tea v1.0 does	1
When to leave tea (escape hatches)	3
Stata cheat sheet (what’s the same, what’s not)	4
Updating from a release tarball	4
Coming from Stata?	4
Manual	4
License	5
Build & run	5
Command-line options	5
Supported file formats	5
Dependencies	5
REPL features	6
Architecture (dependency order)	6
Semantics decisions (locked)	6
Implemented in Milestone 1 (verified by tests/demo.do)	7
Implemented in Milestone 2 (also verified by tests/demo.do)	7
Implemented in Milestone 3 (also verified by tests/demo.do)	8
Implemented in v1.1 (plots, browser, reproducibility)	8
Not yet implemented (later)	8
Implemented in v0.2 (shell & REPL polish)	9
Implemented in v0.3 (REPL polish, shell, spreadsheets)	9
Implemented in v0.4 (econometrics)	9
Architecture	9
Numerical stack	10
Implemented in v0.5.3 (reshape crash hotfix)	10
Implemented in v0.5.2 (round 2 of real-data fixes)	10
Implemented in v0.5 (bug-fix release after real-data testing)	10
Build hygiene	11

tea — tiny econometric assistant

Data prep, brewed cleanly.

A small C17 program for the data-prep phase economists use Stata for: import, clean, define variables, prepare for econometrics. Interactive REPL plus batch do-files for repeatability. Deliberately *not* pandas — the syntax is close to Stata’s and the semantics are faithful where it matters.

tea is an independent open-source project. It is not affiliated with, endorsed by, or sponsored by StataCorp LLC. *Stata* is a registered trademark of StataCorp LLC.

What tea v1.0 does

Data manipulation:

Area	What works
File I/O	.dta (Stata-native, full round-trip of types, formats, var labels, value labels), .tea (native binary), .csv/.tsv (RFC 4180 quoting), .xlsx/.ods (via LibreOffice headless)
Variable creation	generate, re-place (with by:, if, in), egen (mean/sum/total/count/min/max/sd/group/ro
Type conversion	encode, decode, destring, tostring, recode
Reshape & merge	reshape long wide (in-place), merge 1:1 m:1 1:m (m:m rejected), collapse
Calendar	mdy, ym, yq, yw, yh, dofm/mofd/..., %td/%tm/%tq/%tw/%th/%ty
Macros & control flow	formats, date(...) parser `local' and \$global, r() and e() references, foreach, forvalues, if/else, capture, quietly
Time-series ops	L.x, F.x, D.x, S.x with (1/3) ranges; gap-aware via tsset
Factor variables	i.var, ib<n>.var, c.var, # and ## interactions across estimators

Estimators (all support if/in/weights, vce(robust|cluster)):

Command	Notes
regress	OLS with HC1 robust, CR1 cluster, weights (fw/aw/pw/iw)
xtreg, fe	Fixed effects (within transform), with all vce options
xtreg, re	Random effects (Swamy-Arora FGLS)
xtreg, be	Between effects (regression on panel means)
hausman	After running both fe and re

Command	Notes
logit / probit	Newton-Raphson MLE on a reusable driver
poisson	Same MLE driver
ivregress 2sls	Two-stage least squares with first-stage F diagnostic
arima	ARIMA(p,d,q) via conditional likelihood (no Kalman in v1.0)

Post-estimation:

Command	Notes
predict	xb, residuals, panel u/e for xtreg; Pr() and xb for logit/probit
margins	AME and MEM with delta-method SEs (dydx(* varlist) [atmeans])
test, lincom	Linear hypothesis testing
_b[name], _se[name]	Coefficient and SE macros
estimates store/restore/dir/drop/table	Named saved estimates ledger
estout	Publication-ready La- TeX/markdown/plain tables with stars and fit stats

Reproducibility & I/O:

log using FILE for session capture (commands + output), set seed N plus runiform()/rnormal() for reproducible Monte Carlo, shell for arbitrary shell commands, do FILENAME to source nested do-files, --strict-stata flag to enforce Stata-only syntax.

When to leave tea (escape hatches)

tea is excellent at the data-prep-to-quick-estimate loop and the common estimators above. For more specialized work, escape to:

Need	Use
GARCH / EGARCH / TGARCH	Python's arch package, or R rugarch
Exact-ML ARIMA (Kalman)	R's arima() or forecast::Arima
Bayesian inference	PyMC, Stan (via rstan/cmdstanpy)
Mixed-effects models	R's lme4, Python's statsmodels.MixedLM
Survival analysis	R's survival, Python's lifelines
Machine learning	scikit-learn, xgboost, lightgbm
Network / graph models	networkx, R's igraph
Spatial econometrics	R's spdep, Python's pysal
Plotting	Pipe to gnuplot, or export delimited and load in Python/R

The shell command makes the bridge frictionless: write a Python or R script that reads your export delimited output, and call it from a tea do-file with shell python myscript.py. Pull the results back via another import delimited.

Stata cheat sheet (what's the same, what's not)

Same as Stata: syntax, missing-value algebra (including the if $x > 5$ footgun that includes missing), value labels, by:/bysort:, if/in, weights, varlists with */? and ranges, factor variables, TS ops.

Tea-specific differences (call them out): - m:m merge is rejected (design choice — it's almost never what you want). - xtreg, be uses simple OLS on panel means (no T_i weighting variant). - arima uses conditional likelihood; for exact ML use R's arima(). - No .gph graphs — use gnuplot or escape to Python/R for plotting.

Updating from a release tarball

./tools/update.sh path/to/tea-vX.Y.Z.tar.gz extracts a release over this repo from anywhere, detects the tarball layout (never creates a nested tea/tea/), and prints git status so you can review what changed before committing. Or by hand: official tarballs contain a top-level tea/, so cd to the repo's *parent* directory and plain tar xzf lands everything in the right place.

Coming from Stata?

Read STATA-QUICKSTART.md / STATA-QUICKSTART.pdf — a four-page tour that uses your existing Stata knowledge to get you running panel regressions on the bundled IMF World Economic Outlook in about a minute, in your browser, with nothing installed. Rebuild the PDF with make quickstart.

Manual

The user guide and command reference: tea-manual.md (single-file Markdown) and tea-manual.pdf. The master source lives in manual/; the command-reference appendix is generated from the binary itself by tools/gen_cmdref.sh, so it cannot drift from the implementation. Rebuild both with make manual (requires pandoc + texlive).

License

GPLv3 or later. See LICENSE at the repository root for the full text. Per-file copyright headers identify the upstream authors.

Build & run

```
make                # produces ./tea
./tea tests/demo.do # batch: run a do-file
./tea              # interactive REPL (type 'help' or Ctrl-D)
make test          # build + run the full regression suite
make smoke         # quick sanity: build + run the demo
```

Command-line options

```
tea [options] [do-file]
```

```
--strict-stata    reject tea-only extensions (default)
--tea-extensions  allow tea-only extensions
--version         print version and exit
--help           print usage and exit
```

Tea is **strict-Stata by default**: do-files that use tea-only features are rejected. This guarantees every passing do-file runs unchanged in Stata. See COMPATIBILITY.md for the full contract and the (short) list of currently-gated extensions.

Supported file formats

Extension	Read	Write	Notes
.dta	yes	yes	Stata format; readstat-backed; default format.
.tea	yes	yes	Native tea binary; fast internal exchange.
.csv	yes	yes	Comma-separated; via import/export.
.tsv	yes	yes	Tab-separated; via import/export.
.xlsx	yes	no	via import excel; uses LibreOffice headless.
.ods	yes	no	via import ods; LibreOffice.

use `foo` and `save foo` without an extension default to `.dta` (Stata-compatible). Use explicit `.tea` if you want the native format. On read, every Stata storage type (byte/int/long/float/double, str#/strL) is upcast to tea's internal double + string representation; on write, numeric columns are compressed per-column to the smallest fitting storage type, matching what Stata's own `compress` command would produce.

Dependencies

- A C17 compiler (gcc 8+ or clang 7+).
- `libm` (always present).
- `libreadline` — line editing, history, tab completion in the REPL.
- `libopenblas` + `liblapacke` — linear algebra for regress and later econometric estimators.
- `libgsl` — distributions (t, F, chi²) for p-values and CIs.

- **libreadstat** — Stata .dta (and SAS/SPSS) file I/O for use/save. The MIT-licensed C library by Evan Miller, also used under the hood by R’s haven and Python’s pyreadstat.

Install on **Debian/Ubuntu**:

```
apt install libreadline-dev libopenblas-dev liblapack-dev libgsl-dev libreadstat-dev
```

On **macOS** (Homebrew):

```
brew install readline openblas lapack gsl readstat
```

Run `make check-deps` after installing to verify every required header is reachable. If any dep is missing, that target reports exactly which one and the install command to use.

The build is warning-clean under `-Wall -Wextra` (two style/false-positive categories suppressed with documented `-Wno-` flags) and runs clean under AddressSanitizer + UBSan on the full demo (no errors, leaks, or UB).

REPL features

History is persistent at `~/.tea_history` (capped at 2000 lines). Tab completes command names at the start of a line and variable names in the active frame elsewhere. Use `help` to list commands, `help CMD` for one’s syntax.

Architecture (dependency order)

The source is modular; each layer is independently testable before the next sits on it.

- `value.h` — the scalar/missing model. Numeric cells are a single double; missing values (`.`, `.a...z`) are encoded as quiet NaNs carrying a 0..26 code. The module guarantees the three Stata invariants: missing sorts after every real number (`.` < `.a` < ... < `.z`); any arithmetic operand missing makes the result `.`; comparisons treat missing as +infinity-with-code.
- `dataset.{h,c}` — typed columnar store, multiple frames, per-frame sort state and panel (tsset) state.
- `lex.{h,c}` — one tokenizer shared by the command and expression parsers.
- `expr.h` + `parse.c` — recursive-descent parser with Stata operator precedence; parses TS operators (L. D. F. S., with # orders), variable subscripts, and function calls into an AST.
- `eval.c` — the evaluator: `_n/_N`, group-relative subscripts, gap-aware TS operators (resolved through a lazily-built (panel,time)→row index), the full missing/string algebra, and the function library including the calendar suite.
- `cmd.{h,c}` — command layer: varlist sugar (`a-b, v*, _all`), the `if/in/options` split, the physical stable multi-key sort, contiguous by-group construction, and dispatch.
- `commands.c` — the command implementations and the dispatch table.
- `interp.{h,c}` — macros (``local'`, `$global`, `r()`), `=exp` evaluation, control flow (`foreach/forvalues/if-else/capture/quietly`), comment handling (`(* // /* */ ///)`), and the do-file/REPL driver.
- `main.c` — entry point.

Semantics decisions (locked)

- Storage collapses to double + string only (byte/int/long/float/double are cosmetic). Full Stata missing algebra, including the intentional footgun: `if x > 5 includes missing` because missing sorts above every number.
- Multiple frames, each with its own sort and panel state.

- `tsset/xtset` with gap-aware L./D./F./S. resolving against panel+time. Order-breaking commands (sort, drop, merge, a replace on a sort key) invalidate the sort and the panel settings.
- `by:` is faithful — it errors if the data isn't sorted; `bysort` sorts first. `by g (s):` groups by `g` and sorts within by `s`.
- `reshape` is in-place only (the user must `frame copy` first). `merge` supports 1:1 and m:1 only — m:m is deliberately excluded as a footgun.
- C17, no novelties — toolchain stability over newer-standard ergonomics.
- **Backend-independent output (v1.1).** The same do-file produces byte-identical output on every machine, BLAS library, and CPU. To make that true, quantities that are *mathematical zeros* are displayed as exactly 0 rather than as floating-point noise, which differs across numerical backends. Concretely, when a regression is a perfect fit (residual SS below 1e-12 of total SS), tea snaps the residual SS, the residual vector, the SEs, and any coefficient below 1e-8 of the largest (the unique exact solution's true zeros) to exactly 0; the model F is reported as `inf` with `p = 0.0000`. The same relative-tolerance snapping applies to `sigma_u` when panel intercepts are identical and to the hausman difference matrix when FE and RE agree to machine precision. This changes *display of degenerate cases only* — normal estimation results are untouched to the last bit. We consider “identical results everywhere” a core promise for reproducible research; the tradeoff is that raw sub-epsilon values (e.g. a coefficient of `-5.8e-15` in an exact fit) are not shown. If a computed quantity is genuinely tiny rather than a mathematical zero, the relative thresholds leave it alone. Verified by running the full regression suite under two maximally different backends (gcc + OpenBLAS natively; clang + reference CLAPACK/BLAS + musl under WebAssembly) with byte-identical output.

Implemented in Milestone 1 (verified by tests/demo.do)

`import/export delimited, use/save (native TEA1 binary), clear, set obs; generate/replace (with type tokens, by:, if/in), egen (mean/sum/total/count/min/max/sd/group/rowtotal/rowmean, by()), drop/keep/rename/order/label/format; describe/list/count/summarize (sets r())/tabulate; sort/gsort; tsset/xtset; collapse; frames; the calendar suite (ym yq yh yw mdy constructors, %tm %tq %td %tw %th %ty display formats, dofm/mofd-type converters, date() parser); macros and loops; display with mixed strings and expressions.`

The demo exercises and passes the hard cases: the missing-value footgun, gap-aware TS operators across panel units, `_n/_N` and running `sum()` resetting per by-group, `replace ... if` change counts, and the calendar formats.

Implemented in Milestone 2 (also verified by tests/demo.do)

- `merge 1:1 | m:1 | 1:m` keyvars using `FILE` with `_merge` (1/2/3), and options `nogen-erate`, `generate(name)`, `keep(master match using | 1 2 3)`, `assert(...)`, `keepus-ing(varlist)`, `update`, `update replace`. m:m is rejected by design with an explanatory message. The using file may be a native `.tea` or a `.csv/.tsv`. Output is the Stata-style match table.
- `reshape long | wide stub[s], i(idvars) j(jvar)`, in-place (the user must `frame copy` first if they want to keep the original — `frame copy` is a later item). Numeric `j` levels are auto-discovered. Verified lossless on a wide↔long round trip, missing values preserved.

Implemented in Milestone 3 (also verified by tests/demo.do)

- **Value labels:** label define name # "txt" ..., label values varlist name (or . to clear), label list. Labels render in list, tabulate, export, and codebook.
- **recode:** rules (1/3 = 1), (7 = .), (missing = 9), (nonmissing = 1), (else = .); in-place or , gen(newvar).
- **codebook [varlist]:** type, unique-value count, missing count, range and mean (numeric).
- **Weights:** [fw=exp] / [aw=exp] / [pw=exp] / [iw=exp] on summarize, count (fw), tabulate, and collapse. Frequency weights use Σw as N with a $\Sigma w - 1$ variance denominator; analytic/probability weights report n and normalize so the weighted variance uses $n - 1$.
- **t*() literal-date constructors:** td(15jan2020) / td(2020-01-15), tm(2020m3), tq(2020q4), tw(2020w5), th(2020h1), ty(2020) — the literal is captured raw by the parser, not parsed as an expression.
- **frame copy old new [, replace], frame rename old new, frame put varlist, into(name), leak-free frame drop.** This enables the safe pattern for the in-place-only reshape: frame copy default backup → frame change backup → reshape

Implemented in v1.1 (plots, browser, reproducibility)

- **SVG plotting, no dependencies:** scatter, line, histogram (hist) with [if]/[in], title() xtitle() ytitle(), saving(), bins(), freq, sort, noview. A self-contained renderer (src/plot.c) writes publication-grade vector SVG — no gnuplot, no graphics library. In the REPL, plots open in the OS viewer; in do-files they only write files, so batch runs stay deterministic. Golden-SVG regression test included.
- **WebAssembly build** (make wasm): tea runs entirely in the browser — REPL via xterm.js, .dta/.csv upload into an in-memory filesystem, plots rendered inline, exports downloadable. Backed by reference CLAPACK/BLAS through a thin linalg.h shim; readline stubbed out. The full regression suite runs inside the WASM module (web/run_wasm_tests.cjs) and passes 40/40, byte-identical with native. Static bundle in web/ deploys directly to GitHub Pages.
- **Push-mode session API:** the REPL loop is now a state machine (TeaSession in interp.h) fed one line at a time — same semantics (#delimit ;, ///, {} blocks), enabling event-driven front-ends and future embeddings.
- **Backend-independent output:** see “Semantics decisions” above. The suite passes byte-identically under gcc+OpenBLAS and clang+reference-BLAS.
- **Bundled practice datasets** (sysuse): six serious datasets are embedded in the binary itself — weo (the full IMF World Economic Outlook, April 2026: 197 economies + 13 aggregates, 145 indicators, 1980-2031), grunfeld (panel), airline (ARIMA), longley (ill-conditioned regression), nmes1988 (health economics: poisson/logit), pwt (Penn World Table sample, CC BY 4.0), and weo — the complete IMF World Economic Outlook database, April 2026 vintage: 197 economies, 44 indicators, 1980-2031 including projections. sysuse dir lists them; sysuse grunfeld, clear loads one. No files installed, works identically in the browser build. Provenance and citations in data/SOURCES.md; regenerate the embedded source with tools/gen_sysdata.py.

Not yet implemented (later)

pw/aw variance for clustered/survey designs beyond the basic form, reshape with string j(), encode/decode, egen rank/pctile family, tabulate two-way, and merge assertion as a hard error (currently warns). The dispatch and qualifier plumbing accommodates these without structural change.

Implemented in v0.2 (shell & REPL polish)

- **help [cmd]** — lists all commands or shows one’s syntax. Knows aliases (`l` → `list`, `su` → `summarize`, etc.).
- **exit [N] [, clear]** — leaves with optional exit code; `,` `clear` drops data first. Also quit as an alias.
- **version / about** — version, build date, GPL notice, non-affiliation statement.
- **pwd / cd [DIR]** — current directory; `cd` with no argument goes to `$HOME`.
- **log using FILE [, replace append] / log close** — tees stdout to a file (implemented with a pipe + tee subprocess, so all output is captured transparently).
- **GNU readline REPL** — line editing, persistent history, tab completion on command names (start of line) and variable names (elsewhere).

Implemented in v0.3 (REPL polish, shell, spreadsheets)

- **Shell escape:** `! cmd` (start of line) or `shell cmd` runs the command through `$SHELL -c`, returning its exit code in `_rc`. This makes a separate “calculator” command unnecessary — `! echo $((2+2))` works, along with `! ls`, `! wc -l file.csv`, etc.
- **assert exp [if]**: fails loud with `rc=9` if the expression is false for any selected observation. Aborts do-files; in the REPL it prints and lets you continue.
- **#delimit ; / #delimit cr**: switch statement terminator between semicolon and new-line, so long commands can span lines without `///`.
- **Spreadsheet import:** `import excel using FILE [, sheet(name)]` and `import ods using FILE` — shells out to `ssconvert` (gnnumeric) or `libreoffice --headless --convert-to csv`, then loads the resulting CSV. Errors out cleanly if neither is on `$PATH`. Note: `sheet()` is honoured by `ssconvert`; `libreoffice` exports only the first sheet.
- **Context-aware tab completion** in the REPL:
 - command names at the start of a line, and as the argument of `help`;
 - filenames after `using`, `cd`, `save`, `use`, `log using`, and `!`;
 - variable names of the current frame everywhere else.
- **Line-annotated errors** in do-file mode: errors print as `line N: ...` and aborts include the line number and `rc`, making debugging long do-files practical.

Implemented in v0.4 (econometrics)

- **regress y x1 x2 ... [if] [in] [weight] [, noconstant robust cluster(var)]** — OLS via QR (LAPACKE `dgels`), with `dgelsd` + column-pivoting QR fallback to detect and drop collinear regressors, marked (omitted) in the table like Stata. Classical SEs (default), HC1 robust SEs (`,` `robust`), and CR1 clustered SEs (`,` `cluster(var)`). Stata-style coefficient table with t-stats, p-values, and 95% CIs. Sets `e(N)`, `e(r2)`, `e(r2_a)`, `e(rmse)`, `e(F)`, `e(p)`, `e(df_m)`, `e(df_r)`.
- **predict newvar [, xb residuals]** — fitted values or residuals from the last estimation. Listwise-deletion-safe: any obs with a missing X gets a missing prediction, matching Stata.
- **test v1 v2 ...** — joint Wald F-test that listed coefficients are zero. Uses the stored variance, so honours whichever SE option was passed to the original `regress`.
- **lincom <linear combo>** — point estimate and SE of any linear combination of coefficients, e.g. `lincom x1 - x2`, `lincom 0.5*x1`.

Architecture

Estimators populate a workspace-level `Estimates` struct (in `src/estimates.h`). Postestimation commands (`test`, `predict`, `lincom`) consume it without knowing what produced it.

This is the same shape as Stata's `e()` macro family, generalized — and it's exactly the hook `xtreg` will need next.

Numerical stack

- **OpenBLAS** (BLAS) + **LAPACK** (LAPACK C interface) for linear algebra.
- **GSL** for distributions (t, F, chi²) and inverse-CDFs used in p-values and CIs.

Install on Debian/Ubuntu: `apt install libopenblas-dev liblapacke-dev libgsl-dev`. On macOS: `brew install openblas lapack gsl`.

Implemented in v0.5.3 (reshape crash hotfix)

Real-data testing surfaced a segfault in `reshape long` when:

- the stub-discovery loop matched variables of **mixed types** at different j-levels (e.g. `z1`, `z2` numeric but `z4`, `z5` loaded as string because the CSV had mixed values like "jajca"), or
- the user passed `i(z*)` so the `i()` vars overlapped the stub vars.

Both now error cleanly with `rc=109` ("type mismatch in stub levels") or `rc=198` ("variable appears in both `i()` and stub list"). No more crash.

Implemented in v0.5.2 (round 2 of real-data fixes)

After a second round of testing against real Stata workflows:

- **list column widths** now computed as `max(header, max(cell))` over the selection, so columns containing both `0.0001234` and `1.67e-05` align cleanly. Matches Stata's behavior.
- **? and arbitrary wildcards** in varlists. Switched from prefix-only `*` matching to `fnmatch()`, so `list z?`, `describe z*1`, etc. work.
- **dir [pattern] / ls [pattern]** — list files in current directory with optional wildcard.
- **use and import require , clear** if data is in memory (`rc=4`). Match Stata's behavior; prevents accidental data loss.
- **Math functions**: added `sin`, `cos`, `tan`, `asin`, `acos`, `atan`, `atan2`, `sinh`, `cosh`, `tanh`, plus distributions `normal`, `invnormal`, `ttail`, `invttail`, `Ftail`, `chi2tail`, `chi2`. All Stata-named.
- **Log captures commands too** — `log using` now writes both the command (with `.` prefix) and its output, matching Stata.
- **reshape long** parsing of stub specifications fixed. Previously treating wildcards in the stub spec as literal characters; now strips trailing `*/?` to recover the stub. `reshape long z, i(x) j(tt)` and `reshape long z*, i(x) j(tt)` both work and round-trip via `reshape wide`.

Implemented in v0.5 (bug-fix release after real-data testing)

After the first round of testing against real Stata workflows on a Mac, the following issues were fixed:

- **log using FILE no longer hangs the interactive REPL**. Replaced the `fork+pipe tee` with a direct `tea_printf` mirror; output now goes to `stdout` and the log file simultaneously without fighting readline.

- **local v "GDP_N"** now strips the surrounding double quotes, so `\v` expands to `GDP_N`, not `"GDP_N"`. This was the root cause of `apparentforeach` failures — locals containing quoted lists couldn't be re-used to drive loops.
- **bysort g: summarize** honors the `by:` prefix, producing per-group headers and per-group statistics.
- **describe varlist** filters to the requested variables (e.g. `describe GDP*`).
- **encode strvar, generate(newvar) [label(name)]** maps a string variable's distinct values to integers 1..K (alphabetically) and attaches a value label. **decode encvar, generate(newstrvar)** is the inverse. Round-trip is value-preserving.
- **File-system commands:** `mkdir DIR [, recursive]`, `rmdir DIR`, `erase FILE / rm FILE`, `copy SRC DST [, replace]`, and `do FILENAME` (source a do-file from inside a session).
- **macOS spreadsheet conversion** now auto-detects soffice and the LibreOffice app bundle at `/Applications/LibreOffice.app/Contents/MacOS/soffice` if `libreoffice` isn't on `$PATH`.
- **Do-file errors** include the line number and `rc` of the aborting command, e.g. `do-file aborted at line 4 (rc=110)`.

Build hygiene

The Makefile uses `-MMD -MP` so editing any header rebuilds every translation unit that includes it. (An earlier session hit a stale-object ABI mismatch from a header change that only `make clean` would have fixed; that class of bug is now structurally prevented.)