



version 1.6.4

User guide and reference

Mico Mrkaic

July 12, 2026

Abstract. `tea` is a small command-line program for the data-prep-to-quick-estimate loop economists use Stata for: import, clean, generate variables, run regressions, panel methods, MLE, IV, and ARIMA — plus publication-grade SVG graphics and six bundled practice datasets, including the complete IMF World Economic Outlook database. The syntax is close to Stata’s where they overlap; the semantics are faithful where it matters (missing-value algebra, the `by:` prefix, gap-aware time-series operators, factor variables). `tea` also runs entirely in a web browser via WebAssembly, with no installation, and its output is byte-identical across machines and numerical libraries. This manual covers installation, a guided tour, every command, the bundled data, and escape hatches to Python/R for things `tea` does not handle.

Audience. Applied economists, policy analysts, and graduate students familiar with Stata. No Stata experience is strictly required, but the manual assumes you know what `regress` is and what panel data looks like.

Status. v1.6.4 follows the v1.0 field-testing release. It has passed 42 regression tests byte-identically on three configurations (native gcc + OpenBLAS in release and sanitizer builds, and WebAssembly clang + reference BLAS), is clean under AddressSanitizer and UndefinedBehaviorSanitizer, and builds warning-free under `-Wall -Wextra -Werror`. Field testing continues; please report what breaks.

AI-assisted development. This software was developed with the assistance of an AI coding assistant (Anthropic’s Claude). The design decisions, mathematical specifications, scope choices, and final acceptance of each component were the author’s; the AI contributed substantial portions of the implementation, drafted regression tests, and produced the first draft of this manual. All code was reviewed before inclusion. This manual was likewise drafted with AI assistance and edited by the author. Bugs and design flaws are the author’s responsibility.

License. GPLv3 or later. Copyright © 2026 Mico Mrkaic. See the `LICENSE` file in the distribution for the full text. `tea` is an independent open-source project, not affiliated with or endorsed by StataCorp LLC. The bundled datasets retain their original licenses and citations; see the “Bundled datasets” chapter and `data/SOURCES.md`.

Contents

Introduction	7
What <code>tea</code> is	7
What <code>tea</code> is not	7
Design choices, locked	7
Installation	8
Build from source	8
Dependencies	8
Sanity check	8
Running interactively	8
Your first <code>tea</code> session	9
The <code>do-file</code> language	10
Comments	10
Macros	10
Control flow	10
Capture and quietly	11
The <code>by:</code> prefix	11
Data management	11
Importing data	11
Saving data	12
Generating and modifying variables	12
<code>egen</code> extended generators	12
Type conversion	12
Labels	13
Exploring data	13
<code>status</code>	13
Progress indicator	13
<code>describe</code>	13
<code>list</code>	13
<code>summarize</code>	14
<code>tabulate</code>	14
<code>codebook</code>	14
<code>count</code>	14
Reshaping, merging, aggregating	14
<code>sort</code> and <code>gsort</code>	14
<code>reshape</code>	14
<code>preserve</code> / <code>restore</code>	15
<code>merge</code>	15
<code>collapse</code>	15
Frames	15
Time series and panel data	16
Declaring panel / time structure	16
<code>xtdescribe</code>	16
Factor variables	16
Atoms	16

Interactions	17
Coefficient naming	17
Known limitation: factor $\times$ TS-op composition	17
Linear regression	17
Basic syntax	17
Output	18
Robust and clustered SEs	18
Constraints and hypothesis tests	18
Panel regression	18
Fixed effects	18
Random effects	18
Between effects	18
Hausman test	19
Maximum likelihood estimators	19
logit	19
probit	19
poisson	19
Marginal effects via margins	19
Instrumental variables	19
Syntax	20
Example	20
Interpreting the first-stage F	20
Variance estimators	20
Time series models	20
ARIMA	20
Example	21
Post-estimation	21
predict	21
Hypothesis tests	21
Output and reproducibility	21
Saving estimates	21
estout: publication tables	22
log	22
Random numbers and Monte Carlo	22
Reproducibility	23
Functions	23
A Monte Carlo example	23
Graphics	23
Bundled practice datasets (sysuse)	24
The WEO database	24
The other five	24
The browser edition	25
Backend-independent output	25

Escape hatches: when to leave tea	25
When to escape	26
A bridge example	26
Known limitations	26
Data and metadata	26
Econometric estimators	26
Factor variables	27
Programming	27
Graphics	27
Reporting bugs	27
Function reference	28
Arithmetic and rounding	28
Exponential and logarithm	28
Trigonometric and hyperbolic	28
Probability distributions	28
Random number generation	29
Missing-value tests and ranges	29
String construction and inspection	29
String transformation	29
Regular expressions (POSIX extended)	30
Numeric ↔ string conversion	30
Calendar: building dates	30
Calendar: extracting components from a daily date	30
Calendar: converting between frequencies	31
Special: running and time-series-context	31
egen function reference	31
Group-wise aggregators	31
Position and identity	32
Row-wise aggregators (across columns of a single row)	32
Summarize / tabstat statistics keywords	32
Stata cheat sheet	32
Same as Stata	33
Tea-specific differences	33
Stata commands that don't exist in tea	33
Command reference	33
scatter	33
line	34
histogram	34
generate	34
replace	34
egen	34
list	34
summarize	34
tabstat	34
count	34
describe	34
drop	34

keep	35
rename	35
order	35
aorder	35
label	35
format	35
sort	35
gsort	35
tabulate	35
tsset	35
xtset	35
xtdescribe	36
xtreg	36
hausman	36
logit	36
probit	36
ivregress	36
poisson	36
arima	36
margins	36
estimates	37
est	37
estout	37
collapse	37
merge	37
reshape	37
recode	37
encode	37
decode	37
destring	37
tostring	38
codebook	38
import	38
export	38
save	38
outreg2	38
which	38
ssc	38
isid	38
duplicates	38
tempfile	38
tempname	38
pwcorr	39
file	39
confirm	39
history	39
sysuse	39
use	39
clear	39
status	39
frame	39
regress	39
predict	40

test	40
lincom	40
help	40
pwd	40
cd	40
mkdir	40
dir	40
rmdir	40
erase	40
copy	40
do	40
preserve	40
restore	40
version	41
log	41
exit	41
Native statements	41

Introduction

What tea is

tea is a C17 program that runs as either an interactive REPL or a batch processor for do-files. It does the data manipulation and econometrics that most applied economists do most days — and nothing else. No plotting, no Bayesian inference, no machine learning. For those, escape to Python or R; see the “Escape hatches” chapter.

The design priority is *Stata fidelity in the overlapping subset*. A do-file that runs in **tea** should run in Stata with the same results, modulo a handful of small documented differences. Where Stata’s behavior includes a footgun (if `x > 5` includes missing values, for instance, because missing sorts above every number), **tea** keeps the footgun. That choice is deliberate: it makes **tea** a drop-in for the data-prep phase, even if you eventually move to Stata or R for the final analysis.

What tea is not

tea v1.0 does *not* implement: graphics of any kind, exact-ML ARIMA (we use conditional likelihood), seasonal ARIMA, GARCH, mixed-effects models, survival analysis, structural breaks, VAR, VECM, cointegration tests, Bayesian inference, or machine learning. It also does not implement Stata’s full programming language — **tea** has macros, loops, and **capture**, but no **program define**, no **syntax** parser, no **Mata**.

The **shell** command makes it easy to bridge to Python or R when you need any of the above. See the “Escape hatches” chapter.

Design choices, locked

- **Storage** is one of two types: **double** (numeric, with the full Stata missing-value algebra) or string. Stata’s **byte**, **int**, **long**, **float** are cosmetic on read and compressed back on write of **.dta** files.
- **Multiple frames** are supported. Each has its own sort state and panel/time settings.
- **by: is strict**. It errors if the data isn’t sorted; **bysort** sorts first.
- **reshape is in-place**. Use **frame copy** first if you want to keep the original.
- **m:m merges are rejected**. They are almost never what the user actually wants.

- **C17**, no novelties. Toolchain stability over newer-standard ergonomics.

Installation

Build from source

You will have received a tarball, `tea-v1.0.tar.gz`.

```
tar xzf tea-v1.0.tar.gz
cd tea
make
```

The build produces `./tea` in the project root. Move it somewhere on your `$PATH` (e.g. `~/bin`) or invoke directly with `./tea`.

Dependencies

You need a C17 compiler (`gcc 8+` or `clang 7+`) and the following libraries with their development headers:

Library	Why
<code>libreadline</code>	line editing, history, tab completion in the REPL
<code>libopenblas</code>	matrix kernels
<code>liblapacke</code>	LAPACK C interface (regression, ARIMA, IV)
<code>libgsl</code>	distributions for p -values and confidence intervals
<code>libreadstat</code>	Stata <code>.dta</code> (and SAS / SPSS) I/O

On Debian / Ubuntu:

```
apt install libreadline-dev libopenblas-dev liblapacke-dev \
    libgsl-dev libreadstat-dev
```

On macOS (Homebrew):

```
brew install readline openblas lapack gsl readstat
```

Run `make check-deps` after installing to verify every required header is reachable. If any dependency is missing, that target reports which one and the install command to use.

Sanity check

```
./tea --version      # prints "tea 1.0 -- tiny econometric assistant"
./tea tests/demo.do  # batch: runs a short demo (no errors expected)
make test           # full regression suite, 39 tests
```

If `make test` reports `tea regression: 39/39 passed`, your build is correct.

Running interactively

Start the REPL by invoking `tea` with no arguments:

```
$ tea
      ( (
      ) )    tea 1.2.2 - tiny econometric assistant
.....    free Stata-like data analysis at the command line
|         |]  GPLv3 · Mico Mrkaic · github.com/micomrkaic/tea
```



```

\      /
 \_____/  type help · sysuse dir for practice datasets · Ctrl-D exits
. _

```

History is persistent at `~/.tea_history` (capped at 2000 lines). Readline editing is available throughout: arrow keys, Ctrl-A/E/K/U/W, Ctrl-L to clear, Up/Down history. Tab completes command names at the start of a line, dataset names after `sysuse`, and variable names elsewhere. `help` lists commands, `help CMD` shows the syntax for one. Start with `tea -q` to suppress the banner; the browser edition offers the same editing and completion.

Your first tea session

Let’s load a small panel dataset, look at it, and run a regression. Save the following as `quickstart.do`:

```

* quickstart.do
import delimited "tests/WE0.csv", clear
rename INDICATOR_ID indicator_id
keep if indicator_id == "NGDP_RPCH"
rename COUNTRY_ID country_id
keep country_id v*
reshape long v, i(country_id) j(year)
rename v growth
keep if !missing(growth)

* Restrict to G7 countries for tractable output
keep if inlist(country_id, "USA", "DEU", "GBR", "FRA", "ITA", "JPN", "CAN")

* Convert string country codes to numeric so we can use factor vars
encode country_id, gen(cid)
xtset cid year

summarize growth
xtdescribe

* Fixed-effects regression of growth on its own lag
xtreg growth L.growth, fe

```

Run it:

```
tea quickstart.do
```

This loads the IMF *World Economic Outlook* CSV (shipped with the distribution), pivots it from wide to long, restricts to G7 countries, sets up the panel structure with `xtset`, and runs a fixed-effects regression. Output looks like Stata’s; if you’ve used Stata before, nothing here should surprise you.

A few things to notice:

- `rename`, `keep if`, `reshape long` all work the same as in Stata.
- `encode` maps string country IDs to consecutive integers 1, 2, 3, ... and attaches a value label so `list cid` still shows “USA”, “DEU”, etc.
- `xtset cid year` declares the panel structure. After this, `L.growth` means “growth lagged one year, gap-aware within country.”
- `xtreg ..., fe` runs fixed effects (within transformation), with classical SEs by default. Add `vce(robust)` for HC1 or `vce(cluster cid)` for clustered SEs.

The do-file language

A do-file is a sequence of *commands*, one per line, with the following syntax skeleton:

```
[prefix:] command [varlist] [if exp] [in range] [weight] [, options]
```

Comments

Four comment forms are supported:

```
* Whole-line comment starting with asterisk
// Inline comment to end of line
/* Block comment
   spanning multiple lines */
generate x = 1 ///
           + 2    // line continuation: combine the next line with this one
```

Macros

Local macros are referenced as “`name'” (back-tick ... tick); global macros as \$name.

```
local x = 5
display "x = `x'"           // x = 5
```

```
global path "/home/me/data"
use "$path/employment.dta", clear
```

```
local vars "gdp pop unemp"
foreach v of local vars {
    summarize `v'
}
```

After most estimation commands, results are stored in the `e()` namespace:

```
regress y x1 x2
display e(N)           // sample size
display e(r2)          // R-squared
display _b[x1]         // x1's coefficient
display _se[x1]        // x1's standard error
```

Summarize stores results in `r()`:

```
summarize gdp
display r(mean)
display r(sd)
```

Control flow

```
* Conditional
if `n' > 100 {
    display "large sample"
}
else {
    display "small sample"
}
```

```
* Foreach over a list, varlist, or numlist
```

```
foreach v of varlist gdp pop unemp {
    quietly summarize `v'
    display "`v': N = " r(N) ", mean = " r(mean)
}

* Forvalues
forvalues i = 1/5 {
    display "iteration `i'"
}
```

Capture and quietly

`quietly` suppresses output; `capture` additionally suppresses errors (storing the return code in `_rc`):

```
quietly regress y x           // runs, no output
display _b[x]                // coefficient is still available

capture confirm variable z // doesn't error if z doesn't exist
if _rc != 0 {
    display "z is missing"
}
```

The by: prefix

```
sort country year           // required before by:
by country: generate dy = y - y[_n-1]

bysort country (year): generate y_first = y[1] // sorts first
```

`tea` enforces the sort requirement — using `by:` on unsorted data is a hard error, not a silent miscomputation. Use `bysort` when you want it to sort automatically.

Data management

Importing data

`tea` reads CSV, TSV, Stata `.dta`, Excel `.xlsx`, and OpenDocument `.ods` files, plus its own native `.tea` binary format.

```
import delimited "data.csv", clear           // CSV
import delimited "data.tsv", clear          // TSV (auto-detected)
import excel "data.xlsx", firstrow clear
import excel "data.xlsx", sheet("Sheet2") firstrow clear

use "data.dta", clear                       // Stata format
use "data.tea", clear                       // native tea format
```

`tea` auto-detects the delimiter for `.csv` and `.tsv`; if the file extension is ambiguous, pass `delimiter(",")` or `delimiter(tab)` explicitly.

The `firstrow` option (for Excel) tells `tea` to use the first row as variable names; invalid characters are removed (`Country Code` → `CountryCode`) and a header that is empty, a duplicate, or starts with a digit (1960) falls back to the Excel column letter (F, BQ, ...) — exactly Stata's rule. Without `firstrow`, columns are named A, B, ... and row 1 is kept as a data row, again as in Stata; columns mixing a text header with numeric data become string columns.

`import delimited` follows Stata's naming rule for CSV/TSV headers: names are lowercased and invalid characters removed (`Country Code` → `countrycode`); a header that is empty, a duplicate, or starts with a digit becomes the position name `v#`. Use `case(preserve)` (or `case(upper)`) to override the case folding.

CSV import handles RFC 4180-style quoted fields, including embedded commas and newlines.

Saving data

```
save "out.dta", replace          // Stata format (default for .dta)
save "out.tea", replace         // native tea format
export delimited "out.csv", replace
export delimited "out.tsv", replace
```

Format choice. Use `.dta` if the data will be read by Stata. Use `.tea` if it's a temporary file between tea sessions (faster to read/write, no readstat overhead). Use `.csv` or `.tsv` if a non-statistical tool will consume the data.

Generating and modifying variables

```
generate gdp_log = log(gdp)
generate growth = (gdp - gdp[_n-1]) / gdp[_n-1]
replace growth = . if year == 1990           // missing for first year
```

```
generate str8 region = "WEST"
replace region = "EAST" if country_id == "CHN"
```

The `str#` prefix declares a string variable (the number is the maximum width, used by Stata when writing `.dta` files; in tea it's cosmetic on read but compressed back on write).

egen extended generators

`egen` provides aggregate functions that `generate` can't do in a single expression:

```
egen mean_gdp      = mean(gdp), by(country)
egen total_pop     = total(pop)
egen rank_gdp      = rank(gdp)
egen group_id      = group(country region)           // assigns 1..K
egen row_total     = rowtotal(x1 x2 x3 x4)
egen has_any       = anyvalue(x), values(1 5 9)
```

The full list of functions: `mean`, `sum`, `total`, `count`, `min`, `max`, `sd`, `var`, `median`, `group`, `rank`, `rowtotal`, `rowmean`, `rowmin`, `rowmax`, `rowmiss`, `rownonmiss`, `tag`, `seq`, `anyvalue`, `concat`.

Type conversion

```
* Numeric -> string
tostring year, generate(year_str)

* String -> numeric (errors if non-numeric values present)
destring score, generate(score_n)
destring score, generate(score_n) force           // converts non-numeric to .

* String -> integer codes + value label
encode country_id, generate(cid)                 // alphabetical 1..K

* Integer codes -> string (inverse of encode)
```

```
decode cid, generate(country_str)
```

```
* Recoding numeric values
```

```
recode age (0/17 = 1 "minor") (18/64 = 2 "working") (65/max = 3 "senior"), gen(age_cat)
```

Labels

```
label data "World Economic Outlook subset"
```

```
label variable growth "Real GDP growth (%)"
```

```
label define region 1 "North" 2 "South" 3 "East" 4 "West"
```

```
label values region region
```

Exploring data

status

A one-line summary of what's in memory — a tea extension (Stata has no single equivalent):

```
. status
WB_final.dta - 25,808,384 obs, 7 vars, 1.9 GB (sorted: CountryCode year; xtset: ifscode year)
```

The source name is tracked through `use / import / sysuse`, updated by `save` (as Stata's "Contains data from" is), and cleared by `clear`. The memory figure is the exact in-memory data size: 8 bytes per numeric cell, pointer plus string length per string cell.

Progress indicator

Long operations (`import`, `reshape`, `merge`) show a live progress line — another tea extension, useful because tea runs on one CPU and big jobs can otherwise look hung:

```
reshape long: 42% (10,800,000/25,808,384)
```

It is deliberately invisible everywhere it could interfere: nothing is drawn until an operation has run for about a second (short operations produce zero output), it goes to `stderr` only when `stderr` is a terminal (logs, pipes, batch runs, and `capture` never see it), and the line is erased on completion so nothing lands in scrollbar. `set progress off` disables it entirely.

describe

```
describe // all variables
describe gdp pop // just these two
describe gdp* // wildcards work
```

Shows variable name, storage type, display format, and variable label.

list

```
list // all rows (capped at 200)
list in 1/10 // first 10 rows
list country year gdp if year == 2020
```

Column widths auto-fit to the maximum of header and value widths.

summarize

```

summarize                                // all numeric variables
summarize gdp, detail                    // adds percentiles, skewness, kurtosis
summarize gdp if year >= 2000
bysort country: summarize gdp

```

After `summarize`, the following `r()` macros are set: `r(N)`, `r(mean)`, `r(sd)`, `r(min)`, `r(max)`, `r(sum)`, `r(Var)`. With `detail`, also `r(p1)`, `r(p5)`, `r(p10)`, `r(p25)`, `r(p50)`, `r(p75)`, `r(p90)`, `r(p95)`, `r(p99)`, `r(skewness)`, `r(kurtosis)`.

tabulate

```

tabulate region                          // one-way frequency table
tabulate region year, row column cell    // two-way + percentages
tabulate region, summarize(gdp) means    // means of gdp by region

```

codebook

```

codebook                                // all variables
codebook gdp                            // just one

```

Reports type, unique values, missing count, range, and distribution summary. More compact than `summarize`, `detail` for screening a new dataset.

count

```

count                                  // total observations
count if missing(gdp)
count if gdp > 0 & year >= 2000

```

After `count`, `r(N)` holds the count.

Reshaping, merging, aggregating

sort and gsort

```

sort country year                      // ascending on both
gsort -gdp                             // descending on gdp
gsort country -year                    // mixed

```

reshape

`reshape` is in-place — to keep the original you must `frame copy` (or `preserve`) first.

* Wide to long: each `yY` variable becomes `y` at year `Y`

```
reshape long y, i(country) j(year)
```

* Long to wide

```
reshape wide y, i(country) j(year)
```

* Multiple stubs

```
reshape long y z, i(country) j(year)
```

* String `j`: levels are name suffixes (long) or column-name parts (wide)

```
reshape wide y, i(country year) j(indicator_code) string
reshape long y, i(country year) j(indicator_code) string
```

Variables that are not part of a stub and not in `i()` are **carried along**, exactly as in Stata: going long they are replicated across the `j`-rows of each wide observation; going wide they must be constant within each `i()` group (error `r(9)` otherwise — drop them or add them to `i()`). Formats and variable labels survive on `i()` and carried columns.

Going wide, every generated column name (`stub + j` value) must be a valid identifier; if your `j` values contain characters like `.`, clean them first with `strtoname()`:

```
gen code_safe = strtoname(indicator_code)
reshape wide y, i(country year) j(code_safe) string
```

`reshape wide` also insists that `j` is never missing/empty and never repeats within an `i()` group — both are data errors in Stata, and `tea` fails loudly rather than guessing.

preserve / restore

`preserve` snapshots the data in memory; `restore` brings the snapshot back. One level deep, as in Stata. The snapshot lives on disk as a native `.tea` tempfile, so preserving a large dataset does not double peak memory.

```
preserve
keep indicator_code indicator_name
duplicates drop
save "lookup.dta", replace
restore
```

Variants: `restore, not` discards the snapshot without reloading; `restore, preserve` reloads but keeps the snapshot for another `restore`. A `preserve` still pending when a `do`-file concludes — normally or via an `abort` — is restored automatically, with a note (Stata behavior).

merge

```
merge 1:1 country year using "macro.dta"
merge m:1 country      using "country_attr.dta"
merge 1:m country year using "panel_history.dta"
```

After `merge`, the variable `_merge` encodes the match status: 1 = master only, 2 = using only, 3 = matched.

`m:m` is deliberately rejected. If you think you need `m:m`, you almost certainly want a different operation (`joinby`-style cross product, or an explicit deduplication before merging).

collapse

```
collapse (mean) gdp pop (sum) revenue (sd) sd_gdp = gdp, by(country)
collapse (median) gdp [aweight=pop], by(region year)
```

Available aggregators: `mean`, `median`, `sum`, `min`, `max`, `sd`, `count`, `first`, `last`, `p25`, `p50`, `p75`, `p10`, `p90`.

Frames

A frame is a named in-memory dataset. Multiple frames let you keep several datasets loaded at once.

```
frame create macro
frame change macro
use "macro.dta", clear
```

```

frame change default
* now we're back on the original frame

frame copy default backup    // duplicate default into a new frame
frame drop backup

```

Time series and panel data

Declaring panel / time structure

```

tsset year                // pure time series
xtset country year        // panel
xtset country year, yearly // explicit time unit

```

After this, time-series operators are valid:

Operator	Meaning
L.x	first lag (one period back)
L2.x	second lag
L(1/3).x	first through third lags
F.x	first lead
D.x	first difference
D2.x	second difference (Stata's notation)
S.x	seasonal difference

These can mix with factor variables (with one limitation; see the “Factor variables” chapter):

```

regress y L.y x i.region    // OLS with lag and region dummies
xtreg y L(1/3).y x, fe      // panel FE with three lags

```

tea’s time-series operators are *gap-aware*: if a country has data for 1995 but not 1996, L.y for 1997 returns missing (not the 1995 value).

xtdescribe

```
xtdescribe
```

Reports number of panels, observations per panel (min/avg/max), and sample pattern. Useful for checking whether a panel is balanced.

Factor variables

Stata’s factor-variable grammar lets you include categorical regressors and their interactions inline, without manually generating dummies. tea supports the same grammar across every estimator (`regress`, `xtreg`, `logit`, `probit`, `poisson`, `ivregress`).

Atoms

Token	Meaning
i.var	dummies for each non-base level; base = smallest level
ib2.var	same, with base level explicitly set to 2

Token	Meaning
<code>c.var</code>	continuous regressor (mostly for use inside #)

```
regress wage exper i.region          // region dummies
regress wage exper ib2.region        // base = region 2
```

`i.var` requires `var` to be numeric. For string identifiers (e.g. ISO country codes), use `encode` first:

```
encode country_id, gen(cid)
regress y x i.cid
```

Interactions

Token	Meaning
<code>A#B</code>	cross product: $(K_A - 1) \times (K_B - 1)$ columns
<code>A##B</code>	main effects + interaction (equivalent to <code>A B A#B</code>)

```
regress wage i.region#c.exper        // region-specific slopes
regress wage i.region##c.exper       // adds main effects of region and exper
regress wage i.region#i.year         // region-by-year interaction
```

The `c.` prefix tells `tea` to treat the variable as continuous in an interaction — without it, a numeric variable would be expanded as a factor.

Coefficient naming

Factor expansions produce coefficients with structured names that `predict`, `margins`, and `_b[]` all understand:

```
regress wage i.region
display _b[2.region]           // coefficient on region 2 dummy
display _b[3.region#c.exper]   // interaction coefficient
```

Known limitation: factor $\times$ TS-op composition

Tokens like `i.country#c.L.gdp` (factor interaction with a lagged regressor) are **not** supported in v1.0.

Workaround: pre-compute the lag manually and use the bare variable inside the interaction:

```
generate L1_gdp = L.gdp
regress y i.country#c.L1_gdp      // works
```

Linear regression

Basic syntax

```
regress y x1 x2 x3
regress y x1 x2, vce(robust)        // HC1 robust SEs
regress y x1 x2, vce(cluster country) // CR1 clustered SEs
regress y x1 x2 [fweight=n], vce(robust) // frequency weights
```

Weight types: `fweight` (frequency, integer), `aweight` (analytic, rescaled to sum to N), `pweight` (probability, no rescaling), `iweight` (importance, like `aweight` but no rescaling).

Output

A standard ANOVA table with F -statistic and R^2 , followed by the coefficient table with t -statistics, p -values, and 95% confidence intervals. After `regress`, the following are stored:

- `e(N)`, `e(r2)`, `e(r2_a)`, `e(rmse)`, `e(F)`, `e(p)`, `e(df_r)`, `e(df_m)`
- `_b[varname]` and `_se[varname]` for each coefficient, plus `_b[_cons]` and `_se[_cons]`.

Robust and clustered SEs

The implementation uses HC1 for robust and CR1 for clustered, matching Stata's defaults. Cluster variables may be string (panel IDs) or numeric.

Constraints and hypothesis tests

After `regress`, use `test` for joint hypotheses and `lincom` for linear combinations:

```
regress y x1 x2 x3 x4

test x1 = x2                // single restriction
test x1 = 0, x2 = 0         // joint (F-test)
test x1 x2 x3               // all = 0 jointly

lincom x1 + x2               // linear combination + CI
lincom 0.5*x1 + 0.5*x2 - x3
```

Panel regression

Fixed effects

```
xtset country year
xtreg y x1 x2, fe           // within (default for xtreg)
xtreg y x1 x2, fe vce(cluster country)
```

Output includes:

- Within / between / overall R^2
- σ_u , σ_e , ρ (variance components)
- F -test of $u_i = 0$
- `corr(u_i, Xb)` — informally measures the FE–RE departure

Identifies as `e(cmd) = "xtreg"` with `e(model) = "fe"`.

Random effects

```
xtreg y x1 x2, re           // Swamy-Arora FGLS
```

Implementation: Swamy-Arora variant (the Stata default). Two-step FGLS — first estimate σ_u^2 , σ_e^2 from the within and between residuals; then quasi-demean using $\theta = 1 - \sigma_e / \sqrt{T\sigma_u^2 + \sigma_e^2}$.

Between effects

```
xtreg y x1 x2, be           // OLS on panel means
```

Collapses the data to one observation per panel (the panel means of y and each x), then runs OLS. This is the BE estimator that underpins one of the components of the RE-FE decomposition.

Hausman test

```
quietly xtreg y x1 x2, fe           // populates fe_est
quietly xtreg y x1 x2, re           // populates re_est
hausman                             // tests FE vs RE
```

Order of `xtreg fe` and `xtreg re` doesn't matter; `tea` keeps both in dedicated workspace slots and the `hausman` command (no arguments needed) tests their coefficient difference.

Maximum likelihood estimators

`tea` implements logit, probit, and Poisson regression via a shared Newton-Raphson driver. All three support `if`, `in`, weights, and `vce(robust|cluster)`.

logit

```
logit voted income education
logit y x1 x2 i.region, vce(cluster country)
```

Reports the maximum likelihood iteration trace, LR χ^2 test against the constant-only model, pseudo R^2 , and the coefficient table with z -statistics.

probit

Same syntax, normal CDF instead of logistic CDF:

```
probit y x1 x2 i.region
```

poisson

For count outcomes:

```
poisson accidents speed_limit population, vce(cluster state)
```

The link is $\log E[y | x] = x'\beta$; coefficients are log-rate-ratios. $y < 0$ is rejected. Non-integer y is permitted (treated as the GLM extension).

Important: `tea`'s Poisson drops the $\log(y!)$ constant in the log-likelihood, so the absolute `e(11)` value differs from Stata's, but the coefficient estimates, SEs, and likelihood-ratio tests are identical.

Marginal effects via margins

After any of the above, `margins` produces average or at-the-mean marginal effects:

```
logit y x1 x2 i.region
margins, dydx(*)           // average marginal effects
margins, dydx(x1)          // just x1
margins, dydx(*) atmeans   // at-the-means MEs
```

Standard errors use the delta method.

Instrumental variables

`ivregress 2sls` implements two-stage least squares with a first-stage F diagnostic for weak instruments.

Syntax

```
ivregress 2sls y [exog_vars] (endog_vars = instruments) [if] [in] [weight] [, options]
```

The parenthesized group is the key piece: variables to the left of the = are endogenous; variables to the right are excluded instruments.

Example

```
* Wage equation: education is endogenous, parents' education is the IV
ivregress 2sls wage exper exper2 (educ = momeduc dadeduc), vce(robust)
```

Output includes the coefficient table (with z -statistics from asymptotic normality), the list of instrumented variables, the list of instruments (included exogenous + excluded), and the first-stage F -statistic for the weakest endogenous regressor.

Interpreting the first-stage F

A common rule of thumb (Stock-Yogo): if the first-stage F is below 10, the instrument is weak and the IV estimates may have substantial finite-sample bias. `tea` flags this in the output with a caret.

Variance estimators

```
ivregress 2sls y (x = z) // classical
ivregress 2sls y (x = z), vce(robust) // HC1
ivregress 2sls y (x = z), vce(cluster country) // CR1
```

The sandwich form is $V = A \cdot B \cdot A$ where $A = (\hat{X}'\hat{X})^{-1}$ (using fitted endogenous columns) and B uses residuals computed from the *original* X (this is the correct 2SLS variance — not the naive OLS-on-fitted variance).

Time series models

ARIMA

```
tsset year
arima y, arima(p d q) [noconstant]
arima y x1 x2, arima(1 0 1) // ARMAX with two exog
```

Specifies an ARIMA(p, d, q) model: AR order p , difference order d , MA order q . Exogenous regressors can be included (then the model is ARMAX).

Important. v1.0's `arima` uses *conditional likelihood*, not the Kalman-filter exact ML that Stata uses by default. The differences:

- Coefficient estimates are consistent and converge to the same limit, but finite-sample values can differ by a few percent.
- The absolute log-likelihood differs (we drop the initial-state contribution).
- Standard errors come from a numerical-difference Hessian and can be off by 1–5% on the MA components specifically. Point estimates are unaffected.

If you need exact ML, escape to R's `arima()` or Python's `statsmodels.tsa.arima`.

Example

```
* US growth ARIMA(1,0,0)
tsset year
arima growth, arima(1 0 0)
```

The output groups coefficients into “ARMA model” (constant + exog), “AR” (ϕ_i), “MA” (θ_j), and `/sigma` (the innovation standard deviation).

Post-estimation

predict

After any estimator, `predict` generates a new variable with the model prediction or a residual:

```
* After regress
regress y x1 x2
predict yhat                // linear prediction (default: xb)
predict resid, residuals    // residuals
predict resid, r            // same, shorter

* After logit/probit
logit voted income
predict p                   // Pr(y=1)
predict xb, xb              // linear index

* After xtreg fe
xtreg y x1, fe
predict u, u                // estimated u_i
predict e, e                // idiosyncratic residual
predict ue, ue              // u_i + e_it
```

Hypothesis tests

`test` performs Wald tests on linear restrictions:

```
regress y x1 x2 x3
test x1 = 0                // single restriction
test x1 = 1, x2 = 0        // joint restriction
test x1 x2 x3              // x1 = x2 = x3 = 0
```

`lincom` computes linear combinations with their SEs and 95% CIs:

```
lincom x1 + x2 - x3
lincom 0.5*x1 + 0.5*x2
```

Output and reproducibility

Saving estimates

The `estimates` (alias `est`) command suite saves and manages named estimation results:

```
regress y x1
estimates store m1
```

```
regress y x1 x2
```

```

estimates store m2

regress y x1 x2 x3
estimates store m3

estimates dir                // list stored estimates
estimates table m1 m2 m3, stats(N r2 rmse) star
estimates restore m1         // load m1 back to last_est
estimates drop m2            // remove m2
estimates drop _all          // remove all

```

estout: publication tables

For papers, use `estout` to produce LaTeX, markdown, or plain side-by-side tables:

```

estout m1 m2 m3                // default: LaTeX
estout m1 m2 m3, format(latex) stats(N r2 rmse)
estout m1 m2 m3, format(markdown) se star
estout m1 m2 m3, format(plain) t

estout m1 m2 m3 using "table.tex", stats(N r2) // write to file
estout m1 m2 m3, title("Regression results") label(tab:main)

* Subset of coefficients
estout m1 m2, keep(x1 x2 _cons) stats(N r2)
estout m1 m2, drop(_cons)                // drop constant row

```

Options:

Option	Effect
<code>format(latex markdown plain)</code>	output format (default: <code>latex</code>)
<code>se, t, p</code>	show SEs / <i>t</i> -stats / <i>p</i> -values below coefs
<code>star</code>	add significance stars (10%, 5%, 1%)
<code>stats(N r2 r2_a rmse F ...)</code>	footer rows
<code>title(...), label(...)</code>	LaTeX caption + reference label
<code>keep(...), drop(...)</code>	restrict coefficient rows
<code>nomtitles, mtitles(...)</code>	override per-model column labels
<code>using FILE</code>	write to file instead of stdout

LaTeX output produces a stand-alone `tabular` block you can `\input` in your paper.

log

```

log using "session.log", replace
* commands and their output are captured...
log close

```

Captures both the commands you ran (echoed with a `.` prefix) and the output they produced.

Random numbers and Monte Carlo

Reproducibility

```
set seed 12345                                // fix the PRNG state
```

tea uses a 64-bit PCG generator. `set seed N` is deterministic — the same seed always produces the same sequence, across runs, machines, and tea versions.

Functions

```
generate u = runiform()                        // Uniform[0, 1)
generate z = rnormal()                        // Standard normal
generate w = rnormal(5, 2)                    // Normal(mean=5, sd=2)
```

Normal draws use Box-Muller with caching. All three functions are available in any expression context (`generate`, `replace`, `egen rowfirst/rowlast`, etc.).

A Monte Carlo example

```
* OLS slope sampling distribution under known DGP
set seed 42
clear
set obs 1000

* True model: y = 2*x + eps,  eps ~ N(0, 1)
generate x = rnormal()
generate eps = rnormal()
generate y = 2*x + eps

regress y x
display "Estimated beta = " _b[x] " (true is 2.0)"
```

Graphics

Three commands cover the everyday cases, rendered by tea's own dependency-free SVG engine:

```
scatter y x    [if] [in] [, title() xtitle() ytitle() saving(FILE) noview]
line   y x    [if] [in] [, sort ...same options...]
histogram v    [if] [in] [, bins(#) freq ...same options...]
```

Output is a vector SVG — publication quality by default, and identical byte-for-byte across platforms (plots are covered by golden-file regression tests). In the interactive REPL the plot opens in your OS viewer; `saving(FILE)` writes to a named file instead, and `noview` suppresses the viewer. In do-files plots are only ever written to files, never opened, so batch output stays deterministic.

`line` connects points in data order; add `sort` to order by x. `histogram` shows density by default (`freq` for counts); automatic binning is `min(ceil(sqrt(N)), 50)`. In the browser edition plots appear in the panel beside the terminal.

A worked example on bundled data:

```
sysuse grunfeld
scatter invest value, title("Investment vs firm value") ///
      xtitle("Market value") ytitle("Gross investment") saving(grun.svg)
```

SVG drops directly into LaTeX documents (via the `svg` package, or one `rsvg-convert` step to PDF).

Bundled practice datasets (sysuse)

Six datasets are embedded inside the binary itself — no files to install, no network, identical in the browser edition:

```
. sysuse dir
  grunfeld  Grunfeld (1958) investment panel: 10 US firms x 1935-1954 (xtreg, hausman)
  airline   Box-Jenkins airline passengers, monthly 1949-1960 (tsset, arima)
  longley   Longley (1967) US macro, 16 obs: famously ill-conditioned (regress)
  nmes1988  Deb-Trivedi (1997) medical care demand, 4406 persons 66+ (poisson, logit)
  pwt       Penn World Table 10.0 sample: 22 countries x 1950-2019, CC BY 4.0 (growth)
  weo       IMF World Economic Outlook, April 2026: 197 countries + 13 aggregates, 145 indicators
```

Load one with `sysuse NAME[, clear]` — the same unsaved-data guard as `use` applies.

The WEO database

`weo` is the complete published IMF World Economic Outlook database, April 2026 vintage, as a long panel: `country iso year aggregate` plus one variable per WEO subject code, lowercased. The codes are the IMF's own unambiguous names — `ngdp_rpch` is real GDP growth, `pcpipch` CPI inflation (period average), `lur` the unemployment rate, `bca_ngdpd` the current account in percent of GDP, `ggxwdg_ngdp` general government gross debt in percent of GDP, and so on. The full code → descriptor → units table ships as `data/weo_codes.txt`.

World and regional groups (World, Advanced Economies, G7, Euro Area, EU, Emerging Market and Developing Economies, regional aggregates) carry `aggregate==1`; countries carry `aggregate==0`:

```
sysuse weo
keep if aggregate==0          // the 197-country panel
xtset iso year
xtreg ngdp_rpch ggxwdg_ngdp if year<=2025, fe

sysuse weo, clear
keep if aggregate==1          // the 13 groups
line ngdp_rpch year if iso=="G001", sort title("World real GDP growth")
```

101 of the 145 subject codes are published for groups only and are empty on country rows — faithful to how the IMF publishes the database. Years through 2031 are the projections *as published in this vintage*: the bundled copy is a practice snapshot, not a live data service. Cite as *IMF, World Economic Outlook Database, April 2026*.

The other five

`grunfeld` is the canonical panel-methods teaching dataset; `xtreg invest value capital, fe` reproduces the textbook estimates (0.110, 0.310). `airline` is the Box–Jenkins monthly series every ARIMA text uses. `longley` is the famously ill-conditioned regression used to certify least-squares implementations — `regress employed gnpdeflator gnp unemployed armedforces population year` reproduces the certified solution. `nmes1988` is Deb and Trivedi's (1997, *JAE*) medical-care demand microdata: physician-visit counts, health status, insurance and income for 4,406 persons 66+ — real material for `poisson` and `logit`. `pwt` is a Penn World Table 10.0 sample (CC BY 4.0) for growth regressions.

Provenance, licenses, and citations for all six are in `data/SOURCES.md`. To refresh the embedded data (for example a newer WEO vintage): run `tools/curate_weo.py` on the new download, then `tools/gen_sysdata.py`, and rebuild — the curator auto-detects both the classic bulk-file and the data-portal export formats.

The browser edition

The WebAssembly build at <https://micomrkaic.github.io/tea/> is the same program compiled for the browser: same commands, same estimators, same bundled data, same output to the last byte — the full regression suite runs inside the browser build and passes identically. Nothing you load or type leaves your machine; there is no server side.

Practical differences from the native binary:

- **Files** live in an in-browser filesystem. The *Upload data file* button places files where `use / import delimited` can see them; *Download workspace files* retrieves anything you `save` or `export`.
- **Plots** render in the panel beside the terminal instead of opening an OS viewer.
- **shell** is unavailable (browsers have no processes) and fails with a clear message.
- The terminal offers the same readline-style editing and Tab completion as the native REPL.

The first visit downloads about 2 MB (compressed); the browser caches it afterwards. The browser edition is a demonstration and practice channel — for production work, build the native binary.

Backend-independent output

tea promises that **the same do-file prints byte-identical output on every machine, operating system, CPU, and BLAS library**. For well-posed computations this is automatic — results agree to displayed precision everywhere. The danger zone is *mathematical zeros*: the residual sum of squares of a perfect fit, or the standard error of an exactly-determined coefficient, which different numerical backends render as different sub-epsilon noise ($5.9\text{e-}29$ on one machine, $4.1\text{e-}29$ on another), occasionally flipping what gets printed entirely (an SE of 0 versus $1\text{e-}16$ turns `t = 0.00`, `p = 1.000` into `t = 1e+16`, `p = 0.000`).

tea therefore snaps such quantities to exactly zero when they fall below tight *relative* tolerances: residual SS under $1\text{e-}12$ of total SS, and the consequences that follow — residual vectors, standard errors, coefficients under $1\text{e-}8$ of the largest in an exact fit, `sigma_u` when panel intercepts are identical, and the `hausman` difference matrix at machine precision. A perfect fit reports `F = inf`, `p = 0.0000`, deterministically. **Normal estimation results are untouched to the last bit** — the thresholds fire only on genuine degeneracy, and a quantity that is merely small (rather than a mathematical zero) is left alone by the relative tests.

This is verified, not asserted: the regression suite passes byte-identically under native gcc + OpenBLAS (release and sanitizer builds) and under WebAssembly clang + reference BLAS + musl — two maximally different numerical stacks. The tradeoff, stated plainly: raw sub-epsilon values in degenerate fits are not displayed. The reasoning is documented in the README’s “Semantics decisions” and in this manual so users of published results can account for it.

Escape hatches: when to leave tea

tea is excellent at the data-prep-to-quick-estimate loop, but many things are out of scope. The escape pattern is:

1. Do data prep in tea.
2. `export delimited` to CSV.
3. `shell python myscript.py` (or R, or anything).
4. `import delimited` the results back.

When to escape

What you need	Where to go
Plotting	gnuplot, Python's matplotlib, R's ggplot2
GARCH / EGARCH / TGARCH	Python arch, R rugarch
Exact-ML ARIMA (Kalman)	R's arima(), forecast::Arima
Seasonal ARIMA	R's arima() with seasonal arg
Bayesian inference	PyMC, Stan (rstan, cmdstanpy)
Mixed-effects models	R's lme4, Python statsmodels.MixedLM
Survival analysis	R's survival, Python lifelines
Cointegration / VAR / VECM	R's urca, vars
Machine learning	scikit-learn, xgboost, lightgbm
Network / graph models	networkx, R igraph
Spatial econometrics	R's spdep, Python's pysal

A bridge example

* Estimate panel FE in tea, then use Python for a residuals plot

```
xtset country year
xtreg gdp_growth L.gdp_growth inflation, fe
predict resid, e
```

* Export for plotting

```
keep country year resid
export delimited "resid.csv", replace
```

* Hand off to Python (assuming plot_residuals.py exists)

```
shell python plot_residuals.py resid.csv resid_plot.png
```

Known limitations

These are not bugs in the strict sense — they're scope decisions for v1.0 that we may revisit based on testing.

Data and metadata

- **.dta round-trip does not preserve xtset state.** After save mydata.dta then use mydata.dta, the panel/time variable assignment is lost. Workaround: re-run `xtset country year` after use. This is the standard Stata-do-file convention anyway.
- **Variable name length cap of 32 characters.** Factor-variable interaction names like `2010.year#1985.country#c.gdp` (28 chars) fit; longer 3-way interactions truncate.

Econometric estimators

- **ARIMA uses conditional likelihood, not Kalman-filter exact ML.** See the “Time series models” chapter for details. For serious work, use R's `arima()`.
- **No seasonal ARIMA** (`sar`, `sma` terms).
- **ARIMA MA standard errors** have a 1–5% precision loss from the numerical-Hessian computation. Point estimates are unaffected.

- **xtreg, be uses straight OLS** on panel means. Stata’s default applies a slight variance correction for unbalanced panels (down-weighting groups with few observations). Differences appear only on unbalanced panels and only in the SEs (point estimates are identical).
- **No xtabond / system GMM** or other GMM estimators.
- **No VAR / VECM.** Escape to R.

Factor variables

- **No composition with TS operators.** Tokens like `i.country#c.L.gdp` are not parsed. Workaround: pre-compute the lag with `gen L1_gdp = L.gdp` and use `c.L1_gdp`.
- **i.var requires numeric var.** Use `encode` on string identifiers.
- **Cell-count cap of 10 000.** Larger factor expansions are blocked.

Programming

- **No program define, no syntax parser, no Mata.** You can write loops and use macros, but you cannot define new commands.

Graphics

`scatter`, `line`, and `histogram` cover the everyday cases (see the Graphics chapter). There is no `twoway` grammar, no faceting/`by()` graphics, no legends or multiple series per plot, and no graph editor; for figures beyond the basics, `export delimited` and pipe to Python / R / gnuplot.

Reporting bugs

If you find a bug:

1. Distill it into a small reproducer — a do-file plus a small input file is ideal. If the input data is sensitive, replace it with `set seed N; gen x = rnormal()` or similar synthetic.
2. Note your `tea --version`, OS, and compiler version.
3. Send to the author at the email associated with your distribution of `tea`. If you don’t have one, ask the person who gave you the software.

What counts as a bug:

- Wrong numerical answer (with hand-calculated expected value).
- Crash, segfault, or hang on input that should work.
- Different behavior from documented Stata semantics (when not in the “known limitations” list above).

What does not count as a bug:

- Missing feature (those go on the feature-request list).
- Disagreement with Stata’s exact numerical output to many decimal places (small differences arise from BLAS vendor variation and are not correctness issues).
- Things listed in the “Known limitations” chapter.

Function reference

Functions available in any expression context (**generate**, **replace**, **if** clauses, **display**, ...). Each entry shows the signature and a one-line description. In all cases: a missing argument propagates to a missing result, unless otherwise stated.

Arithmetic and rounding

Signature	Returns
<code>abs(x)</code>	$ x $
<code>int(x)</code>	truncate toward zero
<code>floor(x)</code>	largest integer $\leq x$
<code>ceil(x)</code>	smallest integer $\geq x$
<code>round(x)</code>	round to nearest integer (half away from zero)
<code>round(x, unit)</code>	round to nearest multiple of unit
<code>sign(x)</code>	$-1 / 0 / +1$
<code>mod(x, y)</code>	$x \bmod y$, same sign as x
<code>min(x, y, ...)</code>	minimum of arguments (missing skipped)
<code>max(x, y, ...)</code>	maximum of arguments (missing skipped)
<code>cond(test, a, b)</code>	a if test is true, else b

Exponential and logarithm

Signature	Returns
<code>exp(x)</code>	e^x
<code>ln(x)</code>	natural log; missing if $x \leq 0$
<code>log(x)</code>	alias for <code>ln</code>
<code>log10(x)</code>	base-10 log; missing if $x \leq 0$
<code>sqrt(x)</code>	$\sqrt{x}$; missing if $x < 0$

Trigonometric and hyperbolic

Signature	Returns
<code>sin(x), cos(x), tan(x)</code>	radians
<code>asin(x), acos(x)</code>	inverse trig; range checked
<code>atan(x)</code>	inverse tangent
<code>atan2(y, x)</code>	two-argument inverse tangent
<code>sinh(x), cosh(x), tanh(x)</code>	hyperbolic

Probability distributions

Signature	Returns
<code>normal(z)</code>	standard-normal CDF $\Phi(z)$
<code>normalden(z)</code>	standard-normal PDF $\phi(z)$
<code>normalden(z, mu, sd)</code>	normal PDF with mean and SD
<code>lnnormal(z)</code>	$\log \Phi(z)$, stable in the lower tail

Signature	Returns
<code>lnnrmalden(z)</code>	$\log \phi(z)$, stable in extremes
<code>invnormal(p)</code>	inverse CDF $\Phi^{-1}(p)$
<code>ttail(df, t)</code>	upper-tail Student- t : $P(T > t)$
<code>invttail(df, p)</code>	inverse: t such that $P(T > t) = p$
<code>F(df1, df2, F)</code>	lower-tail F CDF
<code>Ftail(df1, df2, F)</code>	upper-tail F : $P(F > f)$
<code>chi2(df, x)</code>	lower-tail χ^2 CDF
<code>chi2tail(df, x)</code>	upper-tail χ^2 : $P(X > x)$

Random number generation

Reproducible from `set seed N`. See the “Random numbers and Monte Carlo” chapter.

Signature	Returns
<code>runiform()</code>	uniform $[0, 1)$
<code>rnormal()</code>	standard normal
<code>rnormal(mu)</code>	normal with mean <code>mu</code> , sd 1
<code>rnormal(mu, sd)</code>	normal with mean <code>mu</code> and SD <code>sd</code>

Missing-value tests and ranges

Signature	Returns
<code>missing(x)</code>	1 if <code>x</code> is any missing code, else 0
<code>inrange(x, lo, hi)</code>	1 if $lo \leq x \leq hi$
<code>inlist(x, v1, v2, ...)</code>	1 if <code>x</code> equals any of <code>v1</code> , <code>v2</code> , ...

String construction and inspection

Signature	Returns
<code>strlen(s), length(s)</code>	character count of <code>s</code>
<code>substr(s, start, len)</code>	substring; <code>start</code> is 1-based; -1 for last
<code>strpos(s, sub)</code>	1-based index of <code>sub</code> in <code>s</code> , 0 if absent
<code>strrpos(s, sub)</code>	like <code>strpos</code> but search from the right
<code>word(s, n)</code>	the n th whitespace-separated word
<code>wordcount(s)</code>	number of whitespace-separated words
<code>abbrev(s, n)</code>	abbreviate <code>s</code> to width <code>n</code>
<code>char(n)</code>	single-character string for ASCII code <code>n</code>

String transformation

Signature	Returns
<code>strupper(s), upper(s)</code>	uppercase
<code>strlower(s), lower(s)</code>	lowercase
<code>strproper(s), proper(s)</code>	title case (first letter of each word)

Signature	Returns
<code>strtrim(s), trim(s)</code>	strip leading and trailing whitespace
<code>strtoname(s)</code>	make a valid variable name: invalid chars to <code>_</code> , leading digit gets <code>_</code> prefix ("CC.EST" -> CC_EST)
<code>ltrim(s), rtrim(s)</code>	strip leading / trailing whitespace
<code>itrim(s)</code>	collapse internal whitespace runs to one space
<code>strreverse(s)</code>	reverse character order
<code>substr(s, from, to)</code>	replace all occurrences of <code>from</code> with <code>to</code>
<code>substr(s, from, to, n)</code>	replace at most <code>n</code> occurrences
<code>subinword(s, from, to)</code>	replace whole-word matches only
<code>strmatch(s, pattern)</code>	glob-style match (<code>*</code> , <code>?</code>); 1 / 0

Regular expressions (POSIX extended)

Signature	Returns
<code>regexm(s, pat)</code>	1 if <code>pat</code> matches anywhere in <code>s</code> , else 0; populates <code>regexs()</code>
<code>regexs(n)</code>	<i>n</i> th captured group from last <code>regexm</code> (0 = full match)
<code>regrex(s, pat, repl)</code>	replace first match in <code>s</code> ; <code>repl</code> may use <code>\1 ... \9</code>

Numeric ↔ string conversion

Signature	Returns
<code>string(x)</code>	default decimal string
<code>string(x, fmt)</code>	format <code>x</code> with the given Stata format
<code>stofreal(x), stofreal(x, fmt)</code>	alias for <code>string</code>
<code>real(s)</code>	parse <code>s</code> as numeric; missing on failure

Calendar: building dates

Signature	Returns
<code>mdy(month, day, year)</code>	daily date (%td) for (month, day, year)
<code>ym(year, month)</code>	monthly date (%tm)
<code>yq(year, quarter)</code>	quarterly date (%tq)
<code>yh(year, half)</code>	half-yearly date (%th)
<code>yw(year, week)</code>	weekly date (%tw)
<code>date(s, fmt)</code>	parse string <code>s</code> per <code>fmt</code> ("YMD", "DMY", "MDY", etc.)

Calendar: extracting components from a daily date

Signature	Returns
<code>year(d)</code>	calendar year
<code>month(d)</code>	month, 1–12
<code>day(d)</code>	day of month, 1–31
<code>quarter(d)</code>	quarter, 1–4
<code>halfyear(d)</code>	half-year, 1 or 2

Signature	Returns
<code>week(d)</code>	ISO week number
<code>dow(d)</code>	day of week, 0 (Sun) – 6 (Sat)
<code>doy(d)</code>	day of year, 1–366

Calendar: converting between frequencies

Signature	Returns
<code>mofd(d)</code>	monthly date corresponding to daily date <code>d</code>
<code>dofm(m)</code>	first day of month <code>m</code> (as daily)
<code>qofd(d)</code>	quarterly index of daily <code>d</code>
<code>dofq(q)</code>	first day of quarter <code>q</code>
<code>hofd(d)</code>	half-year index of daily <code>d</code>
<code>dofh(h)</code>	first day of half-year <code>h</code>
<code>wofd(d)</code>	weekly index of daily <code>d</code>
<code>ty(y)</code>	year-as-yearly-date (just <code>y</code> , for display)
<code>dofy(y)</code>	first day of year <code>y</code>
<code>tm, tq, tw, th, td</code>	identity converters used for format-tagging

Special: running and time-series-context

Signature	Returns
<code>sum(x)</code>	running cumulative sum (resets per by-group)
<code>_n</code>	current observation index (1-based)
<code>_N</code>	total observation count in current scope
<code>x[_n-1], x[_n+1]</code>	lag / lead within current sort order
<code>x[1], x[_N]</code>	first / last observation of current scope

egen function reference

Functions usable as `egen newvar = func(...)`; see the **egen** part of the Data-management chapter.

Group-wise aggregators

When combined with `by(grouplist)`, these compute one value per group and broadcast that value back to every row in the group. Without `by()`, they collapse to a single scalar broadcast everywhere.

Signature	Returns
<code>mean(exp)</code>	arithmetic mean
<code>sum(exp), total(exp)</code>	sum (missing values skipped)
<code>count(exp)</code>	number of non-missing observations
<code>min(exp), max(exp)</code>	extrema
<code>sd(exp)</code>	sample standard deviation, $n - 1$ denominator
<code>var(exp)</code>	sample variance
<code>median(exp)</code>	50th percentile (linear interpolation)
<code>iqr(exp)</code>	interquartile range, 75th – 25th
<code>pc(exp)</code>	percentile; option <code>p(N)</code> sets which (default 50)

Signature	Returns
-----------	---------

Position and identity

Signature	Returns
<code>rank(exp)</code>	1-based rank within group; ties handled by mean rank
<code>group(varlist)</code>	1...K integer ID per unique combination of <code>varlist</code>
<code>tag(varlist)</code>	1 on the first row of each <code>varlist</code> group, 0 otherwise

Row-wise aggregators (across columns of a single row)

These ignore `by()` (the row is the unit).

Signature	Returns
<code>rowtotal(varlist)</code>	sum across columns; missing skipped
<code>rowmean(varlist)</code>	mean across columns
<code>rowmin(varlist)</code>	minimum across columns
<code>rowmax(varlist)</code>	maximum across columns
<code>rowsd(varlist)</code>	sample SD across columns
<code>rowmiss(varlist)</code>	count of missing cells in the row
<code>rownonmiss(varlist)</code>	count of non-missing cells

Summarize / tabstat statistics keywords

Available with `tabstat varlist, statistics(...)`, and many also as percentile selectors after `summarize, detail`.

Keyword	Meaning
<code>n</code>	non-missing count
<code>mean</code>	arithmetic mean
<code>sum</code>	total
<code>sd</code>	standard deviation (sample, $n - 1$)
<code>variance</code>	variance
<code>min, max</code>	extrema
<code>range</code>	<code>max - min</code>
<code>median</code>	50th percentile
<code>p1, p5, p10, p25, p50</code>	percentiles
<code>p75, p90, p95, p99</code>	percentiles
<code>iqr</code>	75th – 25th percentile
<code>skewness, kurtosis</code>	after <code>summarize, detail</code>

Stata cheat sheet

A quick reference for Stata users adapting to `tea`.

Same as Stata

- Syntax: `[prefix:] command [varlist] [if] [in] [weight] [, options]`
- Missing-value algebra, including `if x > 5` including missing
- Macros: ``local'` and `$global`, `r()` and `e()` returns
- `by:`, `bysort:`, the `[_n-1]` subscript, `_N`
- Value labels: `label define / values / list`
- Factor variables (`i.`, `c.`, `ib<n>.`, `#`, `##`)
- Time-series operators (`L.` `F.` `D.` `S.`)

Tea-specific differences

- `m:m` merge is rejected with an explanatory message.
- `xtreg, be` uses simple OLS on panel means.
- `arima` uses conditional likelihood.
- `i.country#c.L.gdp` doesn't parse; pre-compute the lag.
- No `.gph` graphics — use `export` and external tools.
- `program` `define`, `syntax`, and `Mata` are absent.

Stata commands that don't exist in tea

- Graphics: `graph`, `twoway`, `histogram`, etc.
- GMM: `gmm`, `xtabond`, `xtdpd`.
- Time series: `var`, `vec`, `dfgls`, `dfuller`, `xcorr`, `wntestq`.
- Survival: `stcox`, `streg`, `stset`.
- Mixed: `mixed`, `xtmixed`, `meqrlogit`.
- Bayesian: `bayes:`, `bayesmh`.
- Multivariate: `mvreg`, `factor`, `pca`.
- Multiple imputation: `mi`.

For each of these, see the “Escape hatches” chapter for the recommended external tool.

Command reference

Generated from the tea binary by `tools/gen_cmdref.sh` — this text is exactly what `help CMD` prints at the prompt, and regenerating it after any change keeps the manual and the implementation in agreement.

scatter

```
scatter yvar xvar [if] [in] [, title() xtitle() ytitle() saving() noview]
    SVG scatter plot; saving(FILE) writes to FILE instead of tea_graph.svg
    e.g.  scatter gdp_growth inflation if year>2000, title("Growth vs inflation")
```

line

```
line yvar xvar [if] [in] [, sort title() xtitle() ytitle() saving() noview]
    SVG line plot; connects points in data order (use sort to order by x)
    e.g. line gdp year if country=="US", sort
```

histogram

```
histogram var [if] [in] [, bins(#) freq title() saving() noview]
    SVG histogram; density by default, freq for counts, auto bins = min(ceil(sqrt(N)),50)
    e.g. histogram wage, bins(30) freq
```

generate

```
generate [type] newvar = exp [if] [in]          create a new variable
    e.g. gen logy = log(income) if year >= 2000
```

replace

```
replace var = exp [if] [in]                    modify existing variable
    e.g. replace income = . if income < 0
```

egen

```
egen newvar = fn(args) [, by(varlist)]          extended generate (mean/sum/group/...)
    e.g. egen meanY = mean(y), by(country)
```

list

```
list [varlist] [if] [in]                      print observations
    e.g. list country year gdp if year>2020 in 1/10
```

summarize

```
summarize [varlist] [if] [in] [weight] [, detail]      N/mean/sd/min/max; sets r()
    e.g. summarize gdp_growth, detail
```

tabstat

```
tabstat varlist [if] [in], statistics(stat ...) [by(var)] [columns(stat|var)]
    stats: mean sd var cv min max range sum count median p1-p99 iqr
    e.g. tabstat gdp pop, stats(mean sd p25 p50 p75) by(region)
```

count

```
count [if] [in] [fw=]                        count matching observations
```

describe

```
describe [varlist]                          show variables, types, formats
```

drop

```
drop varlist | if exp | in range                remove vars or observations
    e.g. drop if gdp == . | drop tempvar1 tempvar2
```

keep

keep varlist | if exp | in range complement of drop
 e.g. keep country year gdp

rename

rename oldname newname rename a variable
 e.g. rename _C* C* (wildcard rename: strip leading underscore)

order

order varlist reorder variables (listed ones go first)
 e.g. order country year

aorder

aorder [varlist] alphabetize variable order

label

label variable v "text" attach a description to a variable
 label define lblname # "text" # "text" define a value-label set
 label values varlist lblname attach value labels to variables

format

format varlist %fmt set display format
 e.g. format date %td | format gdp %12.2f

sort

sort varlist ascending stable sort
 e.g. sort country year

gsort

gsort [+/-]var [+/-]var ... sort with per-key direction
 e.g. gsort -gdp +country (gdp desc, then country asc)

tabulate

tabulate var [if] [in] [weight] one-way frequency table
 e.g. tabulate region if year == 2020

tsset

tsset timevar [, format(%fmt)] declare time series
 e.g. tsset year

xtset

xtset panelvar timevar declare panel (gap-aware L./F./D./S.)
 e.g. xtset country year

xtdescribe

xtdescribe summarize panel structure
 e.g. xtdescribe (requires xtset first; shows n, T, balance)

xtreg

xtreg y x1 x2 ... [if] [in], fe [vce(robust|cluster v)] panel FE OLS
 e.g. xtreg growth L.growth pop, fe
 (requires xtset; within estimator; reports R-w/b/o, sigma_u/e, rho)

hausman

hausman FE vs RE specification test
 e.g. xtreg y x, fe
 xtreg y x, re
 hausman (Ho: cov(u_i, x) = 0; rejection means use FE)

logit

logit y x1 x2 ... [if] [in] [weight], [vce(robust|cluster v)] logistic regression
 e.g. logit voted income education, vce(cluster county)

probit

probit y x1 x2 ... [if] [in] [weight], [vce(robust|cluster v)] probit regression
 e.g. probit voted income education, vce(robust)

ivregress

ivregress 2sls y [exog] (endog = instruments) [if] [in] [weight] [, vce(robust|cluster v)]
 Two-stage least squares with first-stage F diagnostic.
 e.g. ivregress 2sls wage educ (exper = age momeduc daddyEduc)

poisson

poisson y x1 x2 ... [if] [in] [weight] [, vce(robust|cluster v)] poisson regression
 e.g. poisson accidents speed_limit population, vce(cluster state)

arima

arima y [exog] [if] [in], arima(p d q) [noconstant] ARIMA(p,d,q) via conditional ML
 e.g. arima gdp, arima(1 1 1) ARIMA(1,1,1) on gdp
 arima cpi unemp, arima(2 0 1) ARMAX(2,1) with exog regressor

margins

margins , dydx(*|varlist) [atmeans] average marginal effects
 e.g. logit y x1 x2
 margins, dydx(*) AME for all regressors
 margins, dydx(x1) atmeans MEM for x1 at sample means

estimates

```
estimates store|restore|dir|drop|table      named estimates ledger
e.g. regress y x1 x2
     estimates store m1
     regress y x1 x2 x3
     estimates store m2
     estimates table m1 m2, se star stats(N r2 rmse)
```

est

```
est = estimates (short form)                see `help estimates`
```

estout

```
estout [names] [, format(latex|markdown|plain) se|t|p stars stats(...)]
Side-by-side LaTeX/markdown/plain table of stored estimates.
e.g. estout m1 m2, stats(N r2 rmse) using table.tex
```

collapse

```
collapse (stat) v1 v2 ... , by(g) [weight]   aggregate to groups
e.g. collapse (mean) gdp pop, by(region year)
```

merge

```
merge 1:1|m:1|1:m keyvars using FILE [, ...] join master with using file
e.g. merge 1:1 country year using gdp.tea (m:m is rejected by design)
```

reshape

```
reshape long|wide stubs , i(idvars) j(jvar) pivot wide<->long (in-place)
e.g. reshape long v, i(country) j(year)
     takes v1980 v1981 ... v2031 -> rows of (country, year, v)
```

recode

```
recode varlist (rules) [, gen(stub)]          map values
e.g. recode score (1/3=1)(4/6=2)(7/10=3) (missing=.), gen(score_g)
```

encode

```
encode strvar, generate(newvar) [label(name)]  map strings to integers + value label
```

decode

```
decode intvar, generate(newvar)                inverse of encode (codes back to strings)
```

destring

```
destring varlist, {replace|generate(stub)} [force] [ignore("chars")] str -> num
e.g. deststring z4, replace force (junk values become missing)
```

tostring

tostring varlist, {replace|generate(stub)} [force] [format(%fmt)] num -> str

codebook

codebook [varlist] type, uniques, missing, range

import

import delimited|excel|ods using FILE [, sheet(name) clear] read CSV/TSV/XLSX/ODS
e.g. import delimited "data/WEO.csv", clear

export

export delimited using FILE write CSV/TSV
e.g. export delimited using out.csv

save

save FILE [, replace] write Stata .dta (default) or .tea
e.g. save mydata.dta, replace - emits Stata-compatible .dta
e.g. save mydata.tea, replace - native tea binary

outreg2

outreg2 using FILE [, replace|append ctitle() dec() bdec() se label
symbol() alpha() addstat("Name", expr, ...) addtext() addnote()]
regression-table exporter (tab-separated; opens in Excel)

which

which CMD - report whether CMD is a built-in tea command

ssc

ssc install PKG - accepted and skipped (no package system)

isid

isid varlist - error 459 unless varlist uniquely identifies obs

duplicates

duplicates report|drop [varlist]

tempfile

tempfile NAME... - set local macros to fresh temp-file paths

tempname

tempname NAME... - set local macros to fresh scratch names

pwcorr

pwcorr varlist - pairwise correlation matrix

file

file open H using F, write [replace|append] | file write H "... " _n | file close H

confirm

```
confirm [new] file FILENAME | confirm [new|numeric|string] variable NAME
error (601/602/111/110/7) unless the condition holds; use with
capture: capture confirm file f.dta
         if _rc { <create it> }
```

history

```
history [N] | history save FILE [, replace] | history clear
list, export, or clear this session's interactive commands
(native REPL arrow-key history persists in ~/.tea_history;
the browser edition persists history in the browser)
```

sysuse

```
sysuse dir | sysuse NAME [, clear]
load a practice dataset bundled inside the tea binary
e.g. sysuse grunfeld, clear
     xtset firm year
     xtreg invest value capital, fe
```

use

```
use FILE [, clear] read Stata .dta, .tea, .csv, or .tsv
e.g. use mydata.dta, clear - extension dispatch decides format
'use foo' with no extension looks for foo.dta (Stata default).
```

clear

clear drop all data in current frame

status

status one-line summary: source, obs, vars, memory, sort/xtset

frame

```
frame create|change|copy|rename|put|drop|dir multiple datasets
e.g. frame create alt | frame change alt | frame put x y, into(alt)
```

regress

```
regress y x1 x2 ... [if] [in] [weight] [, noconstant robust cluster(var)] OLS
e.g. regress gdp_growth investment trade if year>=2000, cluster(country)
```

predict

predict newvar [, xb residuals]

predict from last fit

test

test v1 v2 ...

Wald F-test (joint zero)

lincom

lincom <linear combo of coefs>

point estimate + SE of L'b

help

help [cmd]

list commands or show one's syntax

pwd

pwd

print working directory

cd

cd DIR

change working directory

mkdir

mkdir DIR [, recursive]

create directory (recursive = mkdir -p)

dir

dir [pattern]

list files in current directory

rmdir

rmdir DIR

remove empty directory

erase

erase FILE | rm FILE

delete a file

copy

copy SRC DST [, replace]

copy a file

do

do FILENAME

run another do-file

preserve

preserve

snapshot the data in memory

restore

restore [, not preserve]

bring the preserve snapshot back

version

```
version                                tea version & build info
```

log

```
log using FILE [, replace append] | log close    tee output to a file
```

exit

```
exit [, clear]           leave tea (clear drops data first)
```

Native statements

Handled by the interpreter before command dispatch: `display`, `assert`, `shell` (and the `!cmd` escape), `#delimit`, `local`, `global`, `foreach`, `forvalues`, `while`, `if/else`, `capture`, `quietly`, `program` `define`, and comments (`*`, `//`, `///`).